

BEYOND INDIVIDUAL CLAIMS: PROBABILISTIC LOGICAL COHERENCE FOR LONG-FORM LLM EVALUATION

Anonymous authors

Paper under double-blind review

ABSTRACT

Long-form factuality assessment decomposes a response into atomic claims and verifies each against retrieved evidence. The paradigm is effective and now standard, but it embeds an assumption — that the truth of each claim can be settled in isolation — which discards the structure that makes prose an argument. A response can therefore pass per-claim verification while contradicting itself, and two responses built from *identical* claims can differ in coherence because of the order in which they assert them. We formalize *logical coherence* as a property of the joint distribution over claim truth values induced by a Markov random field whose pairwise factors are the relations the response itself draws. The relation vocabulary is a two-level taxonomy: twelve interpretable discourse senses, grounded in PDTB 3.0 and RST, that compile through a deterministic map onto five inferential couplings, each of which fixes a factor table by naming the joint assignment it penalizes. We characterize the factor tables coherence requires by a row-stochasticity property and show that the tables used for claim–evidence factuality are wrong here, collapsing the measure’s dynamic range. We then define four candidate coherence readouts over the model — the mean posterior marginal, a two-term consistency probability, a reified coherence node, and a normalized log-partition — give the exact auxiliary-variable construction each requires, and evaluate all four against five stated criteria by exact inference. Only two are monotone in conflict resolution, and the two that fail do so for one shared and instructive reason: they measure whether a conflict is *active* rather than what the reader should *believe*. A third natural candidate is worse still, since a score asking whether asserted relations are honoured is maximized by asserting nothing. Claim priors come from an existing factuality pipeline, making coherence and factuality two stages over one decomposition rather than competing scores. Empirically, on a corpus of response families whose relative coherence ordering is fixed by construction, the readouts recover 82.7% of the declared orderings from gold relation graphs and 34.7–37.1% from graphs mined out of the response text; since the two conditions differ only in the relation graph, the gap localizes the open problem in relation extraction rather than in the probabilistic model.

1 INTRODUCTION

The dominant paradigm for assessing the factuality of long-form generation is decompose-then-verify: split a response into atomic claims, check each against retrieved evidence, and aggregate the verdicts (Min et al., 2023; Wei et al., 2024; Song et al., 2024). It is effective, it is now standard, and it embeds an assumption. Treating each claim as independently checkable discards precisely the structure that distinguishes an *argument* from a *list*.

Consider two failure modes that per-claim verification cannot see. A report may assert both that no passengers were harmed in an incident and that a wire service reported three deaths. One of these claims is true and one is false, and per-claim verification will say so — but it will not register the more basic defect, that the response contradicts *itself*, asserting two things that cannot both hold. Separately, the inferential links a response draws may fail even when their endpoints are individually true: a chain of accurate statements joined by a spurious “therefore” is not a sound argument, and every atom in it verifies. Xie et al. (2026) quantify how often this matters: on logic-related samples,

SAFE, VeriScore and FactCheck-GPT omit the logical dimension on 60%, 78% and 48% of cases respectively.

Factuality and coherence are orthogonal axes. A response can be a bag of true facts that is incoherent, and it can be internally impeccable while being entirely false. The sharpest demonstration in our fixtures is a pair of summaries of one appellate judgment, built from the *same* set of claims — one summary’s atom list is a permutation of the other’s — and differing only in whether self-serving assertions are stated before or after the findings that qualify them. Every atom is identical, so every per-atom check returns an identical verdict; yet one summary is a faithful account and the other is not, and our measure separates them (0.95 against 0.90). No amount of grounding individual claims against evidence can make that distinction, because the distinction is not about the claims.

Coherence is a statement about a joint distribution. This paper measures the second axis, and the central modelling commitment is that the right object is a joint distribution rather than a tally. Given what a response asserts and how it links its assertions, we ask: how much of what it asserts does the argument, taken as a whole, uphold? That question is answered by a *marginal* in a probabilistic graphical model whose variables are claim truth values and whose factors are the asserted relations — not by counting how many rules a configuration violates. The distinction is not stylistic. A conflict between two claims should propagate: it should depress belief in both endpoints, and through them in whatever those endpoints support. Counting violated rules cannot propagate, and as we show in §5, scoring rule satisfaction directly produces a measure that a response can maximize by asserting no relations at all.

Our contributions are as follows.

- **A coherence MRF and a formal two-level relation taxonomy** (§4). Twelve discourse senses compile through a deterministic map onto five inferential couplings, each defined by the joint assignment it penalizes, which fixes its factor table. Separating the interpretable layer from the inferential one lets the annotation vocabulary be rich while the probabilistic model stays a closed three-parameter pairwise family that no relation type can extend.
- **A characterization of why factuality’s factor tables are wrong for coherence** (§4.3). The tables that are correct for claim–evidence factuality charge a premise for being a premise. We characterize the tables coherence requires by a row-stochasticity property, and show the alternative collapses the measure: the root of an entailment chain falls to $\mathbb{P}(a_1=1) = 0.19$ and the mean marginal moves only from 0.490 to 0.498 when every conflict edge is deleted, against 0.607 to 0.620 for the correct tables.
- **Four readouts, their exact constructions, and a selection argument** (§5). We give the deterministic auxiliary-variable constructions that make two of them tractable without a $2^{|\mathcal{R}|}$ table, evaluate all four by exact inference across a coherence-ordered family, and establish why the mean posterior marginal is the one to report: it alone satisfies all five criteria, and the two failures share a diagnosable cause (Prop. 3).
- **Two-stage composition with factuality** (§6), taking posterior marginals from an existing factuality pipeline as the coherence model’s priors, so one decomposition and one inference stack serve both axes and either can be read alone.
- **Measurements** (§8) that separate relation-extraction error from scoring error by holding the readouts fixed and swapping only the relation graph, plus a worked incoherent/coherent pair with exact values (§10) and a Markov-logic correspondence establishing what the pairwise encoding does and does not buy (Appendix C).

2 RELATED WORK

Atomic factuality. FActScore (Min et al., 2023) introduced per-atom precision against a knowledge source, decomposing a biography into atomic facts and reporting the fraction supported. SAFE (Wei et al., 2024) added search-augmented verification with a multi-step reasoning agent per claim and the $F_1@K$ metric that trades precision against a recall target. VeriScore (Song et al., 2024) restricted attention to *verifiable* claims, on the argument that many decomposed atoms are not checkable in principle. FactVerify (Bayat et al., 2023) reasons over retrieved evidence with an undecided verdict

class. All four aggregate independent per-claim verdicts, and all four are therefore blind to inter-claim structure by construction.

FactReasoner (Marinescu et al., 2025) replaced verdict counting with probabilistic inference: it builds a factor graph over claim and evidence variables with pairwise factors derived from natural-language-inference relations, and reports posterior marginals rather than a tally. This is the closest antecedent of our model, and we reuse its decomposition, its relation extractor and its inference stack (§3), and take its marginals as our priors (§6). The question it asks is nonetheless different: its factors run between a claim and the *evidence* retrieved for it, so its graph has no claim–claim edges and its marginals cannot express internal consistency. Coherence is a different question asked of the same decomposition.

The logic gap. Xie et al. (2026) states the problem we address: decompose-then-verify pipelines verify atoms individually and overlook the logical dependencies among them, so a text combining true facts through a false dependency is misjudged as factual. Their remedy classifies dependencies as a labelling task over claim pairs. Our departure is to treat the dependencies as *factors in a joint distribution*, which is what allows a local conflict to propagate to distant claims and what makes the resulting score a functional of the whole graph rather than a per-edge average.

Discourse coherence. Entity-grid models (Barzilay & Lapata, 2008) score coherence from the distribution of entity mentions across sentences, and their neural successors (Li & Jurafsky, 2017) learn the same signal with distributed representations. Both are trained and evaluated against *synthetic perturbations* — sentence shuffling and, in later work, contradiction injection — a methodology we extend in §7 by making the perturbation an operator on a declared relation graph rather than on raw text, so its intended effect on the score is known per item. DiscoScore (Zhao et al., 2023) scores generated text using BERT-based discourse representations and correlates well with human coherence judgements. What these approaches share is that coherence is read off surface statistics or learned representations; none exposes an inspectable relation graph, and none yields per-claim attributions. Our relation inventory takes its senses from PDTB 3.0 (Prasad et al., 2019) and RST (Mann & Thompson, 1988), but treats them as an interpretable layer that compiles to a smaller inferential vocabulary — which is what makes discourse senses usable as factor types at all.

Consistency checking and LLM judges. SelfCheckGPT (Manakul et al., 2023) detects hallucination by sampling several responses and running pairwise NLI between them; its NLI variant is the natural ablation isolating what global inference adds over purely local contradiction detection, and we propose it as such in §7.1. G-Eval (Liu et al., 2023) is the practical incumbent: an LLM judge scoring a coherence dimension directly from a rubric. Judges of this kind are strong on aggregate correlation and weak on exactly what a graphical model provides — attribution to specific claims, and a guarantee that a meaning-preserving edit cannot change the score.

Relational probabilistic models. Markov logic (Richardson & Domingos, 2006) is the natural alternative encoding, replacing factor tables with weighted first-order rules. Appendix C works out the correspondence and finds it instructive in both directions: the obvious one-clause-per-relation encoding is *not* equivalent to our model and is degenerate at its mode, while a three-clause expansion reproduces it exactly on the three-coupling fragment. We retain the pairwise MRF for the measure because its factor family is closed, and retain the Markov logic view because the beyond-pairwise rules that would address our main modelling limitation are expressible there.

3 STAGE ONE: PROBABILITIES FOR INDIVIDUAL CLAIMS

The pipeline in one pass. Figure 1 is the whole measure, and reading it left to right fixes what each stage owns. A response y is atomized *once*, into the claim set both stages then share. Stage 1 takes that set outward: it retrieves evidence $\mathcal{C}$ for each claim, labels claim–evidence and evidence–evidence pairs by NLI, and runs inference over a factor graph of claim *and* evidence variables, returning a posterior $q_i = \mathbb{P}(a_i=1 \mid \mathcal{C})$ per claim. Stage 2 takes the same set inward: it mines relations *between* claims from the response text — each arriving already typed by an inferential coupling, so nothing intervenes between mining and model construction — and builds a second MRF over the claims alone, one pairwise factor per relation. The two stages meet at exactly one place, the dashed arrow:




Figure 1: The two-stage pipeline. Stage 1 (§3) is FactReasoner (Marinescu et al., 2025): it grounds each atom against retrieved evidence and returns a posterior marginal q_i . Stage 2 (§4) builds a second MRF over the atoms *alone* and adds one pairwise factor per mined atom–atom relation; the mined relation already carries its Level-1 coupling, so no separate compile step intervenes. The dashed arrow is the only coupling between the stages, and carries the composition $\pi_i \leftarrow q_i$ into the $\phi_i = [1 - \pi_i, \pi_i]$ of Def. 1. Atomization is shared, so the response is decomposed once and both stages score the same claim set; setting $\pi_i \equiv \frac{1}{2}$ instead recovers a pure coherence score that consults no external evidence (Rem. 3).

Stage 1’s posterior becomes Stage 2’s prior, $\pi_i \leftarrow q_i$, occupying the free ϕ_i slot of Def. 1. Everything else runs one way.

That single point of contact is what buys the properties we rely on later. Because atomization is shared, a factuality verdict and a coherence verdict are statements about the same objects, and the composition is a substitution into an existing slot rather than a new model (§6). Because the coupling is only the priors, the axes stay separable in both directions: Stage 1 remains independently reportable, and setting $\pi_i \equiv \frac{1}{2}$ recovers a coherence score that consults no evidence at all (Rem. 3) — which is the condition under which every evidence-free comparison of §8 is run. And because priors are applied when the network is scored rather than by mutating the mined graph, one mined graph can be rescored under uniform, factuality-derived or ablated priors without re-running the estimator, making the prior ablations essentially free.

The coherence model consumes claim-level probabilities and reuses the inference machinery that produces them, so we describe that stage before defining the model itself. It is FactReasoner (Marinescu et al., 2025); the account here is a summary sufficient to make the composition of §6 precise, and it also fixes notation and states the estimator whose label-span construction the relation miner reuses in §4.5.

3.1 DECOMPOSITION, RETRIEVAL, AND RELATION EXTRACTION

A response y is decomposed by an LLM *atomizer* into atomic claims $\mathcal{A} = \{a_1, \dots, a_n\}$, each a single independently verifiable proposition. A *reviser* then decontextualizes each one, resolving pronouns and anaphora so that the claim is checkable outside its original sentence; near-duplicates are collapsed. For each claim a retriever returns supporting passages, which are deduplicated into an evidence set $\mathcal{C} = \{c_1, \dots, c_m\}$.

An LLM-based NLI extractor then labels admitted pairs. For a claim–evidence pair the label is drawn from {entailment, contradiction, neutral}; for an evidence–evidence pair both directions are scored and reconciled, with entailment in both directions relabelled *equivalence*. Neutral verdicts are dropped rather than encoded: a claim about which the evidence is silent contributes no factor and is left at its prior, which is what makes the factor set of §4 contain only the relations someone actually asserted.

3.2 READING A PROBABILITY OFF THE GENERATION

Each retained relation carries a probability p^* . Rather than asking the model to verbalize a confidence, we read it from the distribution over the tokens it actually emitted. The extractor is prompted to end with a bracketed label; let T be the set of generated tokens overlapping the span of that label in the

decoded text. Then

$$p^* = \exp\left(\frac{1}{|T|} \sum_{t \in T} \log p_t\right), \quad (1)$$

the per-token geometric-mean probability of the label the model produced. Three details of this construction are load-bearing, and we record them because each was a source of silent error.

Span alignment. The probability is measured over exactly the span the label *parser* reports, reconstructing the decoded string from the token list and tracking each token’s character offsets. Measuring the same span the label came from makes the label and its probability structurally unable to disagree. Requiring standalone bracket-delimiter tokens instead does not work: tokenizers routinely fuse [,] and " into subwords with adjacent characters, so a delimiter-based rule silently fails on some backends and not others.

Token-array conventions. An earlier version of this code stripped the final entry of the logprob array as an end-of-sequence token. OpenAI-compatible and vLLM logprob arrays contain only emitted content tokens — the stop is signalled out of band — so dropping the last entry deleted a real content token, and for a bracketed label that token is the one closing the label whose confidence is being measured.

Degenerate cases. When the logprob array is empty, no label span is found, or no token aligns to it, the estimator returns 0.5 rather than 0.0. Returning zero would be read downstream as a *degenerate*, *impossible* relation for a label the model in fact generated, which is a stronger claim than “we could not measure the confidence”.

Backends without logprobs. Some serving stacks expose no token logprobs at all. There p^* instead comes from SIMBA-UQ (Bhattacharjya et al., 2025), which samples the extractor repeatedly across a temperature schedule, builds a pairwise similarity matrix over the samples, and aggregates each sample’s similarity row into a self-consistency confidence. The winning sample supplies the label and its aggregated confidence supplies p^* , so the downstream model is unchanged.

3.3 THE FACTUALITY NETWORK AND ITS VARIANTS

The labelled graph becomes a pairwise MRF with one binary variable per claim and per evidence passage. Unary factors are $\phi_X(x) = [1 - \pi_X, \pi_X]$ with π_X the prior on X ; each retained relation contributes one pairwise factor from Table 2 below, oriented so the evidence passage is the premise. Marginals $\mathbb{P}(A_i=1)$ are computed by weighted mini-bucket elimination (Liu & Ihler, 2011) with i-bound 6, and a claim is reported supported when $\mathbb{P}(A_i=1) > \frac{1}{2}$.

Three graph variants differ in which factors are admitted: FR1 connects each claim only to the passages retrieved for it; FR2 deduplicates passages and scores every claim against all of them; FR3 adds evidence–evidence factors. We use FR2. Because scoring every claim against every passage is $O(nm)$ LLM calls, a candidate-pair policy can restrict the pairs actually scored; the calibration of that policy is instructive for the analogous choice in §4.5. Replaying policies against recorded verdicts on a 20-claim narrative with 60 passages, a provenance-restricted policy with an embedding similarity gate holds recall at 1.000 through a threshold of 0.20, slipping to 0.985 at 0.30 and 0.833 at 0.40. Substituting token-Jaccard similarity for sentence embeddings, however, loses 22 of 72 non-neutral relations and moves the factuality score by 0.05 — and *no* threshold recovers full recall, even 0.05. Cheap surrogates for semantic similarity are not safe here, and the saving is workload-dependent rather than constant: roughly $5\times$ on unrelated subtopics but $1.2\times$ on a narrative where every claim shares characters.

Conventions. We use the with-priors factor tables (Table 2) throughout, with $\pi_C = 0.9$ for evidence and $\pi_A = 0.5$ for claims. The published factuality model uses the no-priors tables and a higher evidence prior. The difference is immaterial to the coherence results below, which set claim priors explicitly, but the convention must be stated because the two table families are not interchangeable — that is the subject of §4.3.

4 STAGE TWO: THE COHERENCE MARKOV RANDOM FIELD

4.1 THE MODEL

Definition 1 (Coherence MRF). Let $\mathcal{A} = \langle a_1, \dots, a_n \rangle$ be the atomic claims of a response in assertion order, each a binary variable with $a_i = 1$ read as “the claim holds”. Let $\mathcal{R} = \{(s_r, t_r, \tau_r, p_r)\}_{r=1}^{|\mathcal{R}|}$ be a set of relations, where $s_r \neq t_r$ index claims, τ_r is an inferential coupling (Def. 2), and $p_r \in [0, 1]$ is the relation’s weight. Given priors $\pi \in [0, 1]^n$, the coherence MRF induced by $(\mathcal{A}, \mathcal{R}, \pi)$ is

$$\mathbb{P}(\mathbf{a}) = \frac{1}{Z} \prod_{i=1}^n \phi_i(a_i) \prod_{r \in \mathcal{R}} \psi_{\tau_r, p_r}(a_{s_r}, a_{t_r}), \quad Z = \sum_{\mathbf{a} \in \{0,1\}^n} \prod_i \phi_i(a_i) \prod_r \psi_r, \quad (2)$$

where $\phi_i(a_i) = [1 - \pi_i, \pi_i]$ over the states $(0, 1)$ and $\psi_{\tau, p}$ is given by Table 2.

Three structural properties of Definition 1 deserve statement, because each is a deliberate commitment rather than an implementation detail.

Remark 1 (The factor family is closed). A coupling is nothing more than a choice of which cells of a 2×2 table are down-weighted and by how much. Every $\psi_{\tau, p}$ in Table 2 is determined by p , the source prior π_s , and the penalized set; the whole inferential vocabulary is therefore a three-parameter pairwise family. This is a guarantee about what the model cannot do: no relation type, however elaborately named at the discourse level, can smuggle in expressiveness beyond a pairwise potential. Appendix C contrasts this with Markov logic, whose rule vocabulary is open by design.

Remark 2 (Order-sensitivity enters only through the relations). Nothing in Eq. 2 references claim position. Assertion order enters exclusively through which pairs are related and in which direction. It follows that an edit which preserves the relation set — reordering claims whose relations are symmetric, or exchanging a temporal **Precedence** for a **Succession**, which Table 4 maps to no factor at all — provably cannot change any readout. This is a falsifiable prediction about the measure and not merely a design note; §7 describes the ladder built to test it.

Remark 3 (Coherence is evaluated without evidence). The graph of Definition 1 contains claim variables only. No retrieved passage appears, and with uniform priors $\pi_i \equiv \frac{1}{2}$ the model consults no external knowledge whatever: it scores how the response hangs together, not whether it is true. This is what makes the two axes separable, and §6 shows how to couple them when a joint reading is wanted.

4.2 LEVEL ONE: INFERENTIAL COUPLINGS

Definition 2 (Inferential couplings). A coupling τ is specified by the set $V(\tau) \subseteq \{0, 1\}^2$ of joint assignments (a_s, a_t) it penalizes — equivalently, the configurations at which the relation is violated. The vocabulary is:

Coupling	Reading	Directed	Penalized set $V(\tau)$
ENTAILMENT	$a_s \Rightarrow a_t$	yes	$\{(1, 0)\}$
CONTRADICTION	$a_s \Rightarrow \neg a_t$	yes	$\{(1, 1)\}$
EQUIVALENCE	$a_s \Leftrightarrow a_t$	no	$\{(0, 1), (1, 0)\}$
EXCLUSIVE	exactly one of a_s, a_t holds	no	$\{(0, 0), (1, 1)\}$
CO_NECESSITY	at least one of a_s, a_t holds	no	$\{(0, 0)\}$
NONE	no coupling	—	$\emptyset$ (no factor)

Write $\mathcal{T}$ for the five edge-producing couplings and $\mathcal{T}_{\text{conf}} = \{\text{CONTRADICTION}, \text{EXCLUSIVE}\}$ for the conflict couplings, those whose penalized set contains the both-true world $(1, 1)$.

Table 1 gives a response excerpt for each coupling, set against the world that coupling denies; since a coupling is its penalized set, reading the two together is the quickest way to see what each one commits the response to. The first three couplings are exactly the NLI edge types of §3, which is what lets a mined relation reuse the existing graph machinery unchanged. The last two are additions required by argumentative prose, and the distinction between them and CONTRADICTION is the one the taxonomy exists to draw.

Table 1: One realistic response excerpt per Level-1 coupling of Def. 2. As in Table 5, each excerpt is prose containing both claims — A is the source, underlined, B **the target, bold**, the *italic* words drawing the link — and $V(\tau)$ is repeated from Def. 2 so the excerpt can be read against the world it rules out. The *Why* column names the world the response denies, which is what distinguishes each coupling from the one above it; the three conflict-adjacent couplings CONTRADICTION, EXCLUSIVE and CO_NECESSITY are illustrated on one two-part weld failure so that only exhaustiveness varies. NONE produces no factor, so its excerpt is one where the text relates the claims without constraining their truth.

Coupling	Reading	$V(\tau)$	Response excerpt, and the world it denies
ENTAILMENT	$a_s \Rightarrow a_t$	$\{(1, 0)\}$	“The alloy is chemically identical to the certified <u>reference</u>, and therefore it meets the airframe specification.” <i>Why</i>: the response rules out the world where the source holds and the target fails — $(1, 0)$ — and says nothing about the case where the alloy differs, which is why the $a_s=0$ row returns the target to its prior.
CONTRADICTION	$a_s \Rightarrow \neg a_t$	$\{(1, 1)\}$	“The weld passed the tensile test, yet the same joint was porous along its full length.” <i>Why</i>: only the both-true world $(1, 1)$ is denied. Both can still fail together — if the wrong joint was tested, neither claim holds — so $(0, 0)$ must stay available.
EXCLUSIVE	exactly one holds	$\{(0, 0), (1, 1)\}$	“The joint was destructively examined, so <i>exactly one of the weld met the tensile minimum and the weld fell short of it must be the finding</i> — there is no third outcome.” <i>Why</i>: both same-value worlds are denied. The added $(0, 0)$ is the whole difference from CONTRADICTION: once the joint has actually been examined the pair cannot both fail, so leaving that world open — as CONTRADICTION does — would under-penalize the conflict.
CO_NECESSITY	at least one holds	$\{(0, 0)\}$	“The fracture began at the joint, so <i>at least one of the weld and the fastener must have given way</i> — conceivably both.” <i>Why</i>: the dual of EXCLUSIVE. Only the both-false world is denied, so asserting the disjunction is evidence for both endpoints (Example 1).
EQUIVALENCE	$a_s \Leftrightarrow a_t$	$\{(0, 1), (1, 0)\}$	“The shipment was certified to specification — <i>that is to say, every contractual requirement had been met</i>.” <i>Why</i>: the two disagreeing worlds are denied, in both directions; unlike ENTAILMENT there is no vacuous row, since a false source now also forces a false target.
NONE	no coupling	$\emptyset$	“The replacement units were installed in March, <i>some three months before the first in-flight failures were reported</i>.” <i>Why</i>: no world is denied. The response genuinely relates the claims, but ordering them in time constrains neither one’s truth, so the pair contributes no factor and both claims stay at their priors.

Table 2: Pairwise potentials $\psi_{\tau,p}(a_s, a_t)$, listed row-major over the assignments $(0,0), (0,1), (1,0), (1,1)$; π_s is the source claim’s prior and bold marks the penalized set $V(\tau)$ of Def. 2. The *no-priors* family is correct for claim–evidence factuality, where the source passage’s own truth is at stake; the *with-priors* family is what coherence requires (Prop. 1) and what Def. 1 uses. The two symmetric couplings EQUIVALENCE and EXCLUSIVE coincide in both families, having no source-prior term to correct.

Coupling τ	No priors	With priors	Symmetry
ENTAILMENT	$[p, p, \mathbf{1-p}, p]$	$[1-\pi_s, \pi_s, \mathbf{1-p}, p]$	asymmetric
CONTRADICTION	$[p, p, p, \mathbf{1-p}]$	$[1-\pi_s, \pi_s, p, \mathbf{1-p}]$	asymmetric
EQUIVALENCE	$[p, \mathbf{1-p}, \mathbf{1-p}, p]$	$[p, \mathbf{1-p}, \mathbf{1-p}, p]$	symmetric
EXCLUSIVE	$[\mathbf{1-p}, p, p, \mathbf{1-p}]$	$[\mathbf{1-p}, p, p, \mathbf{1-p}]$	symmetric
CO_NECESSITY	$[\mathbf{1-p}, p, p, p]$	$[\mathbf{1-p}, \pi_s, \pi_s, p]$	symmetric

Why exhaustiveness needs its own coupling. CONTRADICTION forbids only the both-true world: two claims that cannot both hold, but might both fail. EXCLUSIVE additionally forbids the both-false world, which is what *exhaustive* alternatives assert. “No one was harmed” and “three people died” can neither both hold nor both fail — either someone was harmed or no one was — so a model that treats them as mere opposition has mis-described the situation in a way that matters for scoring: it leaves available a world in which *neither* claim is true, and so under-penalizes. Conversely CO_NECESSITY is the dual, penalizing only the both-false world, and is what a disjunction asserts (“at least one of these two findings must hold”). Adjudicated disputes — incident investigations, litigation, competing hypotheses — are precisely the domains where exhaustive pairs occur, because the function of a finding is to close one. The CONTRADICTION, EXCLUSIVE and CO_NECESSITY rows of Table 1 state one two-part weld failure three times over, so that the only thing varying across them is which same-value worlds the response leaves standing.

Example 1 (A CO_NECESSITY factor in isolation). *Take a single CO_NECESSITY relation with $p = 0.9$ over two claims with priors $\pi = 0.5$, and no other factors. From Table 2 the pairwise table is $[1-p, \pi_s, \pi_s, p] = [0.1, 0.5, 0.5, 0.9]$, and multiplying in the unaries and normalizing gives world masses*

$$\mathbb{P}(0,0) = 0.05, \quad \mathbb{P}(0,1) = 0.25, \quad \mathbb{P}(1,0) = 0.25, \quad \mathbb{P}(1,1) = 0.45.$$

The both-false world is suppressed from the 0.25 it would receive under independent fair priors to 0.05, while the three worlds satisfying “at least one holds” share the recovered mass in proportion to how many endpoints they assert. Each marginal rises from 0.5 to 0.7: asserting that at least one of two claims holds is evidence for both, which is the intended reading.

4.3 THE PAIRWISE POTENTIALS, AND WHY FACTUALITY’S TABLES ARE WRONG HERE

Proposition 1 (Row-stochasticity of the asymmetric couplings). *The with-priors tables of the two asymmetric couplings are row-stochastic in the source state: for $\tau \in \{\text{ENTAILMENT}, \text{CONTRADICTION}\}$ each row of $\psi_{\tau,p}$ sums to one, so the row indexed by a_s is the conditional distribution $\mathbb{P}(a_t | a_s)$, with*

$$\mathbb{P}(a_t=1 | a_s=0) = \pi_s, \quad \mathbb{P}(a_t=1 | a_s=1) = \begin{cases} p & \tau = \text{ENTAILMENT} \\ 1-p & \tau = \text{CONTRADICTION}. \end{cases}$$

Consequently such a factor constrains the target only when the source is true; when the source is false it returns the target to the prior. The symmetric couplings are not of this form, and are not claimed to be. In particular CO_NECESSITY is not row-stochastic: its rows sum to $(1-p) + \pi_s$ and $\pi_s + p$, which both equal one only in the degenerate case $p = \pi_s = \frac{1}{2}$.

Proof. For both asymmetric tables the $a_s=0$ row is $[1-\pi_s, \pi_s]$, which sums to one; the $a_s=1$ row is $[1-p, p]$ for ENTAILMENT and $[p, 1-p]$ for CONTRADICTION, each summing to one. Dividing a row by its sum is therefore the identity, so the rows already are conditionals and may be read off directly. For CO_NECESSITY $= [1-p, \pi_s, \pi_s, p]$, row-stochasticity would require $(1-p) + \pi_s = 1$ and $\pi_s + p = 1$, i.e. $\pi_s = p$ and $\pi_s = 1-p$ simultaneously, hence $p = 1-p$. $\square$

What row-stochasticity means, and what it does not. The condition is a statement about the *table*: it says each row of ψ , read across the target’s two states, is already normalized, so $\psi(a_s, \cdot)$ can be read as $\mathbb{P}(a_t | a_s)$ rather than as an unnormalized compatibility score. Two consequences follow, and a third does not. First, the $a_s=0$ row being the prior $[1 - \pi_s, \pi_s]$ is what makes a false source *uninformative* about the target, which is the formal content of “an implication says nothing when its antecedent fails”. Second, writing $\ln \psi = \alpha + \beta a_s + \gamma a_t + \delta a_s a_t$, row-stochasticity *pins* the source-only weight: $\beta = \ln \frac{1-p}{1-\pi_s}$ for ENTAILMENT, with the other couplings tabulated in App. C. The source term is not zero, but neither is it free — it is exactly the amount needed to hold the $a_s=0$ row at the prior, so no residual degree of freedom is left with which the relation could bid the source’s own plausibility up or down. That is what distinguishes these tables from the no-priors family below, whose $(0, \cdot)$ row is $[p, p]$ and so trades the source’s truth against the target’s. What does *not* follow is any invariance of the marginals in the assembled model: once a factor is multiplied against its neighbours and renormalized, both endpoints can move, and for an isolated ENTAILMENT relation the source’s own marginal already departs from π_s whenever $\pi_t \neq \frac{1}{2}$. Row-stochasticity constrains how a relation *enters* the joint, not where inference then lands.

Table 3: Proposition 1 made concrete, at $p = 0.9$ and $\pi_s = 0.5$ so the three couplings are directly comparable. Tables are row-major over $(0, 0), (0, 1), (1, 0), (1, 1)$, so the first pair is the $a_s=0$ row and the second the $a_s=1$ row. The two asymmetric couplings have both rows summing to one and so admit the conditional reading; CO_NECESSITY does not, and its rows are the entry marked in bold. Every value is produced by `fact_reasoner.factors.edge_factor_values`.

τ	Table	Row sums		Conditional reading		Reading
		$a_s=0$	$a_s=1$	$\mathbb{P}(a_t=1 a_s=0)$	$\mathbb{P}(a_t=1 a_s=1)$	
ENTAILMENT	[.5, .5, .1, .9]	1.0	1.0	$0.5 = \pi_s$	$0.9 = p$	source true $\Rightarrow$ target up
CONTRADICTION	[.5, .5, .9, .1]	1.0	1.0	$0.5 = \pi_s$	$0.1 = 1 - p$	source true $\Rightarrow$ target down
CO_NECESSITY	[.1, .5, .5, .9]	0.6	1.4	0.833	0.643	source <i>false</i> $\Rightarrow$ target up

Table 3 states the proposition in numbers, and the three rows are worth reading against the prose they encode. For ENTAILMENT — “Voyager 1 crossed the heliopause, so its instruments measured the density rise” — a true source pulls the target up to $p = 0.9$, while a false source leaves it at the 0.5 prior: the row is flat, which is exactly a vacuous antecedent. For CONTRADICTION — “the probe is in interstellar space, so it has *not* left the Solar System” — a true source pushes the target down to $1 - p = 0.1$, and a false source is again silent. Both rows normalize, so both readings are conditional probabilities.

CO_NECESSITY inverts the direction of information and thereby breaks the property. Reading the weld example of Table 1 — “at least one of the weld and the fastener must have given way” — it is the weld *holding* that forces the fastener to have failed, so the informative row is $a_s=0$, where the reading is 0.833, well above the 0.5 prior. The $a_s=1$ row is the weak one: once the weld has failed the disjunction is already satisfied and the fastener is left near its prior. A source-conditioned normalization cannot express this, since it reserves the flat row for the false source; the rows sum to 0.6 and 1.4 instead of to one. This is not a defect of the CO_NECESSITY table but a property of disjunction, and Example 1 shows the same fact from the joint side: both endpoints rise together, from 0.5 to 0.7, because the factor is symmetric and pushes neither endpoint preferentially.

Proposition 1 is the substantive modelling choice of this section, and the contrast with the no-priors family is worth being precise about. In the no-priors ENTAILMENT table the $(1, 0)$ cell is $1 - p$ and the $(0, \cdot)$ cells are both p , so the table charges the *source* being true whenever the target is uncertain. That is appropriate for claim–evidence factuality: there the source is a retrieved passage whose own reliability is in question, and an entailment from a passage to a claim the model doubts is evidence against the passage. It is wrong for coherence, where the source is another claim of the same response, and where a well-supported premise must not be penalized for being a premise.

The consequence is not a matter of degree. On the response of §10.3 — sixteen claims joined by a twelve-edge entailment backbone — the no-priors tables drive the root of the backbone to $\mathbb{P}(a_1=1) = 0.19$, far below its 0.5 prior, purely because it is the premise of a chain whose later members are uncertain. The mean marginal sits at 0.490 and moves only to 0.498 when every conflict edge is deleted, a range of 0.008. Under the with-priors tables the same comparison runs 0.607 to

Table 4: The compile map $\mathcal{C}$ of Def. 3. *Ordering-only* senses record assertion order but induce **no factor**, which is what makes an order-only edit provably score-invariant (Rem. 2). **Concession** is a conflict the text may itself resolve, handled by Eq. 3. The nine senses marked $\star$ form the admissible vocabulary of §4.5.

Level-2 sense	$\mathcal{C}$ coupling	prior	dir.	flag	gloss
Cause-Effect $\star$	ENTAILMENT	$\perp$	yes		source brings target about
Effect-Cause $\star$	ENTAILMENT	$\perp$	yes		source is the effect, target the cause
Evidence $\star$	ENTAILMENT	$\perp$	yes		source is evidence for target
Condition	ENTAILMENT	$\perp$	yes		source conditions target
Instantiation	ENTAILMENT	$\perp$	yes		target instantiates source
Restatement $\star$	EQUIVALENCE	0.90	no		definitional identity
Contrast $\star$	CONTRADICTION	$\perp$	no		opposed, <i>not</i> exhaustive
Concession $\star$	CONTRADICTION	$\perp$	yes	conc.	a tension the text resolves
Alternative $\star$	EXCLUSIVE	$\perp$	no		exhaustive: exactly one holds
Disjunction $\star$	CO_NECESSITY	$\perp$	no		at least one holds
Precedence $\star$	NONE	—	yes	ord.	temporal order only
Succession	NONE	—	yes	ord.	temporal order only

0.620: a range more than 1.5 times larger, from a base that is not already depressed by the backbone itself. All four values are exact, by 2^{16} enumeration. A coherence measure built on the factuality tables does not merely give different numbers; it spends its dynamic range penalizing premises and has little left for conflicts, which are the thing being measured.

4.4 LEVEL TWO: DISCOURSE SENSES AND THE COMPILE MAP

Couplings are the right vocabulary for a factor table and the wrong one for annotation. Asked whether two claims stand in ENTAILMENT, an annotator — human or model — has been handed a question about logic, when the fact actually observable in the text is a question about discourse: does this sentence give a cause, offer evidence, concede a tension, or lay out alternatives? We therefore mine an interpretable layer and compile it down.

Definition 3 (Compile map). *Let $\mathcal{S}$ be the set of Level-2 discourse senses. The compile map*

$$\mathcal{C} : \mathcal{S} \longrightarrow (\mathcal{T} \cup \{\text{NONE}\}) \times ([0, 1] \cup \{\perp\}) \times \{0, 1\}^3$$

assigns to each sense a coupling, a strength prior ($\perp$ meaning “use the estimator’s own value”), and the flags directed, concession and ordering-only. It is total, deterministic, and given in Table 4.

Table 5 grounds each sense in a response excerpt that contains both claims and the wording that licenses the link, since what an annotator must actually recognize is a form of words rather than a truth-functional pattern. Several entries encode judgements worth defending. **Restatement** carries a fixed strength prior of 0.90 because a marked paraphrase (“in other words”, “equivalently”) is a claim about identity of content, and a model asked to estimate its strength has little to add; the estimator’s value is used when available, and the prior seeds it otherwise. Five senses collapse onto ENTAILMENT, which is a deliberate loss: the model does not distinguish causation from evidence, because a pairwise potential over truth values cannot. The grouped ENTAILMENT block of Table 5 shows exactly what compilation discards — five distinct discourse facts, one factor table. What survives compilation is the inferential force, and the sense is retained alongside it for interpretability and for error analysis.

The Contrast/Alternative boundary is exactly the exhaustiveness question of §4.2, and the Contrast, Alternative and Disjunction rows of Table 5 are that question in concrete form: one two-part weld failure written up three ways — as mere opposition, as an exhaustive choice, and as a disjunction — differing in nothing but which same-value worlds the response leaves available. The two ordering-only senses are the taxonomy’s way of representing a relation the text genuinely draws while asserting that it has no truth-functional content: a narrative that says one event preceded another has related the two claims, and has said nothing about whether either is true.

Concession: a conflict the text resolves. A Concession asserts a tension and then overrides it. “The airline argued pilot error, but the final report found a metallurgical defect” is not an incoherence

Table 5: One realistic response excerpt per Level-2 sense of Table 4, in the same row order and grouped by the coupling each compiles to. Each excerpt is prose of the kind the estimator actually sees and *contains both claims*: *A* is underlined, *B* is **bold**, and the *italic* words are the ones that draw the link — the span §4.5 obliges the model to quote, which is why it must be findable in the response and not in either claim. Senses are marked * when admissible; the three unmarked ones are included because their breadth is the argument for excluding them (§4.5). The Contrast, Alternative and Disjunction rows describe one two-part weld failure three ways, separated only by exhaustiveness (§4.2); they reuse the scenario of Table 1 deliberately, so that a discourse sense and the coupling it compiles to can be compared on the same prose.

Level-2 sense	<i>C</i>	Response excerpt, and why this sense rather than a neighbour
Cause-Effect*	ENTAILMENT	“<u>The company shipped a batch it knew to be out of tolerance</u>, and <i>as a direct result</i> its share price fell 15% that quarter.” Why: the response itself presents <i>A</i> as bringing <i>B</i> about; the marker names the direction.
Effect-Cause*	ENTAILMENT	“<u>The entire fleet was grounded on 14 March</u>. The order came <i>because</i> inspectors had found a cracked wing spar on two airframes.” Why: the same support runs from the stated effect to its cause. Direction of <i>explanation</i> is reversed; direction of inferential support is not.
Evidence*	ENTAILMENT	“<u>The maintenance log records no inspection between January and June</u>, <i>which shows that</i> the mandatory 500-hour check was skipped.” Why: <i>A</i> is offered in support of <i>B</i> without any claim that it caused it — a record is evidence, not a mechanism.
Condition	ENTAILMENT	“A part is cleared for airframe use <i>only if</i> <u>its alloy meets the reference specification</u>.” Why: <i>B</i> is asserted contingent on <i>A</i>. Note how weakly the cue constrains: nearly any dependence can be reread as a condition, which is why the sense is inadmissible.
Instantiation	ENTAILMENT	“<u>Several suppliers failed the 2023 audit</u> — <i>for instance</i>, the Tucson plant failed on two counts.” Why: <i>B</i> is one case falling under <i>A</i>’s generalization. Like Condition, broad enough to fit almost any pair that shares a topic.
Restatement*	EQUIVALENCE	“<u>The shipment was certified to specification</u> — <i>in other words</i>, every requirement in the contract had been met.” Why: a <i>marked</i> paraphrase asserts identity of content, which is why the strength prior is fixed at 0.90 rather than estimated.
Contrast*	CONTRADICTION	“<u>The weld passed the tensile test</u>, <i>but</i> the same joint was porous along its full length.” Why: opposed and <i>not</i> exhaustive. A third possibility — the wrong joint was tested — leaves both false, so the both-false world must stay available.
Concession*	CONTRADICTION	“<i>Although</i> <u>the airline argued throughout that pilot error was to blame</u>, the final report identified a metallurgical defect; the tribunal <i>ultimately held</i> for the latter.” Why: a tension the text itself settles. The resolving holding triggers the discount of Eq. 3 rather than counting as a live conflict.
Alternative*	EXCLUSIVE	“The two accounts cannot be reconciled: <i>either</i> <u>no one aboard was harmed</u>, or three passengers died, and <i>exactly one of them is true</i>.” Why: exhaustive. Unlike Contrast the pair can neither both hold nor both fail, so the both-false world is penalized too.
Disjunction*	CO_NECESSITY	“Metallurgy confirms the fracture began at the joint, so <i>at least one of</i> <u>the weld</u> and the fastener must have failed — possibly both.” Why: at least one holds and both may. The dual of Alternative: only the both-false world is penalized.
Precedence*	NONE	“The replacement units were installed in March, <i>some three months before</i> the first in-flight failures were reported in June.” Why: the text genuinely relates the two claims, and says nothing about either one’s truth — so the sense induces no factor.
Succession	NONE	“The first in-flight failures were reported in June, <i>fully three months after</i> the replacement units had been installed.” Why: the mirror of Precedence — the same two claims, the same absence of truth-functional content. Offering both senses invites an arbitrary choice on a relation that produces no factor either way.

if the report’s finding is presented as authoritative; treating it as a live contradiction would penalize a response for doing the very thing careful argument requires. When a resolving *holding* claim h is present we discount the conflict’s weight,

$$p_r^{\text{eff}} = p_r (1 - \lambda \pi_h), \quad (3)$$

with $\lambda = 0.45$; absent a resolver the relation stands at full strength and the tension counts against the response. Holdings are identified by adjudicative cue phrases (*held, ruled, concluded, found that, the tribunal, ultimately*) on a claim that is not itself an endpoint of the relation and that maximizes lexical overlap with the two endpoints. This is a heuristic, and with λ it is the one place in the design where hand-set constants enter; §11 returns to the point.

4.5 ESTIMATING RELATIONS FROM THE RESPONSE

Definition 1 takes $\mathcal{R}$ as given. In use it must be estimated from the text, and the estimator is the weakest link in the pipeline (§8), so we specify it carefully.

The estimand, decomposed. A relation’s weight factors into a type term and a strength term:

$$p_r = \underbrace{\mathbb{P}(\tau \mid a_s, a_t, y)}_{\text{type confidence}} \cdot \underbrace{\mathbb{P}(a_t \mid a_s, \tau, y)}_{\text{conditional strength}}. \quad (4)$$

The first factor asks how confident we are that this is a τ -relation *at all*; the second, given that it is, how strongly the source determines the target. Conflating them is a mistake: a confidently-identified weak causal link and a doubtfully-identified strong one are different epistemic situations that a single elicited number cannot distinguish. Two prompts estimate them separately (full texts in Appendix A).

Prompt A elicits the discourse sense and its coupling, with chain-of-thought *preceding* the label so the model commits to a sense only after reasoning — and so that the label’s token logprobs, read exactly as in Eq. 1, measure a decision already made rather than one being made. *Prompt B* elicits the conditional strength as a surrogate token: the answer’s first word must be *Yes* or *No*, and the strength is the renormalized $\mathbb{P}(\text{Yes}) / (\mathbb{P}(\text{Yes}) + \mathbb{P}(\text{No}))$ (Kadavath et al., 2022), or the affirmative fraction over N samples where logprobs are unavailable. The question is deliberately graded — “is the implication at least plausible?” rather than “is it certain?” — so that a weak but real link is not forced to answer *No*; its weakness surfaces in the renormalized probability instead. A verbalized variant asking for [p=0 . NN] directly is retained only as a comparison, verbalized confidence being weakly calibrated (Xiong et al., 2024).

Grounding is mandatory, not optional. Prompt A always receives the full response, and is instructed to assert a coupling only where the response *itself* draws the connection. This is the single most consequential decision in the estimator. Judging a claim pair in isolation invites the model to accept any abstractly plausible relation, and the failure is not a mild over-generation: measured relation density reaches six to nine relations per claim, including spurious contradictions on paragraphs that contain none. Because Eq. 2 is a product over relations, a graph of many soft mutually-constraining edges drives every marginal toward the prior, so an over-connected graph does not merely add noise — it destroys the measure’s ability to distinguish a coherent response from an incoherent one, in the direction of calling everything mediocre.

Two further precision mechanisms are used, one requested and one checkable. The *admissibility filter* restricts output to the nine (sense, coupling) combinations marked $\star$ in Table 4. Condition and Instantiation are the notable exclusions: both compile to ENTAILMENT and both are semantically broad — their rows in Table 5 illustrate the breadth, since almost any dependence can be read as a condition and almost any topical pair as an instantiation — so a model reaching for “these two claims are related somehow” lands on them readily; they accounted for 13% of one measured arm’s edges while appearing zero times in gold. The *evidence-span requirement* obliges the model to quote, verbatim from the response, the words that draw the link; the parser then verifies that the quotation is actually a substring of the response and drops the relation otherwise. This is the only precision mechanism in the estimator that is *checked* rather than merely asked for. It must quote the response and not the claims, because claims are decontextualized rewrites and only about 27% of them appear in the response verbatim even after normalization.

Algorithm 1 SELECTCANDIDATEPAIRS: response-anchored candidate selection.

Require: claims $\mathcal{A} = \langle a_1, \dots, a_n \rangle$ in assertion order; response y ; policy; window w ; gate threshold θ_g ; discourse thresholds θ_d, δ

Ensure: ordered candidate pairs P , plus a coverage record

- 1: **if** policy = ALL_PAIRS **then return** every ordered pair $(a_i, a_j), i \neq j$
- 2: $W \leftarrow \{(a_i, a_j) : 0 < j - i \leq w\}$ ▷ forward order window
- 3: **if** policy = BIDIRECTIONAL **then** $W \leftarrow W \cup \{(a_j, a_i) : 0 < j - i \leq w\}$
- 4: $G \leftarrow \emptyset$
- 5: **if** policy = GATED **then**
- 6: $G \leftarrow \{(a_i, a_j) : j - i > w, \text{sim}(a_i, a_j) \geq \theta_g\}$
- 7: align each claim to its best-matching sentence of y
- 8: $Prom \leftarrow \{(a_i, a_j) \notin W \cup G : \text{DISCOURSEADJACENT}(i, j)\}$
- 9: $Dem \leftarrow \{(a_i, a_j) \in W : \neg \text{DISCOURSEADJACENT}(i, j)\}$
- 10: **return** (DEDUP($W \cup G \cup Prom$) $\setminus Dem$), ($|W|, |G|, |Prom|, |Dem|$)

Two prompt generations. The first-generation Prompt A over-asserts badly: against a corpus in which 7.2% of candidate pairs are genuinely related, it answers “related” on 72% of the pairs a windowed policy hands it. Five properties push it that way, and the revision addresses each. (i) Its few-shot balance was three related to one unrelated; class balance in the exemplars is a strong prior on the answer distribution, and this one pointed the wrong way by an order of magnitude. The revision is two related to three unrelated. (ii) Its single negative exemplar had the response *explicitly disavow* a link (“the two processes were not related this year”), which teaches that NONE requires denial, whereas real prose simply stays silent; the revision’s negatives are all silence. (iii) It asked for grounding but gave the model nowhere to show it and nothing to check it — hence the evidence span. (iv) Its first reasoning step was a leading question listing eight affirmative relation types before mentioning NONE; the revision asks the discriminating question first, whether the text draws the link or the claims are merely near each other. (v) It gave NONE one line against nine for relation-bearing senses; the revision gives it first position and equal prose. The revision also names *transitivity* explicitly as a non-relation, since roughly 85% of the first generation’s false positives are the model closing a chain the text laid out pairwise, and 96.6% of them touch an endpoint of a real relation. Both generations are retained permanently so that any claim about the prompt can be measured rather than asserted; §8 reports both.

Candidate pairs. Querying all $n(n - 1)$ ordered pairs is quadratic in LLM calls and, as argued above, actively harmful. We define four policies:

- ALL_PAIRS: every ordered pair $i \neq j$.
- WINDOWED: forward pairs with $0 < j - i \leq w$, default $w = 4$.
- GATED: the window, plus long-range forward pairs whose embedding or entity similarity exceeds a threshold, to catch discourse callbacks.
- BIDIRECTIONAL: *both* arc directions within $|j - i| \leq w$.

The last exists because of a structural limitation, not a tuning preference. A forward-only policy *cannot emit* an arc that runs backward in claim order, so a gold relation that is both directed and backward is unreachable under it no matter what the model says. In our corpus 105 of 603 edge-producing gold relations are of that kind, capping forward-only recall at 82.6% before the model is consulted; §8 measures the effect.

Every policy is then refined against the response itself (Algorithm 1). Each claim is aligned to its originating sentence; a pair is *discourse-adjacent* if the two claims share content tokens above a Jaccard threshold, if their sentences are within a short span, or if the target’s sentence opens with a discourse connective. Out-of-window pairs that are discourse-adjacent are *promoted* into the candidate set; in-window pairs that are not are *demoted* out of it. The refinement is a correlated filter — it selects for the same surface adjacency that the sense prompt reads as evidence of relatedness — which is why we report it as a policy component and measure its effect rather than treating it as a neutral preprocessing step.

Failures must be counted, not absorbed. Mining issues one LLM call per candidate pair for Prompt A and one more for each pair that yields an edge, run concurrently under a rate limit. A throttled or failed call returns no parse, and a parser that reads “no parse” as “no relation” converts an infrastructure failure into a sparse graph — which then scores as a *more* coherent response, since fewer conflicts are asserted. We therefore tally call errors separately from genuine NONE answers and report them, and distinguish an explicit NONE from an unparseable coupling for the same reason.

5 FOUR COHERENCE READOUTS

Definition 1 specifies a distribution; a coherence score is a functional of it, and several are defensible. We state five criteria, define four candidates with their exact constructions, and then decide between them by exact inference rather than by preference.

Definition 4 (Criteria). *A coherence readout should be: (R1) a genuine functional of the joint $\mathbb{P}(\mathbf{a})$, not of the relation list alone; (R2) valued in $[0, 1]$ and interpretable; (R3) monotone — resolving or removing a conflict must not decrease it; (R4) sensitive to conflict without hand-tuned constants; (R5) computable from standard inference queries.*

R1 deserves comment, since it is the criterion that does the most work. A readout that depends only on $\mathcal{R}$ — how many relations there are, what types, what weights — is not measuring the response; it is measuring its own annotation. Proposition 2 below shows this is not a hypothetical concern.

5.1 (A) MEAN POSTERIOR MARGINAL

$$\text{LCS}_{\text{mm}}(y) = \frac{1}{n} \sum_{i=1}^n \mathbb{P}(a_i = 1). \quad (5)$$

The marginals come directly from a single MAR query. The interpretation is the expected fraction of the response’s claims that the argument, taken as a whole, upholds. A conflict suppresses the marginals of its endpoints; a satisfied entailment chain propagates support along itself; a claim no relation touches stays at its prior and contributes neutrally.

Alongside the scalar we report three diagnostics that come free with the same query: the per-claim marginals themselves, which localize an incoherence to specific claims; the count of claims dragged *below their own prior*, which is the sharpest single indicator of an active conflict; and the mean normalized binary entropy $\mathcal{E} = \frac{1}{n} \sum_i H_2(q_i)$, lower being more decisive.

5.2 (B) CONSISTENCY: A TWO-TERM PROBABILITY

The obvious alternative to averaging marginals is to ask for the probability that the response is consistent. Making that precise turns out to require two terms:

$$\text{LCS}_{\text{cn}} = \frac{1}{2}(\mathcal{K} + \mathcal{S}), \quad \mathcal{K} = \mathbb{P}\left(\bigwedge_{\tau: \tau = \text{CONTRADICTION}} \neg(a_{s_r} \wedge a_{t_r})\right), \quad \mathcal{S} = \frac{\sum_{r \in \mathcal{R}_S} p_r U_r}{\sum_{r \in \mathcal{R}_S} p_r}, \quad (6)$$

where $\mathcal{R}_S$ ranges over $\{\text{ENTAILMENT}, \text{EQUIVALENCE}, \text{EXCLUSIVE}\}$ and U_r is the probability that relation r is *actively upheld*:

$$U_r = \mathbb{P}(a_{s_r}=1 \wedge a_{t_r}=1) \text{ for ENTAILMENT, EQUIVALENCE,} \quad U_r = \mathbb{P}(a_{s_r} \neq a_{t_r}) \text{ for EXCLUSIVE.}$$

The two coupling sets are deliberately different. The conflict event $\mathcal{K}$ keys on CONTRADICTION alone, even though $\mathcal{T}_{\text{conf}} = \{\text{CONTRADICTION}, \text{EXCLUSIVE}\}$. The reason is that an EXCLUSIVE spreads its incoherence over *both* same-value worlds, so reading only its both-true half would describe half the coupling and silently report the other half as satisfactory. EXCLUSIVE therefore earns credit in $\mathcal{S}$ instead, where its exactly-one world is simultaneously where the exclusion is honoured and where it is informative. CO_NECESSITY enters neither term: its defect is the both-false world, which (a)’s marginals already see directly.

Computing $\mathcal{K}$ exactly. $\mathcal{K}$ is an event probability over a conjunction of $|\{r : \tau_r = \text{CONTRADICTION}\}|$ pairwise events, which a naive encoding would represent as a factor of exponential arity. Instead we augment the network with deterministic auxiliary variables: one u_r per conflict edge with a ternary factor enforcing $u_r = a_{s_r} \wedge a_{t_r}$, then a chained running disjunction $U_k = U_{k-1} \vee u_k$ with ternary factors, and finally $\mathcal{K} = \mathbb{P}(U_{|\mathcal{R}|}=0)$. Every auxiliary factor is deterministic and at most ternary, so no $2^{|\mathcal{R}|}$ table is ever built and the augmented network is read with one MAR.

Computing $\mathcal{S}$ exactly. $\mathcal{S}$ is not an event probability but a weighted sum of *pairwise joints* — a linear functional of the joint — so it cannot be read off the same accumulator chain. It needs one deterministic indicator variable per supported edge and its own MAR. Because the indicators are mutually independent given the base network, they may be batched across several queries without approximation, which matters only for inference backends that enumerate.

Why “actively upheld” and not “satisfied”. The definition of U_r credits a relation only in the informative world, not in every world where it is technically satisfied. An entailment $a_s \Rightarrow a_t$ is satisfied vacuously whenever $a_s = 0$, and crediting that is fatal:

Proposition 2 (The vacuous-truth trap). *Let $\mathcal{K}' = \mathbb{P}(\text{every relation in } \mathcal{R} \text{ is satisfied})$, the apparently natural extension of $\mathcal{K}$ to all couplings. Then a response with no mined relations satisfies $\mathcal{K}' = 1$, the maximum attainable value.*

Proof. $\mathcal{K}'$ is the probability of an intersection of $|\mathcal{R}|$ events; for $\mathcal{R} = \emptyset$ the intersection is the whole space. $\square$

The proposition is immediate, but its consequence is not, and it is the reason LCS_{cn} has the form it does. Measured on the family of §10.3, a disconnected bag of claims scores $\mathcal{K}' = 1.000$ while a response whose causal spine is fully satisfied scores 0.221 — an inversion, not a tie. Nor is the defect repaired by count normalization: the geometric mean over relations gives 0.883 for the base response against 0.882 for the coherent one, non-monotone *and* still 1.000 on the bag; the expected fraction of satisfied relations gives 0.857 against 0.852, likewise inverted; and normalizing against a prior ceiling leaves $[0, 1]$ altogether (8.82 against 5.46) while still running backwards. The common cause is that “are the asserted relations honoured?” is a property of $\mathcal{R}$ rather than of y — a violation of R1 — and it is maximized by asserting nothing. Crediting only the informative world removes the free lunch: a bag of claims upholds nothing and scores $\mathcal{S} = 0$.

The two terms are averaged rather than multiplied so that neither can annihilate the other. A live contradiction should not erase a well-supported spine, and a relation-free response should not inherit a perfect score from an empty conflict set.

5.3 (C) A REIFIED COHERENCE NODE

The most direct reading of “is this response coherent?” as a single probability is to introduce a variable for it. Add a binary R with Bernoulli prior ρ and, for each relation, a ternary *vote* factor

$$h_r(R, a_s, a_t) = \begin{cases} 1 & R = 0 \quad (\text{the branch is flat}) \\ 1 & R = 1, (a_s, a_t) \notin V(\tau_r) \\ 1 - p_r & R = 1, (a_s, a_t) \in V(\tau_r), \end{cases} \quad \text{LCS}_{\text{re}} = \mathbb{P}(R = 1). \quad (7)$$

In the $R = 1$ branch each violated relation charges $1 - p_r$, so R behaves as a noisy-AND over relation satisfaction; in the $R = 0$ branch nothing is charged, so $R = 0$ carries no commitment. With no relations R is decoupled and $\text{LCS}_{\text{re}} = \rho$, correctly signalling that nothing was measured. Note that the violated set is $V(\tau)$ from Def. 2, so ENTAILMENT is treated as satisfied at $a_s = 0$ — vacuously — which is where (c) inherits a weaker form of Prop. 2’s problem. Appendix B works a three-claim subnet through by hand, giving $\mathbb{P}(R=1) = 0.453$.

5.4 (D) NORMALIZED LOG-PARTITION

The partition function measures how much probability mass the constraints leave standing, which is a global coherence signal requiring no marginals at all. It is unnormalized, so we grade it between two

Table 6: All four readouts across a coherence-ordered family of five variants of one response, by exact inference over 2^{16} worlds. The variants run from least to most coherent left to right in the intended ordering. Bold marks a value that moves the *wrong* way relative to its predecessor. Only (a) and (d) are monotone throughout.

Variant	(a) LCS_{mm}	(b) LCS_{cn}	(c) LCS_{re}	$\log Z$	(d) LCS_{lp}
<i>worse</i>	0.586	0.778	0.116	-10.54	0.68
<i>base</i>	0.607	0.781	0.148	-9.64	0.78
<i>concession</i>	0.612	0.772	0.134	-9.64	0.80
<i>fixcasualty</i>	0.613	0.764	0.156	-8.95	0.89
<i>coherent</i>	0.620	0.752	0.181	-8.25	1.00
monotone?	yes	no	no	yes	yes

Table 7: The four readouts against the criteria of Def. 4, with the inference queries each requires.

Readout	Queries	R1	R2	R3	R4	R5	Note
(a) mean marginal	MAR	✓	✓	✓	✓	✓	selected
(b) consistency	$\text{MAR} \times 2$	✓	✓	×	✓	✓	concession dip
(c) reified	MAR	✓	✓	×	×	✓	free ρ ; dip
(d) norm. $\log Z$	PR, MAP	✓	✓	✓	✓	✓	needs 2 references

references built from the same edge skeleton:

$$\text{LCS}_{\text{lp}} = \frac{\log Z - \log Z_{\min}}{\log Z_{\max} - \log Z_{\min}}. \quad (8)$$

For the ceiling, $Z_{\max}$ is the partition function of the same network with every conflict edge removed: deleting constraint factors only adds mass, so $\log Z \leq \log Z_{\max}$ always. For the floor we take the MAP world mass of the base network, $Z_{\min} = \max_{\mathbf{a}} \prod_i \phi_i \prod_r \psi_r$.

Lemma 1 (The MAP world mass is a valid, coherence-graded floor). $\log Z_{\min} \leq \log Z$ for every network, and $Z_{\min}$ is non-increasing in conflict strength.

Proof. $Z = \sum_{\mathbf{a}} \text{mass}(\mathbf{a})$ is a sum of non-negative terms, so its single largest term is a lower bound. For the second claim, increasing p_r on a conflict edge multiplies the mass of every world in $V(\tau_r)$ by a factor below one and leaves other worlds unchanged, so the maximum over worlds cannot increase. $\square$

Lemma 1 is not a technicality; it replaces two floors that look more natural and are wrong. Retyping every edge to CONTRADICTION, or saturating conflicts to $p = 1$, both fail to bound $\log Z$ below for the row-stochastic tables of Prop. 1: an isolated conflict removes no mass at all when its endpoints sit at the uniform prior — the setting in which both candidates were checked — and little away from it, whereas a mass-concentrating entailment backbone concentrates a great deal, so a network combining a strong backbone with several soft conflicts can have a base $\log Z$ below either candidate floor. We observed exactly this on real graphs, which put Eq. 8 outside $[0, 1]$ until the MAP floor replaced them. In practice we clamp the ratio to $[0, 1]$ regardless, since approximate inference computes $\log Z$ and the MAP bound with finite i-bound and small tolerance violations are possible.

5.5 DECIDING BETWEEN THE FOUR

Table 6 settles the question empirically and Table 7 summarizes. Readouts (b) and (c) both fail R3, and they fail it at the *same* variant — the one where a conceded tension is resolved — for a shared reason that is worth stating precisely, because it looks like an implementation bug and is not.

Proposition 3 (Belief versus event activity). *Softening a resolved conflict via Eq. 3 increases the probability of its both-true world. For a single conflict edge with uniform priors, discounting p from 0.80 to 0.55 moves $\mathbb{P}(a_s=1, a_t=1)$ from 0.10 to 0.225.*

Algorithm 2 COMPUTELCS(y) — the end-to-end coherence score.

Require: response y ; relation estimator $\mathcal{M}$; discount λ ; priors π
Ensure: $\text{LCS} \in [0, 1]$, marginals $\{q_i\}$, diagnostics

- 1: $\mathcal{A} \leftarrow \text{ATOMIZEANDDECONTEXTUALIZE}(y)$
- 2: $P \leftarrow \text{SELECTCANDIDATEPAIRS}(\mathcal{A}, y, \dots)$ ▷ Algorithm 1
- 3: $\mathcal{R} \leftarrow \emptyset$
- 4: **for all** $(a_i, a_j) \in P$ **do** ▷ concurrent, rate-limited
- 5: $(\text{sense}, \tau, c) \leftarrow \text{PROMPTA}(a_i, a_j, y)$ ▷ $c = \mathbb{P}(\tau \mid \cdot)$ from Eq. 1
- 6: **if** $\tau = \text{NONE}$ **then continue**
- 7: $\sigma \leftarrow \text{PROMPTB}(a_i, a_j, \tau, y)$ ▷ surrogate-token strength
- 8: $\mathcal{R} \leftarrow \mathcal{R} \cup \{(i, j, \tau, \text{clip}_{[0,1]}(c \cdot \sigma))\}$
- 9: **for all** $r \in \mathcal{R}$ with $\tau_r \in \mathcal{T}_{\text{conf}}$ sensed as **Concession** **do**
- 10: **if** a resolving holding h exists **then** $p_r \leftarrow p_r(1 - \lambda\pi_h)$
- 11: **else** mark r unresolved
- 12: $\mathcal{N} \leftarrow$ network of Eq. 2 over $(\mathcal{A}, \mathcal{R}, \pi)$ ▷ with-priors tables
- 13: $\{q_i\} \leftarrow \text{INFER}(\mathcal{N}, \text{MAR})$
- 14: **return** $\frac{1}{n} \sum_i q_i, \{q_i\}, (\#\{q_i < \pi_i\}, \mathcal{E}, \log Z)$

Proof. Under both conflict couplings the $(1, 1)$ cell of Table 2 is $1 - p$, which increases as p falls. Take a lone CONTRADICTION edge with $\pi_s = \pi_t = \frac{1}{2}$, so the with-priors table is $[\frac{1}{2}, \frac{1}{2}, p, 1 - p]$ and each world carries an equal unary factor of $\frac{1}{4}$. The unnormalized masses are then proportional to the table itself, giving $\mathbb{P}(1, 1) = (1 - p) / (\frac{1}{2} + \frac{1}{2} + p + (1 - p)) = (1 - p)/2$, which is monotone decreasing in p . Evaluating at $p = 0.80$ and $p = 0.55$ gives 0.100 and 0.225. $\square$

The consequence sorts the four readouts into two kinds. A readout that asks *is the conflict active* — (b) counts exactly that event, (c) charges $1 - p_r$ for it — must register a softened conflict as *more* often active, and therefore dips. A readout that asks *how strongly should each claim be believed* — (a) through marginals, (d) through mass — sees less suppression and rises. Since resolving a concession is by construction an increase in coherence, a coherence score must track belief rather than event activity. That is the substantive finding of this section, and it is not specific to our parameterization: any readout defined as the probability of a conflict-free configuration will invert under any discount-based treatment of resolved concessions.

A second failure mode of (b). Table 6 shows (b) failing worse than a single dip, and the reason is instructive about $\mathcal{K}$ and $\mathcal{S}$ jointly. In this family both conflicts are EXCLUSIVE rather than CONTRADICTION, so $\mathcal{K} \equiv 1.000$ throughout by the coupling-set decision above, and the score is carried entirely by $\mathcal{S}$. But removing an exclusion removes *support* credit as well as conflict, so the score runs backwards from 0.778 to 0.752 across the whole ladder. On three-coupling variants of the same response, where the conflicts are CONTRADICTION and $\mathcal{K}$ is live, (b) instead runs $0.576 \rightarrow 0.654 \rightarrow \mathbf{0.582} \rightarrow 0.672 \rightarrow 0.752$: correctly ordered apart from the concession dip. Either way it is non-monotone, and the lesson is that $\mathcal{K}$ and $\mathcal{S}$ should be read separately as diagnostics rather than averaged into a headline.

Selection. We report (a) as the headline. It satisfies all five criteria, needs a single MAR query, and is the only one of the four whose value decomposes into per-claim contributions — which is what makes an incoherence *locatable* rather than merely detectable. We compute and report all four, since (d) is a genuine alternative satisfying the same criteria through a different route (mass rather than marginals), and $\mathcal{K}, \mathcal{S}$ carry diagnostic information the mean marginal does not expose.

5.6 ALGORITHM AND COST

Algorithm 2 is dominated by steps 4–9: $|P|$ Prompt-A calls plus one Prompt-B call per surviving edge, which the candidate policy keeps near-linear in n rather than quadratic. Inference is negligible by comparison.

Answering all four readouts requires seven inference calls: the base MAR and PR, the $\mathcal{K}$ accumulator MAR, the $\mathcal{S}$ indicator MAR, the reified MAR, the conflict-free ceiling PR, and the base MAP. Computed independently they would cost fourteen, since each readout re-runs the shared base pair. Production inference is weighted mini-bucket (Liu & Ihler, 2011) with i-bound 6; every *exact* value reported in this paper instead comes from explicit 2^n enumeration, which we use as a regression oracle for $n \leq 20$ and which is the source of Tables 6 and the values in §10.

6 COMPOSING COHERENCE WITH FACTUALITY

The priors π of Definition 1 are a free slot, and the natural thing to put in them is what Stage One already computed (Figure 1):

$$\pi_i \longleftarrow q_i = \mathbb{P}(a_i = 1 \mid C). \quad (9)$$

A claim is then credited when it is both externally supported and internally upheld, and the two readings compose in the natural direction: an unsupported claim enters the coherence graph already doubted, so an argument resting on it is penalized twice — once for the evidence and once for what the argument does with it. Setting $q_i \equiv \frac{1}{2}$ recovers the pure coherence score of Remark 3, so the composition is an option rather than a commitment and the two axes remain separately reportable.

Two stages rather than one joint network. One could instead build a single MRF over claim and evidence variables with claim–claim edges added. We do not, for three reasons. Inference runs over n variables rather than $n(1+k)$ for k retrieved passages per claim, which matters because the augmented networks of §5.2 add further variables on top. The mined relation graph can be rescored under several prior sets — uniform, factuality-derived, ablated — without re-running the estimator, which makes prior ablations essentially free. And the factuality stage stays independently cacheable and independently reportable, which a fused model forfeits. Priors are applied at scoring time and never by mutating the mined graph, so one mined graph supports every prior condition.

Aligning claims across stages. The two stages must agree on which claim is which, and the obvious key is the wrong one. Matching proceeds by object identity first (the shared-atomization path, where the response is decomposed once and both stages see the same objects), then by normalized text (casefolded, whitespace-collapsed, terminal punctuation stripped), and only then by identifier. Text precedes identifier deliberately: near-duplicate removal keeps the first-seen claim and leaves the surviving identifier set *sparse* — $a_0, a_1, a_3, \dots$ — so two independent decompositions of the same response can disagree about which claim occupies which index. Matching on identifiers alone would then attach a prior to the *wrong* claim, which is strictly worse than attaching none, since it silently corrupts the score rather than degrading it. When coverage falls below one half the run is flagged degraded, and the configured policy is to warn, to raise, or to discard every prior and fall back to a clean coherence-only score rather than report a half-primed mixture.

7 EVALUATION PROTOCOL

Validating a coherence measure requires responses that differ in coherence while holding content fixed. Natural corpora do not supply this: two documents that differ in coherence almost always differ in topic, length and claim set as well, so any score difference is unattributable. Human coherence ratings supply a target but not a controlled contrast. And no existing resource supplies what separates *scoring* error from *estimation* error — a gold relation graph, against which the estimator of §4.5 can be measured directly.

We therefore evaluate on synthesized *coherence ladders*. A ladder is a family of five responses over one claim set: a base response realizing a declared relation plan, and four variants produced by applying perturbation operators to that plan — adding a spurious relation, resolving a conceded tension, dropping conflicts — each variant carrying *its own* gold relation graph derived from the base graph by the same operators. Each family declares the ordering its rungs must exhibit, in three classes: **C1**, a strict increase between consecutive rungs; **C2**, a *predicted* decrease or invariance — which is where Prop. 3’s concession dip becomes a falsifiable prediction rather than an excuse, and where an ordering-only edit must be exactly flat by Remark 2; and **C3**, endpoint separation between the extreme rungs. A measure is scored on the fraction of declared constraints it satisfies.

Arms. An *arm* fixes where the relation graph comes from while holding the claims, the priors and the readouts identical. The GOLD arm takes each item’s own labelled relations; GOLD_VALID keeps only those labelled valid; a MINED arm runs the estimator of §4.5 on the response text. This is the central experimental device of the paper. Because the readouts are byte-identical across arms, any shortfall of a mined arm below the gold arm is attributable to relation estimation and not to the probabilistic model, and the two error sources never have to be disentangled after the fact.

Match levels. Mined relations are compared to gold at three nested levels: **pair** (were these two claims related at all), **coupling** (pair, plus the Level-1 coupling, with direction required for the asymmetric couplings), and **sense** (plus the Level-2 sense). The coupling level is the one that matters, since it is the level at which Table 2 selects a factor, and it is materially stricter than pair level.

7.1 BASELINES

A coherence measure should be compared against controls chosen to fail in informative ways. We describe five families and state what each is diagnostic of. Four are implemented and share the LCS’s own decomposition — each baseline is handed the same atoms the LCS was scored on and, like the coherence MRF, consults no retrieved evidence — so the comparison is an ablation rather than a separate experiment. Metrics requiring a source document or a reference are therefore excluded by construction, and we say so where the exclusion is not obvious.

1. *Ladder-blind controls*: claim count, relation count, response length, and number of perturbation operators applied. These cost nothing and they catch the most likely way a coherence measure is secretly wrong — by tracking edit count or graph density rather than coherence. A measure that cannot beat them on C3, or that fails to stay flat where they do, has not earned its complexity. Steen & Markert (2022) make the same argument for coherence measures generally, and introduce intra-system correlation and bias matrices to identify measures whose apparent performance comes from system-level confounders; we adopt the controls in that spirit. A control that *succeeds* is bad news for the item, not good news for the control, and §7.2 reports one such case.
2. *Local contradiction detection*: pairwise NLI over claim pairs with no global model, scored as one minus the fraction of pairs labelled contradictory, in the spirit of SelfCheckGPT’s NLI variant (Manakul et al., 2023). This isolates precisely what joint inference adds over local detection. The prediction is that it does creditably where the defect is a *present* contradiction and poorly where the defect is a *missing* entailment, since a broken chain has no contradictory pair to find. We run it with the same NLI extractor as §3, over the same atoms, so the only differences from the LCS are relation *typing* and *propagation*.
3. *Structured reasoning consistency*: ROSCOE’s Self-Consistency (Golovneva et al., 2023), $1 - \max_i \max_{j < i} p_{\text{contr}}(h_i, h_j)$, which is evidence-free like ours but differs in three separable ways — it is a *maximum*, so a second conflict changes nothing; it is *forward-only*, so a relation running backward in claim order is unreachable (§4.5 measures 105 of 603 gold relations as both directed and backward); and it is *untyped*. We therefore also report a mean-aggregated and a symmetric variant, which turn one baseline into an ablation identifying *which* of the three differences accounts for any gap. ROSCOE targets step-by-step reasoning chains, and a long-form response is not a chain, so we label it adapted.
4. *LLM judges*: G-Eval’s coherence dimension (Liu et al., 2023), with the SummEval coherence rubric (Fabbri et al., 2021), and a direct “rate the logical coherence from 1 to 5” prompt as the no-rubric control. These are the practical incumbent and the comparison reviewers will want. The ladder is a demanding test for them, since consecutive rungs differ by a few words, so a judge must be sensitive at very small edit distances while remaining flat where meaning is preserved. We report each judge as a mean over five runs with its standard deviation, because the determinism of Remark 2 is only an advantage next to a measurement of the judge’s instability, and because a spread wider than the gap being discussed bounds what the comparison can conclude.
5. *Discourse representations*: DiscoScore (Zhao et al., 2023), with an entity-grid model (Barzilay & Lapata, 2008) as the classical reference. DiscoScore is presented as reference-based and most of its variants consume a reference, which a standalone response does not have; three of them (RC, LC and ENTITYGRAPH) ignore the reference argument entirely and score a single document, and those are the three we use — the last of which *is* an entity grid, so the classical

reference comes from the same implementation. All three reduce to noun-lemma repetition across sentences, with no representation of entailment or contradiction, which makes them a sharp test of whether coherence is separable from *cohesion* rather than a weak competitor: §7.2 reports a pair of responses with identical nouns and identical repetition structure, one self-consistent and one self-contradictory, on which all three are exactly flat.

6. *Internal ablations*: uniform versus factuality priors (Eq. 9); the four readouts of §5; the four candidate-pair policies; the two Prompt-A generations; and the gold-relation ceiling. These separate the model’s design choices from the estimator’s accuracy.

Beyond ladder ordering we recommend Spearman correlation against human coherence ratings on the same items, reported alongside per-family ordering agreement. Ordering agreement is the sharper instrument, being a controlled contrast with a known answer; human correlation is what establishes that the construct being ordered is the one readers care about. Neither substitutes for the other.

A baseline we cite but do not run. The closest existing work to our relation-mining stage is Liu & Strube (2025), who also build a graph over sentence pairs labelled with PDTB 3.0 senses — via a discourse parser rather than an LLM — and add entity links. We do not report a head-to-head number, and the reason is a difference in output type rather than a convenience: their model emits a *three-way* label (low, medium, high), and a three-way label cannot be checked against the ordering constraints of §7, which ask whether consecutive rungs move in a declared direction. The comparison would also require their trained classifier and a PDTB parser, neither of which is a pretrained scorer one can apply to a new response. Three differences are worth recording regardless, because they are the axes on which the two approaches diverge: their model performs no probabilistic inference, so a local conflict cannot propagate; it retains only forward-order pairs ($i < j$), so a backward relation is unreachable, which is the same structural cap we measure in §4.5; and it produces no per-claim attribution, so an incoherence is detected but not located.

7.2 WHAT THE BASELINES SHOW

Two properties of these baselines are worth establishing before any ladder sweep, because both bear on what the comparison in §8.5 can support.

First, *cohesion is not coherence, and the discourse floor cannot tell the difference*. On a pair of two-sentence responses with identical nouns and identical repetition structure — “The weld passed the tensile test. The weld held under load.” against “... The weld *failed* under load.” — all three single-document DiscoScore metrics return *exactly* equal values, while the pairwise-NLI and ROSCOE baselines separate the pair. The metrics are defined over noun repetition, so this is a property of their definitions rather than an accuracy failure; it is also the cleanest available demonstration that the axis we measure is not the one entity-based coherence models measure.

Second, *a controlled contrast has to be checked, not assumed*. A measure can reproduce a declared ordering for reasons that have nothing to do with coherence, so a contrast is only informative once the model-free proxies have been shown *not* to reproduce it. The claim-identical ordering pair of §10.4 is the clean case: claim count is exactly flat there by construction, response length differs by four tokens (265 against 269) in the direction *opposite* to the declared ordering, and all three discourse metrics rank the adversarial summary *above* the faithful one, the self-serving-first ordering carrying more noun repetition. A cohesion metric prefers the less coherent response of the two. The ladder corpus of §8 generalizes exactly this property: every rung of a family shares one claim set, so a claim-count proxy is flat by construction across all seventy items.

8 EXPERIMENTS

8.1 THE CORPUS

Evaluation is on a single body of data: a generated corpus of coherence ladders. §7 defines what a ladder, an arm and a match level are; this subsection gives the concrete instance.

The corpus holds **fourteen families and seventy items** — fourteen ladders of five rungs, each rung a distinct response over one claim set. It was produced by `claude-opus-5` and admitted

by a three-model cross-vendor validation committee (gpt-5.6-terra, gemini-3.1-pro, claude-sonnet-5), which recovers the relation graph from prose without seeing the plan (V1, any-of), audits the prose for leakage and hedging (V3, majority) and checks atom coverage (V4, majority). Items carry 15–16 claims and 609–698 words, for 1,115 atoms and 670 gold relations in total at a 0.61 valid fraction, together with 342 claim pairs the corpus *explicitly declares unrelated* — a labelled negative set rather than an unlabelled remainder. Subject matter spans fourteen canonical topics across all four domains, so no single topic or domain carries a result.

All four ladder types are represented, and the mix matters for reading the results: eight CONFLICT families (rungs resolve or remove conflicting relations), two CHAIN (entailment links are broken and repaired), three ORDER and one CONTROL. The last two are the negative controls. Every rung of an ORDER or CONTROL family applies only meaning-preserving edits — sentence reordering, or a PRECEDENCE $\leftrightarrow$ SUCCESSION swap that couples no truth values — so those ladders assert *invariance*, and a measure that trends across them is tracking surface form rather than coherence. CONFLICT and CHAIN assert a monotone increase. A corpus with only the increase-type ladders cannot distinguish a coherence measure from an edit-count proxy, which is why both kinds are present.

8.2 PROTOCOL

Each family declares the ordering its rungs must exhibit, per readout, and those declarations are the unit of measurement: 202 **constraint checks** in total, decomposing as 112 over CONFLICT, 30 over CHAIN, 45 over ORDER and 15 over CONTROL. A check names a readout and a rung pair and asserts increase or invariance; a missing score fails rather than passes.

An arm fixes where the relation graph comes from and holds everything else — the claims, the atom priors, the readouts, the scorer settings — byte-identical, so a difference between two arms is attributable to relation sourcing alone. The GOLD arm compiles each item’s own labels; GOLD_VALID drops the deliberately-planted invalid ones. The mined arms run the estimator of §4.5 over the response prose with two miners, llama-3.3-70b-instruct and gpt-oss-120b, both read through token logprobs, under two candidate-pair policies: WINDOWED (forward-only, within an order window) and BIDIRECTIONAL (both arc directions within the same radius).

Which relation estimator these arms used, and why it matters. The mined arms in this section read the relation type-confidence from the label token’s logprobs under the chain-of-thought prompt — the estimator §9.2 later shows to be *saturated*, returning ≈ 1 for essentially every pair. So in the mined arms below the relation *type* is informative and the relation *probability* is not, and the coupling strengths that enter Eq. 2 are near-deterministic. We report the arms as measured rather than re-mining them under the no-reasoning prompt, and the reader should read the mined percentages as a lower bound for that reason. Two considerations bound how much this matters. First, the gold arms are unaffected: they compile the corpus’s own labels and consume no estimator, so the gold-versus-mined gap — the section’s central claim — is not an artifact of the probability channel. Second, §9.4 finds the readouts’ behaviour is driven by relation *type* rather than by the probabilities attached to it, which predicts that repairing the channel would move these numbers modestly rather than closing the gap. That prediction is untested here and we flag it as the obvious next experiment.

Mining is sampled once per cell at the model’s default temperature, so the mined arms are not deterministic. We measure that rather than estimate it: two *independent* runs of the same two WINDOWED configurations moved every pair, coupling and sense F1 by at most 0.014, while the gold arms reproduced bit-for-bit across the same pair of runs. Per-arm ladder agreement is the noisier statistic, moving by four to six of 202 assertions between those runs, so a single arm’s percentage should not be read to the point. That 0.014 supersedes the ± 0.06 we previously estimated on a ten-item corpus: the larger corpus tightens reproducibility rather than merely enlarging the denominator. These are measurements, not significance claims. All exact values are by 2^{16} enumeration; all corpus-scale values are by weighted mini-bucket with i-bound 6.

8.3 RELATION ESTIMATION IS THE BOTTLENECK

Table 8 measures the estimator of §4.5 against gold, and three findings stand out.

Table 8: Mined-versus-gold agreement over the 603 edge-producing gold relations (670 total, less the ordering-only PRECEDENCE/SUCCESSION edges, which compile to Level-1 none and contribute no factor), micro-averaged over seventy items. **Fwd/Bwd** split coupling-level recall by whether the gold relation runs forward or backward in claim order; 105 of the 603 run backward. **Non-rel** counts violations of the 342 pairs the corpus explicitly declares unrelated — a real negative set, unlike the unlabelled remainder. **Dup** counts unordered pairs carrying two edges, which BIDIRECTIONAL emits by construction and which place two factors on one variable pair.

Miner	Policy	Edges	Pair R	Coup P	Coup R	Coup F1	Fwd	Bwd	Non-rel
llama-3.3-70b	WINDOWED	1106	82.3	34.1	62.5	44.1	71.1	21.9	34/342
llama-3.3-70b	BIDIRECTIONAL	3172	91.7	17.9	79.3	29.3	79.9	76.2	119/342
gpt-oss-120b	WINDOWED	1094	83.4	40.2	73.0	51.9	82.1	29.5	38/342
gpt-oss-120b	BIDIRECTIONAL	2973	96.5	22.4	88.4	35.8	89.0	85.7	147/342

First, the forward-only policy is *structurally* capped, exactly as predicted, and the cap is now measured on a corpus large enough to see it clearly: WINDOWED recovers 21.9–29.5% of backward-running gold relations against 76.2–85.7% for BIDIRECTIONAL, because 105 of the 603 scorable gold relations are both directed and backward and no forward-only candidate set can emit them. This is a property of the candidate policy, not of either model.

Second, admitting both arc directions buys that recall at a precision cost severe enough to invert the F1 ordering. BIDIRECTIONAL roughly triples the emitted edges (1,106 $\rightarrow$ 3,172 for llama), lifts coupling recall by 16.8 points, and *halves* coupling precision (34.1 $\rightarrow$ 17.9), so coupling F1 falls from 44.1 to 29.3. The mechanism is the one §4.5 anticipated for ALL_PAIRS and it appears here in milder form: because Eq. 2 is a product over relations, spurious factors drive marginals toward the prior rather than merely diluting a signal. The 1,231 duplicate unordered pairs are the same defect seen structurally — two factors on one variable pair, which the model has no principled way to combine — and declared-non-relation violations rise from 34/342 to 119/342 in step.

Third, the two miners differ in the direction that matters, and gpt-oss-120b is the stronger relation estimator at every level of the taxonomy: on WINDOWED it beats llama-3.3-70b by 7.8 points of coupling F1 (51.9 vs. 44.1), 9.3 points of sense F1, and 10.5 points of backward recall, while emitting slightly *fewer* edges. The gap widens as the label gets finer, which is what a genuine capability difference looks like rather than a threshold artefact: both models find roughly the same pairs (pair recall 82.3 vs. 83.4) and then disagree about what those pairs mean.

A caveat belongs with that third finding, because it cuts against the natural reading. The gpt-oss strength estimates are almost entirely saturated: 99% of its edge strengths sit at the numerical boundary (25% at the floor, 75% at the ceiling) against 58% genuinely interior values for llama-3.3-70b. Probed directly, the model answers the strength surrogate with all of its top-five mass on one side for easy pairs and with values as small as 10^{-9} for hard ones — so this is the model’s own near-determinism, not a defect of the reader. gpt-oss therefore buys its accuracy advantage while supplying almost no graded uncertainty, and llama is the reverse. A practitioner who needs calibrated edge probabilities rather than accurate edge labels should read this table in the opposite direction.

8.4 THE MODEL ORDERS CORRECTLY WHEN THE GRAPH IS CORRECT

On gold relations the readouts satisfy 82.7% of the declared ordering constraints — 167 of 202 (Table 9). Mined arms reach 34.7–37.1%. Since the readouts, the claims and the priors are byte-identical across arms, the 46–48 point gap is attributable to relation estimation alone, and it is consistent with Table 8: at 29–52% coupling F1 roughly half the factors in a mined network are wrong or missing, and a ladder whose consecutive rungs differ by one or two relations cannot survive that. This is the paper’s central empirical claim, and it is a claim about *where the remaining problem is*: the probabilistic model orders responses correctly when handed a correct graph, and the open problem is extracting one.

The per-ladder breakdown localizes the failure, and it is worth reading before the aggregate: the single percentage conceals a sharp qualitative split.

Table 9: Left: declared ordering constraints satisfied, grouped by ladder type — 202 assertions per arm in total (8×14 CONFLICT, 2×15 CHAIN, 3×15 ORDER, 1×15 CONTROL). Right: each readout averaged over the seventy items. ORDER and CONTROL assert *invariance*: every rung of those ladders applies only meaning-preserving edits (sentence reordering, or a PRECEDENCE $\leftrightarrow$ SUCCESSION swap), which move no factor, so a readout that trends there is tracking surface form rather than coherence. GOLD_VALID is near-flat on CONFLICT by construction — a conflict-fixing rung removes the planted-invalid conflicts this arm has already discarded — so its failures carry no information about the readouts; the GOLD arm is what the constraints test. Averaging a readout across the rungs of a ladder deliberately collapses the ordering the ladder exists to test, so the right-hand block is a level comparison only.

Arm	Policy	Constraints satisfied					Mean readout value			
		Cfl	Chn	Ord	Ctl	%	LCS _{mm}	LCS _{cn}	LCS _{re}	LCS _{ip}
GOLD	—	89/112	18/30	45/45	15/15	82.7	0.795	0.515	0.238	0.410
GOLD_VALID	—	37/112	12/30	45/45	15/15	54.0	0.794	0.576	0.316	0.485
<i>mined by llama-3.3-70b</i>										
	WINDOWED	53/112	17/30	0/45	0/15	34.7	0.796	0.676	0.384	0.224
	BIDIRECTIONAL	57/112	18/30	0/45	0/15	37.1	0.723	0.542	0.205	0.037
<i>mined by gpt-oss-120b</i>										
	WINDOWED	59/112	12/30	0/45	0/15	35.1	0.825	0.500	0.425	0.129
	BIDIRECTIONAL	53/112	17/30	0/45	0/15	34.7	0.808	0.462	0.251	0.033

On gold, ORDER and CONTROL are perfect (45/45 and 15/15). Every rung of those ladders applies only meaning-preserving edits, which move no factor, so exact invariance is the correct answer and the model gives it. CONFLICT is strong at 89/112 (79.5%). CHAIN is the weak gold ladder at 18/30 (60.0%): its rungs differ by removing entailment links from a chain whose remaining entailments already concentrate most of the mass, the smallest perturbation the corpus contains and the one most exposed to i-bound approximation.

Every mined arm scores 0/45 on ORDER and 0/15 on CONTROL. This is the most consequential number in the table and it is not a scoring artefact. Those ladders hold the claim set fixed and permute the prose; a relation graph read off the text should therefore be identical across their rungs. It is not. On one ORDER family the gpt-oss WINDOWED miner emitted 9, 16, 20, 20 and 17 edges across five rungs that differ *only* in sentence order, and LCS_{cn} swung by ± 0.49 between adjacent rungs where the required answer was exact invariance. So the extractor is strongly order-dependent: the same claims, reordered, yield a materially different graph. Human annotators are the obvious reference point here and we do not have them, so we state the weaker claim the data supports — LLM relation extraction is not reorder-invariant, and any coherence measure built on it inherits that sensitivity however sound its probabilistic layer.

The corollary is that the mined arms are *not* uniformly worse than gold. On the two ladders that demand a monotone trend they are respectable — CONFLICT 53–59 of 112 and CHAIN 12–18 of 30, the latter matching or exceeding gold’s 18/30 on llama’s BIDIRECTIONAL arm — and they lose the aggregate comparison almost entirely on the two invariance ladders, where they score nothing at all. A measure evaluated only on increase-type ladders would have reported mined extraction as far closer to gold than it is. That is precisely the false positive the CONTROL ladder was added to catch, and here it caught one.

Two secondary observations. Mined ladder agreement does *not* track coupling F1: llama’s BIDIRECTIONAL arm has the worst coupling F1 of the four (29.3) and the best ladder agreement (37.1%), while gpt-oss’s WINDOWED arm inverts that. Given the four-to-six assertion run-to-run movement measured in §8, the honest reading is that these four arms are not separated by ladder agreement at all, and candidate policies should not be ranked on it. And the mean-readout block shows the BIDIRECTIONAL density cost quantitatively: LCS_{ip} collapses to 0.037 and 0.033 against gold’s 0.410 on the same seventy responses, from adding roughly two thousand spurious factors — the log-partition readout registers added factor mass most directly.

Table 10: Declared ordering assertions satisfied on the ladder corpus, under two different judge models. Only the readout-independent content of a constraint is scorable by a single-score baseline: the four C1 consecutive-increase pairs and the C3 endpoint separation, giving five per family. C2 is excluded because it predicts the internal behaviour of particular readouts (§5.5), which a baseline makes no claim about. The ORDER and CONTROL families assert invariance rather than increase and so contribute *no* scorable assertion, leaving ten families and 50 assertions; the LCS rows are scored on *exactly* those 50, so they are comparable to the rows below them, while the 82.7% of Table 9 is computed over the fuller 202-assertion set including invariance and is not. Judges are the mean of five runs. Model-free baselines are identical under both judges, as they must be; the two judge rows are the entire difference.

Measure	Assertions satisfied				Uses LLM
	llama-3.3-70b		gpt-oss-120b		
LCS _{mm} (gold relations)	37/50	74.0			no
LCS _{lp} (gold relations)	36/50	72.0			no
LCS _{cn} (gold relations)	34/50	68.0			no
RC (cohesion)	30/50	60.0	30/50	60.0	no
LC (cohesion)	28/50	56.0	28/50	56.0	no
LENGTH (control)	20/50	40.0	20/50	40.0	no
ROSCOE-SC (mean)	16/50	32.0	17/50	34.0	yes
CONTRADICTION RATE	12/50	24.0	12/50	24.0	yes
JUDGE, direct rating	4/50	8.0	18/50	36.0	yes
JUDGE, G-Eval	0/50	0.0	21/50	42.0	yes
ROSCOE-SC (max)	0/50	0.0	0/50	0.0	yes
ENTITY GRAPH	0/50	0.0	0/50	0.0	no
CLAIM COUNT (control)	0/50	0.0	0/50	0.0	no

8.5 THE EVIDENCE-FREE BASELINES, MEASURED

The baselines do not all measure the same thing as the LCS, and the comparison is only readable once two constructs are separated. *Logical coherence* is a property of the relations among a response’s claims: whether what it asserts can jointly hold. *Discourse cohesion* is a property of its surface texture — entity repetition, lexical chaining, referential continuity. The two are correlated in ordinary prose and come apart under exactly the edits this corpus applies, which is why the cohesion metrics appear here as baselines rather than as competitors: §7.2 gives a two-sentence pair on which all three single-document DiscoScore metrics return *identical* values while one response contradicts itself and the other does not. A cohesion metric that scores well below is measuring a different axis, not measuring this one badly.

Two results follow.

The probabilistic model separates from cohesion counting. LCS_{mm} reaches 37/50 (74.0%) against the best evidence-free baseline, noun-repetition cohesion, at 30/50 (60.0%). *All three* graded readouts (74.0, 72.0, 68.0) sit above *every* evidence-free baseline, so the ordering is not an artefact of choosing a headline readout, and the margin is larger than the four-to-six assertion run-to-run movement of §8. This supersedes a tie between the two that we previously reported at this comparison’s smaller scale; we flag the direction of the correction rather than restating the number quietly.

An LLM judge’s score on this benchmark is a property of the judge, not of the method. Every model-free baseline is bit-identical across the two judge runs — RC 30/50 twice, LC 28/50 twice, LENGTH 20/50 twice — which is the internal check that the harness is deterministic where it should be. Against that fixed background the two judge rows move enormously: G-Eval scores 0/50 under llama-3.3-70b and 21/50 under gpt-oss-120b, and direct rating 4/50 against 18/50. Same prose, same rubric, same five-run averaging, same 50 assertions; only the judge model changed. A single-judge evaluation of this kind therefore reports the judge as much as the metric, and the 42-point spread is larger than the gap between the best baseline and the worst.

Table 11: Both axes at once, on the ladder corpus. **Increase** is the 50 readout-independent ordering assertions over the eight CONFLICT and two CHAIN families; **Invariance** is the 20 rung pairs over the three ORDER and one CONTROL families, satisfied when the score is unchanged to 10^{-6} . Baselines are shown under both judge models where the two differ; the model-free rows are judge-independent by construction. The three baselines at 20/20 invariance reach it *vacuously* — each is a constant over all seventy items, hence trivially invariant and trivially unable to order anything, which is why the final column takes the weaker of the two axes.

Measure	Increase (of 50)		Invariance (of 20)		min
LCS _{mm} (gold relations)	37	74	20	100	74
LCS _{lp} (gold relations)	36	72	20	100	72
LCS _{cn} (gold relations)	34	68	20	100	68
RC (cohesion)	30	60	16	80	60
LC (cohesion)	28	56	16	80	56
LENGTH (control)	20	40	12	60	40
ROSCOE-SC (mean)	16–17	32–34	5–9	25–45	32
CONTRADICTION RATE		12 24	8–9	40–45	24
JUDGE, G-Eval	0–21	0–42	5–13	25–65	0
JUDGE, direct rating	4–18	8–36	5–14	25–70	8
ROSCOE-SC (max)	0	0	20	100 [†]	0
ENTITY GRAPH	0	0	20	100 [†]	0
CLAIM COUNT (control)	0	0	20	100 [†]	0

[†] vacuous: the measure is constant across all seventy items.

The mechanism is saturation rather than noise, and the two judges fail in opposite ways. llama’s G-Eval is astonishingly stable and useless: mean per-item standard deviation 0.0002, only three distinct values across seventy items, and a mean score of 0.996 — it rates almost everything at ceiling, so it cannot order anything. gpt-oss’s G-Eval uses its range (fourteen distinct values, mean 0.749) and pays in stability (mean sd 0.0800, max 0.187). Reporting judge variance alone would have called the first judge the reliable one, which is exactly backwards for the task.

Finally, four baselines score 0/50 under both judges and it is worth separating the reasons, since they are not one finding. ROSCOE-SC (max) collapses to a constant 0.0 on all seventy items: a max over pairwise contradiction probabilities saturates as soon as one pair is confidently contradictory, which on prose of this length is every item. ENTITY GRAPH returns a constant 1.0. CLAIM COUNT is constant *by construction*: every rung of a family shares one claim set, which is what makes this corpus a controlled contrast and what makes a claim-count proxy vacuous on it. Only the first two are failures of the measure; the third is the corpus doing its job.

8.6 SUMMARY: WHAT THE LCS ADDS, AND WHAT THE BASELINES CANNOT DO

The corpus asks two questions of a coherence measure, and they are not the same question. On the CONFLICT and CHAIN families it must *move in the declared direction* as relations are repaired — 50 assertions readable by any single-score measure. On the ORDER and CONTROL families it must *not move at all*, because those rungs permute prose while holding the relation set fixed — 20 further pairs. A measure is only useful if it passes both: ordering without invariance is a length or edit-count proxy, and invariance without ordering is a constant.

What the LCS adds. Given a correct relation graph it is the only measure here that passes both axes non-vacuously: 37/50 on ordering and 20/20 on invariance, and the same holds for the other two graded readouts (Table 11). The invariance column is where the gap is categorical rather than incremental. The readouts are invariant under prose reordering *by construction*, not by fitting: a PRECEDENCE $\leftrightarrow$ SUCCESSION swap compiles to Level-1 none and contributes no factor, and shuffling sentences does not change which claims stand in which relation, so the Markov network is literally identical and the readout cannot move. No surface statistic can offer that guarantee, because every surface statistic is a function of the surface. This is the property the construction buys, and it is the reason a coherence measure should be defined over relations rather than over text.

Table 12: The increase axis of Table 10, with relations mined from prose instead of taken from the corpus. Scored on the same 50 readout-independent assertions (C1 consecutive increase plus C3 endpoint separation, over the eight CONFLICT and two CHAIN families), so every row here is comparable to every row of Table 10. Gold rows repeated for reference; the best baseline in that table is noun-repetition cohesion at 30/50 (60.0), shown here as the line a mined arm must clear.

Relations	Readout	Assertions satisfied	
GOLD	LCS _{mm}	37/50	74.0
	LCS _{lp}	36/50	72.0
	LCS _{cn}	34/50	68.0
<i>mined by llama-3.3-70b</i>			
WINDOWED	LCS _{mm}	20/50	40.0
	LCS _{lp}	35/50	70.0
	LCS _{cn}	23/50	46.0
BIDIRECTIONAL	LCS _{mm}	28/50	56.0
	LCS _{lp}	25/50	50.0
	LCS _{cn}	34/50	68.0
<i>mined by gpt-oss-120b</i>			
WINDOWED	LCS _{mm}	25/50	50.0
	LCS _{lp}	29/50	58.0
	LCS _{cn}	25/50	50.0
BIDIRECTIONAL	LCS _{mm}	21/50	42.0
	LCS _{lp}	28/50	56.0
	LCS _{cn}	24/50	48.0
RC (cohesion; best baseline, Table 10)		30/50	60.0

What the baselines fail to do, and it is three different failures. They are not weak in one way. *Cohesion metrics* (RC, LC) are the strongest baselines and get respectably far on ordering (60%, 56%) while being nearly flat on invariance (16/20, drifting by at most 0.0096 across a permuted family) — they fail not by being noisy but by measuring a different axis, so they cannot distinguish a response that contradicts itself from one that does not when both repeat the same nouns. *Contradiction-based metrics* (ROSCOE-SC, CONTRADICTION RATE) detect conflict soundly and then destroy the signal in aggregation: the max variant is a constant on prose of this length, and the mean variant divides by $\binom{n}{2}$, so both are dominated by response size rather than by relational structure. *LLM judges* are the least reliable of the three, and in a way that a single-judge evaluation conceals: G-Eval scores 0/50 under one judge and 21/50 under another on identical prose and rubric, and both judges violate invariance freely — gpt-oss’s G-Eval moves by up to 0.55 across rungs that differ only in sentence order. A judge asked for a coherence rating is answering a question about the text, so it inherits the text’s surface variation.

8.7 THE SAME TWO COMPARISONS, ON MINED RELATIONS

Tables 10 and 11 score the LCS on *gold* relations, which is the right way to ask whether the probabilistic model works but not whether the *pipeline* does. Tables 12 and 13 repeat both comparisons with the relation graphs mined from prose by the two miners of §8, on exactly the axes the baselines are held to, so a mined arm can be read against a cohesion metric directly. Everything else — the claims, the atom priors, the readouts, the assertion sets — is unchanged, and the gold rows are repeated in both tables as the reference. These numbers come from the same cached runs as Table 9; the estimator caveat of §8 applies to them equally.

On the increase axis a mined arm is competitive; on invariance it is not. Read Table 12 against the baselines and the mined arms are not the collapse that Table 9’s 34.7–37.1% might suggest — that figure is over the fuller 202-assertion set, which includes the invariance families the mined arms fail wholesale. Restricted to the increase axis the mined arms span 40–70%: two of the twelve mined rows clear the best evidence-free baseline’s 60.0% (llama/WINDOWED LCS_{lp} at 70.0 and llama/BIDIRECTIONAL LCS_{cn} at 68.0) and three more reach 56–58%, while the remaining seven fall below it. But Table 13 is the honest summary, and there every mined arm scores 0/20 on

Table 13: Both axes of Table 11, on mined relations. **Increase** is the 50 readout-independent assertions; **Invariance** is the 20 rung pairs of the three ORDER and one CONTROL families, satisfied when the score is unchanged to 10^{-6} . The final column takes the weaker axis, as in Table 11. Every mined arm scores 0/20 on invariance, and §8.7 shows this is real movement rather than absent scores.

Relations	Readout	Increase (of 50)		Invariance (of 20)		min
GOLD	LCS _{mm}	37	74	20	100	74
	LCS _{lp}	36	72	20	100	72
	LCS _{cn}	34	68	20	100	68
<i>mined by llama-3.3-70b</i>						
WINDOWED	LCS _{mm}	20	40	0	0	0
	LCS _{lp}	35	70	0	0	0
	LCS _{cn}	23	46	0	0	0
BIDIRECTIONAL	LCS _{mm}	28	56	0	0	0
	LCS _{lp}	25	50	0	0	0
	LCS _{cn}	34	68	0	0	0
<i>mined by gpt-oss-120b</i>						
WINDOWED	LCS _{mm}	25	50	0	0	0
	LCS _{lp}	29	58	0	0	0
	LCS _{cn}	25	50	0	0	0
BIDIRECTIONAL	LCS _{mm}	21	42	0	0	0
	LCS _{lp}	28	56	0	0	0
	LCS _{cn}	24	48	0	0	0

invariance and therefore 0 on the weaker-axis column. The gold arms' categorical advantage is entirely on that second axis.

Why the invariance failure is the substantive one. A 0/20 could in principle be an artefact of missing scores, since a missing score fails rather than passes. It is not: the readouts move, and by margins far above the 10^{-6} tolerance. On llama's WINDOWED arm the LCS_{mm} range across the five rungs of an invariance family is 0.084 on f003 (CONTROL), 0.165 on f004, 0.123 on f007 and 0.378 on f011 (all ORDER) — movement of up to a third of the scale on rungs that only permute sentences. The mechanism is the one §8.6 identifies: the readouts are reorder-invariant *given a fixed relation set*, because a PRECEDENCE $\leftrightarrow$ SUCCESSION swap compiles to no factor and shuffling sentences does not change which claims stand in which relation. Mining breaks the premise rather than the theorem — an order-window policy reads adjacency as evidence, so permuting the prose changes which pairs are even considered, and the network is no longer identical. Invariance is a property the model has and the current extractor destroys, which is a sharper localization of the open problem than the aggregate percentage gives.

A row not to over-read. llama/WINDOWED LCS_{lp} at 35/50 (70.0) sits just below the gold LCS_{lp} of 36/50, which invites the reading that mined relations nearly match gold on this readout. They do not. That arm's per-rung sequences are not monotone — on f001 the five rungs read 0.162, 0.169, 0.138, 0.125, 0.130 — so it satisfies the consecutive-increase and endpoint assertions the axis counts while wandering in between, and its mean value (0.224) is roughly half gold's (0.410). The C1/C3 subset is a coarse instrument by construction, which is exactly why it is paired with the invariance axis rather than reported alone.

Two honest qualifications. First, the LCS column above is the *gold-relation* arm. Handed graphs mined from prose the same readouts fall to 34.7–37.1% of the full 202 assertions and to 0/45 and 0/15 on the invariance ladders (§8.5, Table 9) — the guarantee is a property of the model given a correct graph, and mined graphs are not yet reorder-invariant. The contribution is therefore a measure that is right when its input is right, plus a benchmark that localizes the remaining problem in extraction rather than in scoring. Second, the two axes are not equally hard: ten of the thirteen rows clear 40% on invariance while only five clear 50% on ordering, and of those five, three are the LCS readouts and two are the cohesion metrics. Invariance alone is therefore not a discriminating test — it

is a necessary condition most baselines meet, and the three constant baselines meet trivially. The claim rests on passing both, which only the LCS does.

9 BEYOND LADDERS: THE TWO-STAGE MODEL ON FACTUALITY CORPORA

Everything up to here was measured on the ladder corpus, where every item is a controlled perturbation of a claim set and the ground truth is an *ordering* the rungs must exhibit. That design is what makes the relation graph auditable, but it buys auditability at the cost of realism: the responses are generated to a plan, and the claim sets are small and hand-checked. This section asks the complementary question. Run the whole two-stage pipeline — retrieval-grounded factuality (§3) feeding the coherence MRF (§4) — over *naturally occurring* long-form generations, with no planted relations and no declared ordering, and report what each measure actually does there.

There is no ordering ground truth on such a corpus, so the question is not “does the measure move the right way”. It is the prior question, which the ladder design cannot ask: **does the measure distinguish these responses from one another at all?** A coherence measure that returns the same number for every response has no ordering to be right or wrong about. We therefore report, for every measure, the number of *distinct* values it takes across the corpus alongside its mean — and it is this column, not the mean, that separates the measures.

9.1 THE CORPORA

Five long-form generation datasets, each already decomposed into atomic claims with retrieved evidence passages attached per claim: BIO (biographies of rare entities, verified against Wikipedia), ASKHIST and ELI5 (long-form question answering), BOOKS (literary questions) and LFOBJ (long-form objective questions), the last four retrieved from web search. Together they hold 1,137 responses and 28,646 claims.

We evaluate on a **seeded stratified sample of 250 responses and 6,266 claims** — fifty from each dataset, so no dataset dominates by size (they range from 157 to 380 items), drawn by a seed recorded with the sample so it regrows exactly. The sample is a cost decision and we state it plainly: the full corpus prices at roughly 131K relation calls for stage one plus 619K mining calls for stage two *per model*, and the experiment runs two models under two pair policies. Items with fewer than five claims are excluded, having too few pairs to form a coherence graph at all.

BIO carries one thing the others do not: human SUPPORTED/NOT-SUPPORTED labels on every claim. Those labels are never shown to either stage; they are held out as an external check on stage one, reported below.

9.2 PROTOCOL

Nothing is atomized and nothing is retrieved: the corpora already carry both, so both stages read the same fixed claim set and the same fixed passages, and the only variables are the estimator and the pair policy. Stage one is FactReasoner v2 over claims and passages, scoring claim–passage relations under the provenance candidate-pair preset, which resolves the same graph as scoring every pair while issuing roughly $2.8\times$ fewer calls. Stage two is the pairwise coherence MRF of §4 with stage one’s posterior marginals as the claims’ unary priors, read out four ways.

One protocol point is load-bearing for the comparison. **Stage one runs once per response, not once per pair policy.** The marginals do not depend on how candidate claim pairs are selected, so they are computed once and both policies are scored against the identical prior vector. The two policies are therefore an exact ablation: any difference between the WINDOWED column and the BIDIRECTIONAL column is attributable to pair selection alone, not to a re-run of stage one. The baselines are handed the same claim texts, and never atomize for themselves, for the same reason.

Relation probabilities, and a prompt that had to change. The relation estimator of §4.5 reads its confidence from the label token’s logprob. Doing that on these models exposed a defect in how that probability was being obtained, and it is worth reporting because it is invisible to any evaluation that reads only labels. The original prompt asks the model to reason and *then* emit the label. By the time the label token is generated the model has already argued the verdict in prose, so the token is all but

determined by its own preceding reasoning: measured over real claim–passage pairs, the estimate was ≥ 0.999 on 30 of 30 pairs, taking a single distinct value to four decimal places. The channel carried no information, and every relation entered the network as a near-deterministic constraint.

Two candidate repairs do not work, and both are instructive. Raising the sampling temperature changes nothing, because the serving stack reports logprobs computed before temperature scaling — $T = 0$ through $T = 2$ moved one borderline pair’s estimate from 0.9941 to 0.9941. And the apparent spread of the label token’s raw probability is an artifact: renormalizing its mass over the three admissible labels returns exactly 1.0, because the competing mass sits on *casing variants of the same label* rather than on rival labels.

What does work is removing the reasoning. A few-shot prompt that maps premise and hypothesis directly to a bracketed label, with no intermediate reasoning, makes the label the model’s first commitment, and its probability a posterior over the three classes; we renormalize over those classes so that surface-form mass cannot read as uncertainty. Labels are unchanged from the reasoning prompt on the same pairs, and output shrinks from roughly 200 tokens to six, so the informative estimator is also the cheaper one. On deliberately borderline pairs it spreads 0.59–1.00 (standard deviation 0.13), e.g. 0.59 where a premise says “most regions” and the hypothesis says “every region”. All results in this section use this estimator.

The repair is model-dependent, which bounds the claim. `gpt-oss-120b` does not comply with the instruction not to reason: it emits an internal analysis channel regardless, so its label remains post-hoc and its probability saturated even on the borderline pairs. For that model the relation *type* is informative and the relation *probability* is not. We report both models under the same estimator rather than quietly giving them different ones, so this difference is visible in the tables instead of hidden in the configuration.

9.3 STAGE ONE: THE MARGINALS SEPARATE SUPPORTED FROM UNSUPPORTED CLAIMS

Before asking anything of the coherence stage, stage one should be checked against the only external ground truth available, the held-out human labels on BIO. It passes, on the fifty BIO responses (992 supported and 563 not-supported claims): supported claims receive a mean posterior of 0.934 and not-supported claims 0.576, a separation of $+0.357$. The gap is wide but the distributions overlap — both span the full unit interval, and the not-supported claims have the larger spread (0.289 against 0.191) — so stage one is a graded, imperfect signal about support rather than a classifier, which is exactly how the coherence stage consumes it. What matters downstream is that the marginals are genuinely graded rather than saturated: they take 455 distinct values over 6,266 claims. A stage that returned 0.5 everywhere, or 1.0 everywhere, would reduce the two-stage model to the coherence-only one.

The distribution is nonetheless strongly ceiling-weighted, and it is worth reporting exactly rather than merely calling it graded: 82.7% of claims land above 0.99, 10.3% in the informative middle $[0.2, 0.8)$, and 2.3% below 0.01. That concentration is a property of the corpus, not a failure of the estimator, and the human labels are what establish the difference. Bucketing the labelled BIO claims by their marginal, the above-0.99 bucket is 88.1% SUPPORTED against a 63.8% base rate, while the middle bucket inverts to 73.4% NOT-SUPPORTED and the below-0.01 bucket to 74.5%. The ceiling is therefore populated by claims that really are supported, and the graded interior is where the genuinely doubtful claims sit — which is the behaviour a prior should have. It does mean the coherence stage receives a near-uniform prior on five claims in six, so on this corpus the composition is doing most of its work on the remaining sixth. A stage that returned 0.5 everywhere, or 1.0 everywhere, would reduce the two-stage model to the coherence-only one.

9.4 WHAT THE CORPUS SHOWS

Table 14 reports every measure’s mean, spread and distinct-value count over the sample, per miner.

The evidence-free baselines collapse here, and more completely than on the ladders. The failure is not uniform across the baselines, and the distinction matters. Under `llama-3.3` the *LLM judges* are effectively constants: over all 250 responses the direct rating takes **three** distinct values and the G-Eval rubric **five**, with a run-to-run spread across five seeds below 10^{-2} . Within a single

Table 14: The two-stage model and the evidence-free baselines on the factuality corpora: 250 responses, 6,266 claims, five domains, scored under two relation miners with the WINDOWED pair policy. **d** counts the *distinct* values a measure takes over the 250 responses: a measure with *d* near 1 is a constant and cannot rank responses whatever its mean. This corpus has no ordering ground truth, so the column to read is *d*, not the mean.

Miner	Measure	Mean	SD	Range	d
llama-3.3-70b	LCS _{mm}	0.9759	0.0692	[0.27, 1.00]	90
	LCS _{cn}	0.8032	0.2465	[0.26, 1.00]	74
	LCS _{lp}	0.6850	0.4507	[0.00, 1.00]	45
	JUDGE, G-EVAL	0.9998	0.0032	[0.95, 1.00]	5
	JUDGE, DIRECT RATING	1.0000	0.0000	[1.00, 1.00]	3
	ROSCOE-SC (MAX)	0.6252	0.4684	[0.00, 1.00]	51
	ROSCOE-SC (MEAN)	0.9955	0.0112	[0.91, 1.00]	73
	CONTRADICTION RATE	0.9971	0.0078	[0.95, 1.00]	48
	RC (COHESION)	0.5024	0.1808	[0.06, 1.00]	224
	ENTITY GRAPH	0.9724	0.1194	[0.11, 1.00]	17
	LENGTH (CONTROL)	0.2066	0.0461	[0.09, 0.33]	136
	CLAIM COUNT (CONTROL)	0.3916	0.1064	[0.08, 0.78]	37
gpt-oss-120b	LCS _{mm}	0.9475	0.0903	[0.43, 1.00]	148
	LCS _{cn}	0.7107	0.2535	[0.12, 1.00]	134
	LCS _{lp}	0.6713	0.4340	[0.00, 1.00]	85
	JUDGE, G-EVAL	0.9166	0.1017	[0.45, 1.00]	10
	JUDGE, DIRECT RATING	0.8286	0.1619	[0.20, 1.00]	16
	ROSCOE-SC (MAX)	0.5640	0.4969	[0.00, 1.00]	2
	ROSCOE-SC (MEAN)	0.9944	0.0118	[0.89, 1.00]	69
	CONTRADICTION RATE	0.9961	0.0086	[0.93, 1.00]	54
	RC (COHESION)	0.5024	0.1808	[0.06, 1.00]	224
	ENTITY GRAPH	0.9724	0.1194	[0.11, 1.00]	17
	LENGTH (CONTROL)	0.2066	0.0461	[0.09, 0.33]	136
	CLAIM COUNT (CONTROL)	0.3916	0.1064	[0.08, 0.78]	37

dataset they are frequently *literal constants*: the direct rating takes one distinct value on each of BIO, BOOKS and LFOBJ, and G-Eval one on each of ELI5 and LFOBJ — fifty documents, one number. A measure with three distinct values on two hundred and fifty documents has no ordering to be right or wrong about; this is the practical incumbent, and on realistic long-form text it is close to not being a measure at all. The *contradiction-based baselines* are better than that and worse than they look: they are heavily ceiling-skewed — most responses sit at exactly 1.0 for the pairwise contradiction rate — but they do take 48–73 distinct values over the 250, so they rank the minority of responses in which the extractor finds any conflict at all. Their problem on this corpus is not degeneracy but coverage: naturally occurring prose rarely contradicts itself outright, so a measure keyed on contradiction alone is uninformative about most of the corpus by construction.

A saturated estimator degrades ROSCOE’s max to a coin flip. One model–baseline interaction deserves its own note, because it shows how a saturated relation probability propagates into a baseline that appears unrelated to it. ROSCOE’s Self-Consistency aggregates pairwise contradiction with $1 - \max_{ij} p_{ij}^{\text{con}}$. Under llama-3.3 that takes 51 distinct values over the 250 responses. Under gpt-oss-120b it takes **exactly two**: 1.0 on 141 responses and 0.0 on 109, with nothing in between. Its symmetric variant collapses identically (2 values against llama’s 52). The mechanism is immediate once stated — gpt-oss’s contradiction probabilities are pinned at ≈ 1 (§9.2), so the max is 1 whenever *any* pair is labelled a contradiction and 0 otherwise, and the readout degenerates into an indicator for “the extractor found at least one conflict”. The mean variant of the same baseline survives on the same verdicts (1.0 down to 0.99, many values), as does the soft contradiction rate (52 values), because both average over all pairs instead of taking an extremum. The lesson generalizes beyond ROSCOE: an aggregator that reads a single extremal probability inherits that probability’s calibration completely, whereas one that averages is partially protected. It is also a reason to report a baseline under more than one model — the same metric looks serviceable under one estimator and is a coin flip under another.

The cohesion metrics and the controls do vary, and that is the point of including them. Being functions of the surface, they take *more* distinct values than the LCS does — 224 and 209 of 250 for the two cohesion metrics, 136 for response length, against 74 for the consistency readout. That is exactly why variation alone is not evidence of measuring coherence: a character count varies freely and knows nothing. Their spread is the null model a coherence measure must beat on something other than spread, which is what §8.6’s invariance axis tests and this corpus cannot.

The two pair policies disagree, and the disagreement is informative. Windowed mining is forward-only, so it cannot assert a relation that runs backward in assertion order; bidirectional mining can, within the same radius. Because both are scored against identical priors, the gap isolates what the backward arcs contribute — and on this corpus that contribution is small. The bidirectional policy recovers more relations per response (96.4 against 89.8 under `llama-3.3`; 62.8 against 58.6 under `gpt-oss-120b`), yet over 250 responses the headline readout barely separates the two: 0.9796 against 0.9759 for `llama` and 0.9596 against 0.9475 for `gpt-oss`, with the consistency readout moving by 0.0008 and 0.0105 respectively. The direction is consistent — bidirectional scores marginally higher, as more mined support should imply — but the magnitude is a fraction of the between-response spread ($sd \approx 0.25$ on consistency). Backward arcs are therefore not the binding constraint on this corpus: most of what the coherence graph finds in naturally occurring prose is already reachable in assertion order. That is a negative result about the pair policy, not about the model, and it is worth stating because the forward-only policy is the cheaper of the two.

The second model discriminates, and it does not help. The degeneracy above is `llama-3.3`’s `gpt-oss-120b`’s judges are *not* constants: over all 250 responses its direct rating takes sixteen distinct values against `llama`’s three, and it is constant on no domain. Read carelessly that looks like a working judge, and it is worth showing why it is not. Its per-seed ratings resolve to only *four* underlying rungs of the five-point scale (0.25, 0.5, 0.75, 1.0); the apparent resolution comes from averaging five seeds that *disagree with each other*, roughly half the responses showing a 0.25 spread, i.e. the judge moving a full rung on byte-identical input. That is why a judge must be reported with its spread rather than as a point estimate.

Whether the residual variation carries information is a fair question and the answer is: a little, but not reliably. Within a domain the correlation between this judge and the LCS consistency readout is -0.02 on BIO ($n = 50$) and $+0.21$ on ASKHIST ($n = 49$), neither significant; pooled over the two it is $+0.24$ ($t = 2.5$), which is nominally significant but is also what one expects when two measures share a between-domain offset, since ASKHIST is both the lower-scoring domain for the judge and the more conflicted one. We therefore do not claim this judge is pure noise — only that its usable resolution is four rungs, its run-to-run spread is a full rung, and any ordering signal it carries is weak enough to be confounded with domain. The two judges fail in the two different ways available to a saturated rater, one returning a constant and the other returning four rungs with rung-scale noise, and the LCS readouts take 34–48 distinct values on the same items either way.

The readouts respond to relational structure, not to the priors. A measure built by composing two stages invites the worry that the second stage is only echoing the first. The corpus lets us check, because the five domains disagree sharply on the consistency readout, and that ordering is *not* the ordering of the priors. It is the ordering of mined conflict. Writing each domain as (conflict-bearing responses of 50, mean consistency), where a response is conflict-bearing when its contradiction-free ceiling exceeds its partition function: ASKHIST (40, 0.589), ELI5 (23, 0.757), BOOKS (17, 0.820), LFOBJ (7, 0.927) and BIO (6, 0.922). The readout falls monotonically in conflict prevalence across all five, the single inversion being the two least-conflicted domains, which are within 0.005 of each other. Meanwhile the correlation between a response’s mean prior and its consistency is -0.13 , i.e. nil. Stage two is therefore responding to the relational structure it mines rather than to the factuality signal it inherits, which is what the composition was supposed to buy and is not something the ladder corpus can show, its items being matched on claim content by construction.

The per-response form of that test is sharper than the per-domain one, and it replicates across both miners. Splitting responses by whether their mined graph contains any conflict edge at all, conflict-bearing responses score 0.496 on the consistency readout against 0.985 for conflict-free ones under `llama-3.3` — a gap of 0.489 over 250 responses (Welch $t = 55$) — and 0.483 against 0.921 under `gpt-oss-120b`, a gap of 0.438 (Welch $t = 28$). The readout is thus separating responses by the

presence of relational conflict, not by domain, length or prior mass, and it does so under a miner whose relation *probabilities* are uninformative as well as one whose are not — which localizes what the readout needs from the extractor to the relation *type*.

The degeneracy of the judges is not an artifact of pooling domains. Aggregate distinct-value counts could in principle hide a judge that discriminates within a domain while being offset between them. It does not. Per dataset (50 responses each) `llama`’s direct rating takes 1, 1, 1, 2, 2 distinct values on BIO, BOOKS, LFOBJ, ASKHIST and ELI5 — a literal constant on three of the five — while the LCS consistency readout takes 35, 13, 10, 17 and 11. The gap is a property of each domain separately, not of the pooled corpus. `gpt-oss`’s rating is not constant on any domain (9–13 values), which is the behaviour the next paragraph takes apart.

Qualification. This corpus has no ordering ground truth, so nothing here establishes that the LCS orders responses *correctly* — only that it discriminates among them where the incumbent judges do not, that it separates conflict-bearing from conflict-free responses by roughly half the scale under both miners, and that its inputs are sound. The ladder corpus is where correctness of ordering is tested; this corpus is where the measures meet realistic text. Neither result substitutes for the other, and the surface baselines above are a reminder that discrimination alone is cheap: the argument requires both axes.

10 WORKED EXAMPLES

We work five examples. The first is the smallest complete instance of the two-stage model: one claim set, two responses, exact inference; the second takes one of those two responses on its own and follows it from relation graph to MRF to score, which is the pipeline in miniature. The third is a richer sixteen-claim graph that exercises the conflict taxonomy. The fourth is a claim-identical pair that isolates ordering. The first four are hand-authored, which bounds what they can show: they demonstrate that the model reads a graph correctly, not that the estimator finds such graphs in text it has never seen. The fifth is therefore a corpus rather than an example — one hundred machine-generated responses whose incoherence is planted by construction, scored end to end with no hand-drawn relations anywhere.

10.1 A FIVE-CLAIM PAIR, END TO END

The point of this example is that factuality and coherence are separable, so we hold the first fixed and vary only the second. Five claims about the Voyager 1 mission are fixed once and for all, three of them true of the world and two false; stage one is taken as given, contributing the priors $\pi = 0.9$ for each true claim and $\pi = 0.1$ for each false one through Eq. 9:

	π_i	Claim
a_1	true 0.9	Voyager 1 crossed the heliopause in August 2012.
a_2	true 0.9	After the crossing, Voyager 1’s instruments measured a sharp rise in plasma density.
a_3	true 0.9	Voyager 1 is therefore in interstellar space.
a_4	false 0.1	Voyager 1’s plasma wave instrument failed permanently before the crossing.
a_5	false 0.1	Voyager 1 has left the Solar System entirely.

Two responses then assert *these same five claims* and differ only in how they arrange them. Because the claim set is identical, every per-claim factuality check returns the same verdict on both; the relation graphs (Figure 2) and the MRFs they induce (Figure 3) are what differ. Table 15 gives both responses in full, with the span asserting each claim marked $[a_i]$ at its first mention.

The italicized spans are the ones the estimator reads as relations, and they are the whole difference between the two. Table 16 lays the two relation sets side by side, one row per claim pair, with each relation’s discourse sense and the coupling it compiles to. In A the false claims are named, attributed, and then refused on a stated ground, which compiles to a CONTRADICTION edge from a_3 . In B the same falsehoods are named and then *used* — a_4 is offered as the reason a_3 holds — while the response also denies a_3 outright and, in its last sentence, denies a_4 and a_5 too. Mentioning a




Figure 2: Relation graphs for the five-claim worked pair (§10). Blue nodes are claims true of the world ($\pi_i = 0.9$), red nodes false ($\pi_i = 0.1$); the *same* five claims appear in both panels. Solid blue: ENTAILMENT; dashed red: CONTRADICTION. Line width grows with the relation weight p . The panels differ only in *how the response arranges* the claims, which is the whole of what coherence measures — no per-claim factuality check can separate them.




Figure 3: The coherence MRF of Eq. 2 for the two responses, as an explicit factor graph. Circles are the binary claim variables; grey squares the unary factors $\phi_i = [1 - \pi_i, \pi_i]$ that carry the priors; black squares the pairwise factors ψ_{τ_r, p_r} of Table 2, one per mined relation. **The priors enter only through the ϕ_i :** the claim-level potentials are identical across the two panels, and every difference in score comes from the ψ_r — their number, their endpoints, and their couplings.

falsehood is not what costs B its score; asserting it and reasoning from it is. All values below are exact, by $2^5 = 32$ -world enumeration.

Example 2 (Coherent). Response A’s four relations (Table 16, column A) are a two-edge support chain over the true claims and one rejection of each false claim. The model returns $\text{LCS}_{\text{mm}} = 0.590$ and $\text{LCS}_{\text{cn}} = 0.963$ with $\log Z = -0.95$. Every claim finishes on the correct side of its own prior: the true claims sit at 0.939, 0.989 and 0.992 — each above its 0.9 prior, because the chain supplies mutual support — and the false ones at 0.013 and 0.020, pushed below 0.1 by the rejections. Nothing in the response fights anything else, so no claim is penalized.

Response A (coherent) — $\text{LCS}_{\text{mm}} = 0.590$, $\text{LCS}_{\text{cn}} = 0.963$, $\log Z = -0.95$

Voyager 1 crossed the heliopause in August 2012, [a₁] and in the days that followed its instruments measured a sharp rise in plasma density [a₂] — the signature of the denser medium beyond, and *the measurement that established the boundary had been passed*. Voyager 1 is therefore in interstellar space. [a₃] Two persistent misconceptions are worth clearing up. The first is that the probe’s plasma wave instrument had failed permanently before the crossing; [a₄] *it cannot have, since it is precisely that instrument which returned the density profile the conclusion rests on*. The second is that Voyager 1 has left the Solar System altogether. [a₅] *Being in interstellar space is not the same thing: the Sun’s gravitational reach extends through the Oort cloud, far beyond the probe’s present distance*.

Response B (incoherent) — $\text{LCS}_{\text{mm}} = 0.494$, $\text{LCS}_{\text{cn}} = 0.414$, $\log Z = -3.78$


Voyager 1 crossed the heliopause in August 2012, [a₁] and after the crossing its instruments measured a sharp rise in plasma density. [a₂] *That said, the probe cannot really be counted as being in interstellar space*. [a₃] The plasma wave instrument had in fact failed permanently before the crossing, [a₄] *and it is precisely that failure which tells us the probe reached interstellar space*. The instrument failure also means Voyager 1 has left the Solar System entirely. [a₅] *Of course, the August 2012 crossing rules out any such permanent failure, and the density measurement is flatly inconsistent with the probe having left the Solar System*.

Table 15: The two responses in full. They assert the *same* five claims of §10.1; each claim’s first mention is marked [a_i]. Italicized spans are the ones the estimator reads as relations. In A they discharge the two false claims; in B they reason *from* them.


Table 16: The mined relations of both responses, one row per claim pair. S is the Level-2 discourse sense, τ the Level-1 coupling it compiles to (Table 4), p the weight. E abbreviates ENTAILMENT, C CONTRADICTION. These are exactly the edges of Figure 2 and the factors ψ_r of Figure 3; a dash means the response draws no relation over that pair.

Pair	Response A (coherent)				Response B (incoherent)				What differs
	Edge	s	τ	p	Edge	s	τ	p	
{a ₁ , a ₂ }	$a_1 \rightarrow a_2$	Cau-Eff	E	.90	$a_1 \rightarrow a_2$	Cau-Eff	E	.90	The one relation both agree on.
{a ₂ , a ₃ }	$a_2 \rightarrow a_3$	Evidence	E	.85	$a_2 \rightarrow a_3$	Contrast	C	.90	A supports its conclusion; B denies it. Support becomes conflict.
{a ₃ , a ₄ }	$a_3 \rightarrow a_4$	Contrast	C	.90	$a_4 \rightarrow a_3$	Evidence	E	.88	The arrow reverses. A rejects a ₄ ; B makes it a premise.
{a ₃ , a ₅ }	$a_3 \rightarrow a_5$	Contrast	C	.85		—			A refuses a ₅ ; B never draws the distinction.
{a ₄ , a ₅ }		—			$a_4 \rightarrow a_5$	Cau-Eff	E	.85	B alone chains one falsehood into the other.
{a ₁ , a ₄ }		—			$a_1 \rightarrow a_4$	Contrast	C	.90	B denies a ₄ , which it had just used as a premise — this strands a ₃ .
{a ₂ , a ₅ }		—			$a_2 \rightarrow a_5$	Contrast	C	.88	Likewise for a ₅ : asserted, then denied.
Totals		2 E, 2 C				3 E, 3 C			A’s conflicts point <i>outward</i> , at the false claims; half of B’s point <i>inward</i> .

Example 3 (Incoherent). *Response B keeps only $a_1 \rightarrow a_2$ and replaces the rest with six relations that pull against each other (Table 16, column B): it denies its own conclusion, promotes the false a₄ to a premise for that same conclusion, chains a₄ into a₅, and then denies both falsehoods in its closing sentence. Three of its six relations are conflicts, and two of those point at claims the response itself asserts. The score falls to $\text{LCS}_{\text{mm}} = 0.494$ and $\text{LCS}_{\text{cn}} = 0.414$, with $\log Z = -3.78$. The diagnostic that localizes the damage is per-claim: the true a₃ is dragged to 0.513, far below its own 0.9 prior, because its only support arrives from a₄, which the response itself denies — and a premise the text refuses cannot transmit support. Three claims finish below their own priors, against two in Response A, and a₃ is the one that matters: a claim true of the world, demoted to a coin flip by the argument built around it.*



(a) The relation graph: three ENTAILMENT (solid blue), three CONTRADICTION (dashed red); width grows with p .



(b) The MRF it induces: one ψ_r per edge (black), one ϕ_i per prior (grey).


Readout	Value	What it says
LCS _{mm} (selected)	0.494	Mean posterior marginal: the expected fraction of B's claims its own argument upholds. Against a prior-only mean of 0.580, the arrangement <i>costs</i> B support rather than adding any.
LCS _{cn}	0.414	Two-term consistency: little over a coin flip that no asserted conflict is realized, against 0.963 for the coherent arrangement of the same claims.
LCS _{re}	0.348	The reified-node variant, for comparison; it agrees on the ordering.
log Z	-3.78	Global mass surviving the constraints (-0.95 for the coherent arrangement): the six factors are jointly hard to satisfy.
LCS _{lp}	(0.267)	Parenthesized because it is <i>not</i> comparable across responses: its references $\log Z_{\max} = -1.65$ and $\log Z_{\min} = -4.55$ are rebuilt from B's own edge skeleton, so it reports distance to <i>that</i> ceiling. A within-response diagnostic only.

Per-claim marginals (the diagnostic the scalar cannot give): a_1 0.960, a_2 0.973, **a_3 0.513**, a_4 0.006, a_5 0.018. Three of five finish below their own prior. The one that matters is a_3 : a claim *true of the world*, with prior 0.9, dragged to a coin flip. Its only support is ψ_3 , arriving from the false a_4 — which ψ_4 has the response denying. A premise the text refuses cannot transmit support, so a_3 is left holding an argument that undercuts itself. All values exact, by $2^5=32$ -world enumeration.

Figure 4: **Response B, anatomized.** The incoherent arrangement of the five-claim pair (§10.1) at two levels of resolution, and what inference over it returns. Left: the relation graph, blue nodes true of the world and red false. Right: the Markov random field of Def. 1 that graph induces, as an explicit factor graph — the mapping is one-to-one, six edges to six pairwise factors ψ_r , five priors to five unary factors ϕ_i . Below: the four readouts of §5 and the per-claim marginals. The reading to take from the pair is that the ϕ_i are identical to those of the coherent arrangement — the claim set and every factuality verdict are unchanged — so the whole of the score difference is carried by the ψ_r , and the whole of the damage is localized by the marginals to a_3 .

Two things are worth reading off this pair. First, the priors enter *only* through the unary factors ϕ_i of Figure 3, which are identical in both panels; the entire $0.590 \rightarrow 0.494$ drop is produced by the ψ_r , so it is attributable to arrangement alone. Second, a claim being true of the world does not protect it: coherence is a property of the joint distribution, and a_3 is suppressed by *how* it is argued for, not by what it says. This is the mechanism Remark 2 anticipates and the ordering pair of §10.4 exhibits at scale.

A caution on readouts. LCS_{lp} is 0.194 for Response A and 0.267 for Response B, which inverts the order — correctly, and for a reason Eq. 8 makes plain: its references $\log Z_{\max}$ and $\log Z_{\min}$ are recomputed *per response* (-0.49 and -1.06 for A, -1.65 and -4.55 for B), so it reports how close a response comes to the ceiling of *its own* edge skeleton, not how the two responses compare. It is a within-response diagnostic. Raw log Z is comparable across responses and moves the right way, by 2.83 nats.



(a) **Incoherent** (base). Two EXCLUSIVE conflicts (dashed): $a_7 \leftrightarrow a_{10}$ is *unresolved*; $a_{11} \leftrightarrow a_{12}$ is a concession the holding a_{13} (green) resolves. $\text{mm} = 0.607$, $\log Z = -9.64$. (b) **Coherent**. The identical support structure with both conflict edges removed; a_{10} and a_7 are no longer suppressed. $\text{mm} = 0.620$, $\log Z = -8.25$, $\text{lp} = 1.00$.

Figure 5: Relation graphs for the AeroParts recall report (§10). Solid blue: ENTAILMENT; teal: EQUIVALENCE; dashed red, double-headed: EXCLUSIVE. Line width grows with the relation weight p . **The two panels differ only in the two conflict edges**, so the $0.607 \rightarrow 0.620$ lift and the 1.39 nats of recovered mass are attributable to conflict structure alone.

10.2 RESPONSE B AT TWO LEVELS OF RESOLUTION

The pair above varies the arrangement and reads off the difference. It is worth pausing on one of the two responses on its own, because the step from a relation graph to a probability model is the step at which most of the modelling content sits, and on a five-claim instance it can be shown in full rather than described. Figure 4 puts Response B’s relation graph next to the MRF that graph induces, and lists what inference over that MRF returns.

The two panels are the same object at two resolutions, and the translation between them is mechanical: Def. 1 sends each mined relation to exactly one pairwise factor and each claim to exactly one unary, so B’s six edges become $\psi_1, \dots, \psi_6$ and its five priors become $\phi_1, \dots, \phi_5$. Nothing is added in the crossing and nothing is dropped; what the right panel supplies is that the priors, implicit in the node colours on the left, become factors of the same kind as the relations, and therefore quantities the inference trades off against one another. This is what makes the two-stage composition of §6 a substitution rather than a redesign: replacing a uniform ϕ_i with a factuality posterior changes a number in an existing slot.

Three things in the score panel are worth reading off. First, $\text{LCS}_{\text{mm}} = 0.494$ sits *below* the prior-only mean of 0.580, so B’s argument is not merely failing to add support — it is removing some, which a measure that only counted satisfied relations could not report (Prop. 2). Second, the readouts agree on the ordering where they are entitled to: LCS_{mm} , LCS_{cn} , LCS_{re} and raw $\log Z$ all place B below the coherent arrangement, while LCS_{lp} is parenthesized because its references are rebuilt per response and so cannot be compared across the pair. Third, and the reason a scalar is not the whole output, the per-claim marginals localize the damage: a_3 is true of the world and carries a 0.9 prior, yet finishes at 0.513. Reading the mechanism off the right panel, a_3 ’s only support is ψ_3 , which arrives from a_4 ; but ψ_4 encodes the response denying a_4 , and ψ_2 has it denying a_3 directly. A premise the text refuses cannot transmit support, so the claim that should have been B’s conclusion is left at a coin flip — by the argument built around it rather than by anything it says.

10.3 A SIXTEEN-CLAIM RECALL REPORT

The second example is a sixteen-claim aviation recall report, which exercises the conflict taxonomy more fully than the five-claim pair: it carries both an unresolved conflict and a conceded one. Its relation graph is twelve support edges — an entailment spine $a_1 \rightarrow a_2 \rightarrow \dots \rightarrow a_6$ (shipment certified to spec, installation, entry into service, in-flight failures, regulatory probe, grounding order), branches to the pulled-from-service claim, the maintenance-log evidence, the tribunal’s damages award and the regulator’s bulletin, and one EQUIVALENCE edge — plus two conflicts. Every value below is exact, by 2^{16} enumeration, and the graph is reproduced edge-by-edge in Figure 5 with each edge’s weight shown by line thickness.

Example 4 (Incoherent). *The report asserts both that no passengers or crew were harmed and that a wire report claimed three passengers died. These are exhaustive alternatives rather than mere opposition: they cannot both hold and cannot both fail, so the relation is EXCLUSIVE with $p = 0.93$ — and nothing in the response resolves it. A second tension, between the airline’s pilot-error argument and the final report’s metallurgical-defect finding ($p = 0.80$), is a **Concession** that the tribunal’s holding does resolve, and Eq. 3 discounts it accordingly. The model returns $\text{LCS}_{\text{mm}} = 0.607$ with $\log Z = -9.64$. The diagnostic that localizes the incoherence is per-claim: the marginal of the claim on the losing side of the unresolved conflict is dragged below its own prior, which no aggregate score would reveal. Applying the concession discount alone — resolving the second tension while leaving the first — raises the score to 0.612, as C2 requires.*

Example 5 (Coherent). *Removing the two conflict edges and changing nothing else leaves the identical support structure (Figure 5b). The score rises to $\text{LCS}_{\text{mm}} = 0.620$ with $\log Z = -8.25$ and $\text{LCS}_{\text{lp}} = 1.00$: the response is now as coherent as its own edge skeleton permits, which is what the ceiling of Eq. 8 means. The suppressed claim recovers to its prior.*

Because the two panels of Figure 5 differ *only* in the two conflict edges, the $0.607 \rightarrow 0.620$ lift and the 1.39 nats of recovered probability mass are attributable to conflict structure alone — which is the property the ladder construction of §7 generalizes from one pair to a family. The full ladder over this family is monotone throughout: $0.586 < 0.607 < 0.612 < 0.613 < 0.620$ over the rungs *worse, base, concession-resolved, one-conflict-fixed, coherent*.¹

10.4 A CLAIM-IDENTICAL ORDERING PAIR

This pair makes the order-sensitivity claim concretely. Two summaries of one appellate judgment are built from an *identical* set of eighteen claims — the same atom texts, one a permutation of the other — differing only in whether self-serving assertions precede or follow the findings that qualify them. The faithful ordering scores $\text{LCS}_{\text{mm}} = 0.95$ and the self-serving-first ordering 0.90. No per-claim factuality check can separate them, because every claim is byte-identical; what differs is which claims the response leaves standing unqualified at the point where it asserts them.

It is worth being exact about *how* the two scores differ, because Remark 2 might seem to forbid it: nothing in Eq. 2 references position, so an edit preserving the relation set cannot move any readout. The pair does not preserve the relation set. Mining is response-grounded (§4.5), so the relations the estimator draws are a function of the prose, not of the claim list: the faithful summary states each tension beside the truth or holding that settles it, which the estimator reads as a **Concession** the text resolves and Eq. 3 discounts; the adversarial summary leaves the same claims standing unqualified until a later holding overrides them, which reads as a live conflict. The claim set is what factuality sees and it is identical; the relation set is what coherence sees and it is not.

11 LIMITATIONS

The model is pairwise, and some incoherence is not. Definition 1 admits only pairwise factors, so genuinely k -ary structure lies outside the family. This is the sharpest limitation of the measure and it is worth stating exactly rather than in passing, because it is not a matter of degree: the relevant conflicts are not approximated badly, they are unrepresentable.

Proposition 4 (Jointly inconsistent, pairwise consistent). *Let $\mathbb{P}(\mathbf{a})$ factorize as in Eq. 2 over a_1, a_2, a_3 with non-negative factors. If $\mathbb{P}(a_1=1, a_2=1, a_3=1) = 0$ then $\mathbb{P}(a_i=1, a_j=1) = 0$ for some pair $i \neq j$. Consequently no member of the family represents “the three claims cannot all hold, but any two can”.*

Proof. By Eq. 2, $\mathbb{P}(1, 1, 1) = Z^{-1} \prod_i \phi_i(1) \prod_{i < j} \psi_{ij}(1, 1)$, a product of non-negative terms, which vanishes only if some term does. If $\phi_i(1) = 0$ then a_i is false almost surely and both pairs

¹Treating both tensions as plain CONTRADICTION rather than EXCLUSIVE — the three-coupling model that omits the exhaustiveness distinction — gives $\text{LCS}_{\text{mm}} = 0.587$ and $\log Z = -9.75$ for Example 4, with Example 5 unchanged at 0.620 and -8.25 . The distinction of §4.2 is worth 0.020 of score here, and the reason it matters more than that figure suggests is Table 6: with EXCLUSIVE conflicts, readout (b)’s conflict term is pinned at 1.000 and the readout fails outright.

containing i have probability zero. If $\psi_{ij}(1, 1) = 0$ then every world with $a_i = a_j = 1$ carries that factor, so $\mathbb{P}(a_i=1, a_j=1) = 0$ directly. A ternary factor that vanishes only at $(1, 1, 1)$ has neither consequence. $\square$

The gap is quantitative as well as qualitative. With uniform priors, a ternary factor supported off $(1, 1, 1)$ gives $\mathbb{P}(1, 1, 1) = 0$ while leaving each pair jointly true with probability 0.143; the best pairwise attempt that drives $\mathbb{P}(1, 1, 1)$ below 10^{-4} leaves the pairs at 0.028, and a random search over 2×10^5 pairwise parameterizations found nothing better. The pairwise model can only suppress the bad world by also suppressing the good ones that share a $(1, 1)$ cell with it. Every figure in this paragraph is exact by 2^3 enumeration and reproduced by a released script and regression test, so the proposition is checked and not merely asserted.

This is not a hypothetical. Prompting a frontier model to plant a single arithmetic inconsistency in an otherwise ordinary encyclopedic paragraph produced a passage whose decomposition places the conflict across *three* claims: that a league comprises 356 teams, that those teams form 32 conferences, and that each conference has exactly 10 member schools. No two of the three conflict; $32 \times 10 \neq 356$ requires all three. Mining that passage with the estimator of §4.5 returns *zero* conflict edges, and the readouts accordingly report it as coherent ($\text{LCS}_{\text{mm}} = 0.68$, $\text{LCS}_{\text{lp}} = 1.00$). The omission is not an extraction failure: run under the `all_pairs` policy, which scores every one of the $\binom{n}{2}$ candidate pairs, the result is unchanged, and querying the estimator on each of the three pairs in isolation returns NONE or ENTAILMENT — never a conflict. Proposition 4 says no reweighting of those factors could have helped.

Two milder consequences of the same restriction are worth naming. Transitivity is not enforced — the model will happily hold $a \Rightarrow b$ and $b \Rightarrow c$ at high weight while leaving a and c unrelated, since it has no rule tying them. And a claim conceded by two independent holdings is discounted twice rather than once. All three are expressible as first-order rules in the Markov logic view of Appendix C, which is the natural route past this limitation and which we have not taken; Remark 1’s closure guarantee is precisely what buys tractability here and precisely what costs us this class of defect.

Hand-set constants. The design deliberately avoids free parameters — criterion R4 exists for that reason, and it is what disqualifies readout (c) with its prior ρ . Two constants remain. The concession discount $\lambda = 0.45$ in Eq. 3 sets how much a resolved tension is softened, and the resolver-detection heuristic identifies holdings by cue phrases and lexical overlap. Both are defensible and neither is derived. A principled treatment would learn the discount from labelled resolved and unresolved concessions, which the corpus of §7 supplies in principle but not at the scale required.

Relation estimation dominates the error. §8 is unambiguous that the bottleneck is not the model. At 56.3% coupling-level F1, roughly half the factors in a mined network are wrong or missing, and the 25–43 point ladder-agreement gap between mined and gold arms follows from that. Everything the paper establishes about the readouts is established on gold or exact graphs; the mined-arm numbers should be read as a measurement of current extraction quality, not as the measure’s ceiling.

Approximate inference is not characterized. Production inference is weighted mini-bucket with i -bound 6, while every exact value reported here comes from 2^n enumeration, feasible only for $n \leq 20$ claims. The two agree on the fixtures we can check both ways, but we have not characterized the approximation error at length, and the augmented networks of §5.2 — which add one auxiliary variable per relation — are the case where it would most plausibly bite.

Empirical scope. The generated corpus is fourteen families and seventy items spanning all four ladder types, with single-sample mining whose measured run-to-run movement is at most 0.014 F1 across two independent runs of the same configuration. All five baseline families of §7.1 are measured on the ladder corpus (§8.5). On the 50 readout-independent ladder assertions all three graded readouts now sit above every evidence-free baseline, which supersedes the tie we reported on ten assertions; we flag the direction of that correction in §8.5 rather than quietly restating it. Liu & Strube (2025) is cited rather than run, for the output-type reason given above.

Three things remain genuinely open. First, the control and order ladders — the invariance experiments — reveal that *mined* graphs are strongly order-dependent (0/45 and 0/15 on every mined arm, with

a five-rung edge count of 9, 16, 20, 20, 17 under pure sentence reordering), and we cannot yet say how a human annotator would score those same rungs, which is the comparison that would separate “the extractor is unstable” from “the invariance requirement is too strict”. Second, human correlation generally, which is what would establish that the construct being ordered is the one readers care about. Third, judge dependence: the two LLM judges we ran disagree by up to 42 points on the identical assertions, so any single-judge number here — including ours — carries that uncertainty. And the corpus is LLM-generated throughout, gold relation labels included, with LLM-committee validation; a bias shared between generator and committee would not be caught by it.

12 CONCLUSION

Coherence is a property of a joint distribution over what a response asserts, and it can be read off a Markov random field whose factors are the relations the response itself draws. Three things follow from taking that view seriously. First, the factor tables matter: the ones that are right for grounding a claim against evidence charge a premise for being a premise, and a coherence measure built on them spends its dynamic range in the wrong place (Prop. 1). Second, the relation vocabulary can be rich at the annotation layer and must be narrow at the inferential one; twelve discourse senses compiling deterministically onto five couplings, each fixed by the world it penalizes, keeps the model a closed pairwise family while leaving the labels interpretable. Third, and least expected, not every reading of “how coherent is this?” is admissible. Of four candidates, two are non-monotone in conflict resolution for a single diagnosable reason — they measure whether a conflict is active rather than what should be believed (Prop. 3) — and a fifth, apparently the most natural of all, is maximized by asserting no relations whatever (Prop. 2). The mean posterior marginal survives all five criteria, needs one inference query, and localizes an incoherence to the claims responsible.

Measuring such a score requires controlled contrasts, and ladders of responses over one claim set with declared orderings supply them. Given gold relation graphs the readouts recover 82.7% of those orderings; given graphs mined from the response text, 34.7–37.1%. Because the two conditions differ in nothing but the relation graph, the gap localizes the open problem: the probabilistic model is adequate, and grounded relation extraction, at 52% coupling-level F1 at best and not invariant to sentence order, is the bottleneck. That is a more tractable target than “measure coherence”, and it is one that improvements can be measured against.

AI USE STATEMENT

Generative AI is both an object of study and a component of the artifacts released here, and we disclose both roles. The relation estimator of §4.5 is LLM-based by design: its prompts, its logprob-derived confidences and its failure modes are contributions of the paper. The evaluation corpus of §7 is LLM-generated throughout, including its gold relation labels, and validated by an independent LLM committee; §11 states the consequences for treating it as ground truth. We additionally used AI assistants for code implementation and for editorial assistance on this manuscript. All LLM-generated code was reviewed by the authors and is covered by a test suite that pins the exact numerical values reported in §10 and Table 6; every number in §8 was traced to a stored experiment record. The research questions, the taxonomy, the readout definitions, the propositions and their proofs are the authors’ own. We take responsibility for the final content of this work.

ETHICS STATEMENT

This work involves no human subjects and no personally identifying data. The fixtures include a summary of a published appellate judgment, used as a worked example of ordering-dependent coherence; it is a matter of public record and no private individual is the subject of analysis. The principal ethical risk is misreading: a coherence score is not a truth score, and a response can be perfectly coherent and entirely false. That is precisely why §6 keeps the two axes separable, and we caution against deploying LCS alone as a factuality filter — a fluent, internally consistent fabrication is exactly the artifact it scores highest. The released fixtures and corpus deliberately contain false claims, labelled as such, and should not be redistributed without those labels.

REPRODUCIBILITY STATEMENT

The model is fully specified in the paper: Definition 1 gives the joint, Table 2 both families of potentials, Def. 2 the penalized sets, Table 4 the compile map, and Eqs. 5–8 the four readouts together with the auxiliary-variable constructions two of them require. Algorithms 1 and 2 give candidate selection and the end-to-end procedure. Appendix A reproduces all four estimation prompts verbatim, including the trailing token boundary that the surrogate readout depends on. The exact values in §10 and Table 6 are computed by explicit 2^n enumeration — 2^5 for the five-claim pair, 2^{16} for the recall report — and are pinned by regression tests in the released code, so they are reproducible without LLM access. The relation graphs they use are given edge-by-edge in Figures 2 and 5, the five-claim pair’s claim set, priors and gold relations ship as a fixture, and a single script recomputes every number it reports. The k -ary representability figures of §11 (Prop. 4) are likewise exact by 2^3 enumeration, with the seeded random search and the published table both reproduced by a released script and pinned by regression tests. §8 names every model used and states the sampling regime and its measured variance. Code, fixtures and corpus are released.

REFERENCES

- Regina Barzilay and Mirella Lapata. Modeling local coherence: An entity-based approach. *Computational Linguistics*, 34(1):1–34, 2008.
- Farima Fatahi Bayat, Kun Qian, Benjamin Han, Yisi Sang, Anton Belyi, Samira Khorshidi, Fei Wu, Ihab F. Ilyas, and Yunyao Li. FLEEK: Factual error detection and correction with evidence retrieved from external knowledge. In *Proceedings of EMNLP 2023: System Demonstrations*, pp. 124–130, 2023. arXiv:2310.17119.
- Debarun Bhattacharjya, Balaji Ganesan, Junkyu Lee, Radu Marinescu, Katsiaryna Mirylenka, Michael Glass, and Xiao Shou. SIMBA UQ: Similarity-based aggregation for uncertainty quantification in large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025. arXiv:2510.13836.
- Alexander R. Fabbri, Wojciech Kryściński, Bryan McCann, Caiming Xiong, Richard Socher, and Dragomir Radev. SummEval: Re-evaluating summarization evaluation. *Transactions of the Association for Computational Linguistics*, 9:391–409, 2021.
- Olga Golovneva, Moya Chen, Spencer Poff, Martin Corredor, Luke Zettlemoyer, Maryam Fazel-Zarandi, and Asli Celikyilmaz. ROSCOE: A suite of metrics for scoring step-by-step reasoning. In *Proceedings of ICLR*, 2023. arXiv:2212.07919.
- Saurav Kadavath, Tom Conerly, Amanda Askell, Tom Henighan, Dawn Drain, Ethan Perez, et al. Language models (mostly) know what they know. *arXiv preprint arXiv:2207.05221*, 2022.
- Jiwei Li and Dan Jurafsky. Neural net models of open-domain discourse coherence. In *Proceedings of EMNLP*, pp. 198–209, 2017.
- Qiang Liu and Alexander Ihler. Bounding the partition function using Hölder’s inequality. In *Proceedings of the International Conference on Machine Learning (ICML)*, 2011.
- Wei Liu and Michael Strube. Joint modeling of entities and discourse relations for coherence assessment. In *Proceedings of EMNLP*, pp. 21910–21926, Suzhou, China, 2025. Association for Computational Linguistics. <https://aclanthology.org/2025.emnlp-main.1113/>.
- Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of EMNLP*, 2023. arXiv:2303.16634.
- Potsawee Manakul, Adian Liusie, and Mark J. F. Gales. SelfCheckGPT: Zero-resource black-box hallucination detection for generative large language models. In *Proceedings of EMNLP*, 2023. arXiv:2303.08896.
- William C. Mann and Sandra A. Thompson. Rhetorical structure theory: Toward a functional theory of text organization. *Text*, 8:243–281, 1988.

- Radu Marinescu, Debarun Bhattacharjya, Junkyu Lee, Tigran Tchirakian, Javier Carnerero Cano, Yufang Hou, Elizabeth Daly, and Alessandra Pascale. FactReasoner: A probabilistic approach to long-form factuality assessment for large language models. In *Findings of the Association for Computational Linguistics: EMNLP 2025*, 2025. arXiv:2502.18573.
- Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of EMNLP*, 2023.
- Rashmi Prasad, Bonnie Webber, Alan Lee, and Aravind Joshi. Penn discourse treebank version 3.0. LDC2019T05, Linguistic Data Consortium, 2019.
- Matthew Richardson and Pedro Domingos. Markov logic networks. *Machine Learning*, 62(1–2): 107–136, 2006.
- Yixiao Song, Yekyung Kim, and Mohit Iyyer. VeriScore: Evaluating the factuality of verifiable claims in long-form text generation. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024.
- Julius Steen and Katja Markert. How to find strong summary coherence measures? a toolbox and a comparative study for summary coherence measure evaluation. In *Proceedings of COLING*, pp. 6035–6049, 2022. <https://aclanthology.org/2022.coling-1.527/>.
- Jerry Wei, Chengrun Yang, Xinying Song, Yifeng Lu, Nathan Hu, Jie Huang, Dustin Tran, Daiyi Peng, Ruibo Liu, Da Huang, Cosmo Du, and Quoc V. Le. Long-form factuality in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Qiming Xie, Wenjie Zheng, Xiangqing Shen, and Rui Xia. Lorefact: Bridging the logic gap in fact-checking. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 6974–6996, 2026.
- Miao Xiong, Zhiyuan Hu, Xinyang Lu, Yifei Li, Jie Fu, Junxian He, and Bryan Hooi. Can LLMs express their uncertainty? an empirical evaluation of confidence elicitation in LLMs. In *International Conference on Learning Representations (ICLR)*, 2024.
- Wei Zhao, Michael Strube, and Steffen Eger. DiscoScore: Evaluating text generation with BERT and discourse coherence. In *Proceedings of the 17th Conference of the European Chapter of the Association for Computational Linguistics (EACL)*, pp. 3865–3883, 2023. arXiv:2201.11176.

A THE ESTIMATION PROMPTS, VERBATIM

The four prompts below are reproduced exactly as they are issued, extracted programmatically from the released source so that this appendix cannot drift from the code. Placeholders in double braces are substituted at call time; `{{sense_menu}}` is substituted at build time from one of the two menus in §A.5.

A.1 PROMPT A, FIRST GENERATION — DISCOURSE SENSE AND COUPLING

6,148 characters. Placeholders: `{{sense_menu}}`, `{{response}}`, `{{atom_a}}`, `{{atom_b}}`.

```

1
2
3 Instructions:
4 You are given the full RESPONSE a model produced, and two atomic claims, A and B, both
  ↳ taken FROM THAT RESPONSE, in their order of appearance (A comes before B). Your task is
  ↳ to decide the discourse/logical relation FROM A TO B, following the steps below.
5
6 IMPORTANT -- ground your decision in the response. Assert a coupling ONLY if the response
  ↳ ITSELF draws that connection between A and B (as written, or as a clear step in the
  ↳ author's argument/narrative). Do NOT assert a relation that is merely plausible in
  ↳ general but that the response does not actually make. If A and B both appear in the
  ↳ response yet the response draws no logical or discourse dependence between them, the
  ↳ answer is None.
```

2214 7
2215 8 1. Reason step by step, referring to the response: does the response present A as causing,
2216 ↳ enabling, providing evidence for, restating, elaborating, temporally preceding,
2217 ↳ contrasting with, or contradicting B? Is B a claim the response later withdraws or that
2218 ↳ a holding resolves? Consider the direction (A to B). If the response links A and B only
↳ indirectly through other claims, or not at all, that is None.

2219 9
2220 10 2. Name the DISCOURSE SENSE, one of: {{sense_menu}}.
2221 11 - Cause-Effect: A causes/leads to B. Effect-Cause: A is the effect, B its cause.
2222 12 - Evidence: A provides evidence for B. Condition: A is a condition for B.
2223 13 - Restatement: A and B assert the same thing. Instantiation: A is a general claim, B a
↳ specific instance (or vice versa).
2224 14 - Contrast: A and B are in opposition but NOT exhaustive (they need not cover all
↳ possibilities; both could conceivably be false).
2225 15 - Concession: A and B are in tension but the text concedes/resolves it ("although A,
↳ still B", or a holding settles it).
2226 16 - Alternative: A and B are EXHAUSTIVE competing options -- EXACTLY ONE holds (they are
↳ mutually exclusive AND together cover the possibilities: not both, and not neither).
2227 ↳ E.g. "no one was harmed" vs "three people died"; "the cause was pilot error" vs "the
↳ cause was a metallurgical defect".
2228 17 - Disjunction: AT LEAST ONE of A and B holds (they may both hold, but the response rules
↳ out neither being true) -- e.g. two supporting findings at least one of which must
2229 ↳ be present.
2230 18 - Precedence/Succession: A and B are ordered in time with no truth dependence.
2231 19 - None: the response draws no logical or discourse dependence between A and B.

2232 20 3. Map the sense to a COUPLING, one of: entailment, contradiction, equivalence, exclusive,
2233 ↳ co_necessity, none.
2234 21 - Cause-Effect, Effect-Cause, Evidence, Condition, Instantiation -> entailment
2235 22 - Restatement -> equivalence
2236 23 - Contrast, Concession -> contradiction
2237 24 - Alternative -> exclusive (exactly one of A, B is true)
2238 25 - Disjunction -> co_necessity (at least one of A, B is true)
2239 26 - Precedence, Succession, None -> none
2240 27 Prefer "exclusive" over "contradiction" when the two claims are not just incompatible
2241 ↳ but EXHAUSTIVE (one of them must be true); prefer "contradiction" when they merely
2242 ↳ cannot both hold but could both be false.

2243 28 4. Give your final answer as two bracketed tags on ONE line, sense first:
2244 [sense=Cause-Effect] [coupling=entailment]
2245 A JSON object {"sense":"Cause-Effect","coupling":"entailment"} is also acceptable.

2246 29 Use the following examples to better understand your task.

2247 30 Example 1 (the response makes the causal link):
2248 RESPONSE: The company launched a flawed product last quarter. Reviewers panned it, returns
2249 ↳ spiked, and the company's stock price fell 15 percent over the same period.
2250 A: The company launched a flawed product last quarter.
2251 B: The company's stock price fell 15 percent last quarter.
2252 1. Reasoning: the response presents the flawed launch as the head of a chain (panning,
2253 ↳ returns) that ends in the stock decline, so the response itself draws a causal link
2254 ↳ from A to B.
2255 2. Discourse sense: Cause-Effect.
2256 3. Coupling: A causing B is a positive inferential link, i.e. entailment.
2257 4. Final answer:
2258 [sense=Cause-Effect] [coupling=entailment]

2259 31 Example 2 (both claims present, but the response draws NO connection -> None):
2260 RESPONSE: The quarterly report was published in April. Separately, the annual audit was
2261 ↳ scheduled for December. The two processes are run by different teams and were not
2262 ↳ related this year.
2263 A: The quarterly report was published in April.
2264 B: The annual audit was scheduled for December.
2265 1. Reasoning: both claims appear in the response, and one might imagine a
2266 ↳ reporting-to-audit link in general, but this response explicitly treats them as
2267 ↳ separate and unrelated. The response draws no dependence from A to B.
2268 2. Discourse sense: None.
2269 3. Coupling: no dependence the response asserts, i.e. none.
2270 4. Final answer:
2271 [sense=None] [coupling=none]

2272 32 Example 3 (the response states an EXHAUSTIVE alternative -> exclusive):
2273 RESPONSE: The official statement said no one was harmed in the incident. However, the
2274 ↳ coroner's report confirmed that three people died in the incident.
2275 A: No one was harmed in the incident.
2276 B: Three people died in the incident.
2277 1. Reasoning: the response sets A and B against each other ("However, ...") and they cannot
↳ both be true; but they also cannot both be false -- either people were harmed or they
↳ were not -- so exactly one holds. This is exhaustive, not a mere contrast. No holding
↳ resolves it.
2278 2. Discourse sense: Alternative.

2268 62 3. Coupling: exactly one of A, B is true, i.e. exclusive.
 2269 63 4. Final answer:
 2270 64 [sense=Alternative] [coupling=exclusive]
 2271 65
 2272 66 Example 4 (at least one must hold -> co_necessity):
 2273 67 RESPONSE: The defect was caught in review: at least one of the two independent checks --
 2274 68 ↳ the vibration analysis or the metallurgical assay -- flagged it.
 2275 69 A: The vibration analysis flagged the defect.
 2276 70 B: The metallurgical assay flagged the defect.
 2277 71 1. Reasoning: the response asserts the defect WAS caught by at least one check, so A and B
 2278 72 ↳ cannot both be false; but both could hold (both checks may have flagged it). This is a
 2279 73 ↳ disjunction, not an exclusion.
 2280 74 2. Discourse sense: Disjunction.
 2281 75 3. Coupling: at least one of A, B is true, i.e. co_necessity.
 2282 76 4. Final answer:
 2283 77 [sense=Disjunction] [coupling=co_necessity]
 2284 78
 2285 79 Your task:
 2286 80 RESPONSE: {{response}}
 2287 81 A: {{atom_a}}
 2288 82 B: {{atom_b}}

2284 A.2 PROMPT A, SECOND GENERATION — REBALANCED FOR PRECISION

2286 6,361 characters. Same placeholders. The five changes relative to the first generation, and their
 2287 motivation, are given in §4.5: few-shot balance from 3:1 related to 2:3; negative exemplars showing
 2288 silence rather than explicit disavowal; a verifiable evidence span; the discriminating question asked
 2289 first; and NONE given first position and equal prose. Transitivity is named as a non-relation.

2291 1
 2292 2
 2293 3 Instructions:
 2294 4 You are given the full RESPONSE a model produced, and two atomic claims, A and B, both
 2295 5 ↳ taken FROM THAT RESPONSE. Decide the discourse/logical relation FROM A TO B.
 2296 6
 2297 7 Most pairs of claims in a text are NOT related. They are simply near each other. Your
 2298 8 ↳ default answer is None, and you should depart from it only when the response itself
 2299 9 ↳ draws a dependence between A and B.
 2300 10
 2301 11 Three things that are NOT relations, however plausible they look:
 2302 12 - Topical adjacency. A and B being about the same subject, or sitting in neighbouring
 2303 13 ↳ sentences, is not a relation.
 2304 14 - World knowledge. If the link is true in general but the response does not draw it, the
 2305 15 ↳ answer is None.
 2306 16 - Transitivity. If the response links A to some third claim C, and C to B, that does NOT
 2307 17 ↳ make A related to B. Report only the links the text draws directly.
 2308 18
 2309 19 1. First ask the discriminating question: does the response ITSELF assert a dependence from
 2310 20 ↳ A to B -- as written, or as an explicit step in the author's argument? Identify the
 2311 21 ↳ words in the response that do it. If you cannot point to such words, the answer is None.
 2312 22
 2313 23 2. If and only if you can, name the DISCOURSE SENSE, one of: {{sense_menu}}.
 2314 24 - None: the response draws no dependence between A and B. This is the common case:
 2315 25 ↳ unrelated, merely adjacent, related only by world knowledge, or related only through
 2316 26 ↳ a third claim.
 2317 27 - Cause-Effect: the response says A causes/leads to B. Effect-Cause: A is the effect, B
 2318 28 ↳ its cause.
 2319 29 - Evidence: the response offers A as evidence/support for B.
 2320 30 - Restatement: A and B say the same thing -- a paraphrase, a converse, or a restatement
 2321 31 ↳ the text marks ("in other words", "equivalently", "that is").
 2322 32 - Contrast: A and B are opposed but NOT exhaustive (both could be false).
 2323 33 - Concession: A and B are in tension and the text concedes or resolves it ("although A,
 2324 34 ↳ still B"; a holding settles it).
 2325 35 - Alternative: A and B are EXHAUSTIVE competing options -- exactly one holds (not both,
 2326 36 ↳ not neither).
 2327 37 - Disjunction: AT LEAST ONE of A and B holds; both may.
 2328 38 - Precedence: A and B are ordered in time with no truth dependence.
 2329 39
 2330 40 3. Map the sense to a COUPLING, one of: entailment, contradiction, equivalence, exclusive,
 2331 41 ↳ co_necessity, none.
 2332 42 - Cause-Effect, Effect-Cause, Evidence -> entailment
 2333 43 - Restatement -> equivalence
 2334 44 - Contrast, Concession -> contradiction
 2335 45 - Alternative -> exclusive (exactly one of A, B is true)
 2336 46 - Disjunction -> co_necessity (at least one of A, B is true)
 2337 47 - Precedence, None -> none

2322 33 Prefer "exclusive" when the two claims are incompatible AND exhaustive; "contradiction"

2323 34 ↳ when they merely cannot both hold but could both be false.

2324 35 4. Answer on ONE line, with three bracketed tags. The evidence must be a SHORT VERBATIM

2325 36 ↳ QUOTE COPIED FROM THE RESPONSE -- the words that draw the link. Do not quote claim A or

2326 37 ↳ claim B; quote the response. Use [evidence=none] when the sense is None.

2327 38 [sense=Evidence] [coupling=entailment] [evidence="which is why the panel held"]

2328 39 Use the following examples to better understand your task.

2329 40 Example 1 (adjacent and on-topic, but the response draws no link -> None):

2330 41 RESPONSE: The survey was administered in March to a sample of 400 households. Response

2331 42 ↳ rates in rural districts have declined over the past decade. The analysis weighted

2332 43 ↳ results by district population.

2333 44 A: The survey was administered in March to a sample of 400 households.

2334 45 B: Response rates in rural districts have declined over the past decade.

2335 46 1. The two sentences are adjacent and both concern the survey, but the response never says

2336 47 ↳ the March administration bears on rural response rates. There are no words that draw a

2337 48 ↳ dependence.

2338 49 2. Discourse sense: None.

2339 50 3. Coupling: none.

2340 51 4. [sense=None] [coupling=none] [evidence=none]

2341 52 Example 2 (the response makes the causal link explicitly):

2342 53 RESPONSE: The company launched a flawed product last quarter. Reviewers panned it, returns

2343 54 ↳ spiked, and as a direct result the company's stock price fell 15 percent over the same

2344 55 ↳ period.

2345 56 A: The company launched a flawed product last quarter.

2346 57 B: The company's stock price fell 15 percent last quarter.

2347 58 1. The response presents the launch as the head of a chain ending in the stock decline, and

2348 59 ↳ marks it: "as a direct result".

2349 60 2. Discourse sense: Cause-Effect.

2350 61 3. Coupling: entailment.

2351 62 4. [sense=Cause-Effect] [coupling=entailment] [evidence="as a direct result"]

2352 63 Example 3 (true in general, but this response does not draw it -> None):

2353 64 RESPONSE: Higher interest rates raise the cost of corporate borrowing. The firm opened

2354 65 ↳ three distribution centres last year. Its logistics costs fell by eight percent.

2355 66 A: Higher interest rates raise the cost of corporate borrowing.

2356 67 B: The firm opened three distribution centres last year.

2357 68 1. One could argue borrowing costs affect expansion decisions, but that is world knowledge.

2358 69 ↳ This response states the two facts side by side and links neither to the other.

2359 70 2. Discourse sense: None.

2360 71 3. Coupling: none.

2361 72 4. [sense=None] [coupling=none] [evidence=none]

2362 73 Example 4 (the response states an exhaustive alternative -> exclusive):

2363 74 RESPONSE: The two readings are mutually exclusive and exactly one must be right: either no

2364 75 ↳ one was harmed in the incident, or three people died in it.

2365 76 A: No one was harmed in the incident.

2366 77 B: Three people died in the incident.

2367 78 1. The response sets A and B against each other and says exactly one must be right:

2368 79 ↳ "mutually exclusive and exactly one must be right: either ... or".

2369 80 2. Discourse sense: Alternative.

2370 81 3. Coupling: exclusive.

2371 82 4. [sense=Alternative] [coupling=exclusive] [evidence="mutually exclusive and exactly one

2372 83 ↳ must be right"]

2373 84 Example 5 (linked only through a third claim -> None):

2374 85 RESPONSE: The alloy is chemically identical to the certified reference. It therefore meets

2375 86 ↳ the reference specification. Meeting that specification is a precondition for airframe

2376 87 ↳ use.

2377 88 A: The alloy is chemically identical to the certified reference.

2378 89 B: Meeting that specification is a precondition for airframe use.

2379 90 1. The response links A to the middle claim (it meets the specification), and the middle

2380 91 ↳ claim to B. It does not link A to B directly; that would be transitivity, which is not

2381 92 ↳ a relation to report.

2382 93 2. Discourse sense: None.

2383 94 3. Coupling: none.

2384 95 4. [sense=None] [coupling=none] [evidence=none]

2385 96 Your task:

2386 97 RESPONSE: {{response}}

2387 98 A: {{atom_a}}

2388 99 B: {{atom_b}}

A.3 PROMPT B, DEFAULT — SURROGATE-TOKEN CONDITIONAL STRENGTH

2,182 characters. Placeholders: `{{response}}`, `{{atom_a}}`, `{{atom_b}}`, `{{coupling}}` (twice). This is the one prompt that does not end in a newline: it terminates with `Answer:` including the trailing space, so that the model’s next token is the surrogate whose logprobs are read. Reproducing it without that trailing space changes the tokenization of the answer position and therefore the estimate.

```

1
2
3 Instructions:
4 You are given the full RESPONSE, two claims A and B drawn from it, and a {{coupling}}
5   ↳ relation that holds from A to B. Assuming A is TRUE, judge -- IN THE CONTEXT OF THIS
6   ↳ RESPONSE -- whether the relation's implication about B is credible: at least plausible
7   ↳ / more likely than not, NOT whether it is certain.
8
9 - entailment or equivalence: given A and how the response uses it, is B at least plausibly
10  ↳ TRUE (more likely than not)?
11
12 - contradiction or exclusive: given A and how the response uses it, is B at least plausibly
13  ↳ FALSE (more likely than not)? (For "exclusive", A and B are exhaustive alternatives, so
14  ↳ A being true makes B false.)
15
16 - co_necessity: A and B are a pair of which at least one holds. Given the response rules
17  ↳ out "neither", is it at least plausible that B holds when A does NOT?
18
19 Answer with a SINGLE WORD, the very first word of your reply: Yes or No.
20 - Answer "Yes" if the implication is credible/plausible (even if not certain).
21 - Answer "No" only if the implication is implausible or the response does not actually
22  ↳ support it.
23 Do not output anything before the word Yes or No.
24
25 Example 1 (response supports a near-certain entailment):
26 RESPONSE: The new alloy is chemically identical to the certified reference alloy, so it
27   ↳ meets the certified reference specification.
28 A: The new alloy is chemically identical to the certified reference alloy.
29 B: The new alloy meets the certified reference specification.
30 coupling: entailment
31 Answer: Yes
32
33 Example 2 (weak but plausible, and the response draws the link -- still Yes):
34 RESPONSE: The company launched a flawed product last quarter, and its stock price fell 15
35   ↳ percent over the same period.
36 A: The company launched a flawed product last quarter.
37 B: The company's stock price fell 15 percent last quarter.
38 coupling: entailment
39 Answer: Yes
40
41 Example 3 (clear contradiction stated by the response):
42 RESPONSE: The official statement said no one was harmed, but three people died in the
43   ↳ incident.
44 A: No one was harmed in the incident.
45 B: Three people died in the incident.
46 coupling: contradiction
47 Answer: Yes
48
49 Your task:
50 RESPONSE: {{response}}
51 A: {{atom_a}}
52 B: {{atom_b}}
53 coupling: {{coupling}}
54 Answer:

```

A.4 PROMPT B, BASELINE — VERBALIZED CONDITIONAL STRENGTH

2,112 characters. Retained only for the comparison reported in §8; verbalized confidence is weakly calibrated (Xiong et al., 2024).

```

1
2
3 Instructions:
4 You are given the full RESPONSE and two claims A and B drawn from it. Assume claim A is
5   ↳ TRUE and, IN THE CONTEXT OF THIS RESPONSE, estimate how strongly A determines B under a
6   ↳ {{coupling}} relation from A to B.
7
8 - For an entailment or equivalence coupling: how likely is B to be TRUE given A?

```

2430 7 - For a contradiction or exclusive coupling: how likely is B to be FALSE given A?
 2431 ↪ (exclusive = A and B are exhaustive alternatives, so A true forces B false.)
 2432 8 - For a co_necessity coupling (at least one of A, B holds): how likely is B to be TRUE when
 2433 ↪ A is FALSE?
 2434 9
 2435 10 Answer with a single probability in [0, 1] to two decimals, after one short sentence of
 2436 ↪ justification. A near-certain link is close to 1.00; a merely plausible link is around
 2437 ↪ 0.60-0.70. End your answer with the probability wrapped in brackets on its own, exactly
 2438 ↪ in the form: [p=0.NN]
 2439 11
 2440 12 Example 1:
 2441 13 RESPONSE: The new alloy is chemically identical to the certified reference alloy, so it
 2442 ↪ meets the certified reference specification.
 2443 14 A: The new alloy is chemically identical to the certified reference alloy.
 2444 15 B: The new alloy meets the certified reference specification.
 2445 16 coupling: entailment
 2446 17 Justification: chemical identity to a certified reference alloy almost guarantees the
 2447 ↪ specification is met, so B follows very strongly from A.
 2448 18 [p=0.95]
 2449 19
 2450 20 Example 2:
 2451 21 RESPONSE: The company launched a flawed product last quarter, and its stock price fell 15
 2452 ↪ percent over the same period.
 2453 22 A: The company launched a flawed product last quarter.
 2454 23 B: The company's stock price fell 15 percent last quarter.
 2455 24 coupling: entailment
 2456 25 Justification: a flawed product can plausibly drive a stock decline, but many other factors
 2457 ↪ affect price, so the link is only moderate.
 2458 26 [p=0.65]
 2459 27
 2460 28 Example 3:
 2461 29 RESPONSE: The official statement said no one was harmed, but three people died in the
 2462 ↪ incident.
 2463 30 A: No one was harmed in the incident.
 2464 31 B: Three people died in the incident.
 2465 32 coupling: contradiction
 2466 33 Justification: if no one was harmed, it is almost certain that B (three deaths) is false.
 2467 34 [p=0.93]
 2468 35
 2469 36 Your task:
 2470 37 RESPONSE: {{response}}
 2471 38 A: {{atom_a}}
 2472 39 B: {{atom_b}}
 2473 40 coupling: {{coupling}}

2461 A.5 THE TWO SENSE MENUS

2462 The full menu offers all twelve senses of Table 4 plus NONE:

2465 1 Cause-Effect, Effect-Cause, Evidence, Condition, Restatement, Instantiation, Contrast,
 2466 ↪ Concession, Alternative, Disjunction, Precedence, Succession, None

2468 The restricted menu offers only the senses the evaluation corpus labels, which is the admissibility
 2469 filter of §4.5 applied at elicitation time rather than after it:

2471 1 Cause-Effect, Effect-Cause, Evidence, Restatement, Contrast, Concession, Alternative,
 2472 ↪ Disjunction, Precedence, None

2473 Note that the restricted menu omits **Succession** as well as **Condition** and **Instantiation**, leaving
 2474 nine relation-bearing senses. **Succession** is dropped because it is the mirror of **Precedence** and
 2475 both compile to NONE, so offering both invites an arbitrary choice on a relation that produces no
 2476 factor either way. The second-generation Prompt A's body is written against this restricted menu;
 2477 when it is run with the full menu, three senses are offered with no gloss and no coupling mapping in
 2478 the prompt body, and we report which menu each arm used.

2480 B THE REIFIED COHERENCE NODE, WORKED

2481 To make Eq. 7 concrete, take the three-claim subnet consisting of the unresolved casualty conflict of
 2482 Example 4 together with one entailment feeding it: claims a_7, a_{10} joined by a conflict at $p = 0.93$,

and $a_4 \rightarrow a_7$ an entailment at $p = 0.65$, all priors $\frac{1}{2}$, with $\rho = \frac{1}{2}$. Writing out the eight-cell vote factor for each relation from Eq. 7 and enumerating the 2^4 worlds of (a_4, a_7, a_{10}, R) gives $\mathbb{P}(R=1) = 0.453$: below its prior, because both relations have violating worlds that carry appreciable mass, and the conflict’s violating world $(1, 1)$ carries the most.

Two properties of the construction are worth noting. The $R = 0$ branch is flat, so $R = 0$ asserts nothing and the node cannot be driven up merely by adding relations that happen to be satisfied. And the factor is a *noisy* AND rather than a hard one: a relation with $p_r < 1$ leaves some mass on its violating worlds even when $R = 1$, so a single strongly-violated relation reduces $\mathbb{P}(R=1)$ without pinning it to zero. The reason we do not adopt the readout despite this reasonable behaviour is Table 6: it dips at the concession variant for the reason of Prop. 3, and it requires the free parameter ρ , violating R4.

C CORRESPONDENCE WITH MARKOV LOGIC

The coherence MRF has a natural first-order encoding (Richardson & Domingos, 2006), with $\mathbb{P}(\mathbf{a}) = \frac{1}{Z} \exp(\sum_k w_k n_k(\mathbf{a}))$ for $n_k(\mathbf{a})$ the number of true groundings of formula F_k . Evidence predicates *Entail*(i, j), *Contradict*(i, j), *Equiv*(i, j) and *Resolves*(h, i, j) carry the mined graph; *Holds*(i) is the query predicate. Working out the correspondence is instructive in both directions.

C.1 THE NAIVE ENCODING IS A DIFFERENT MODEL

The obvious encoding gives each relation one weighted clause, e.g. $\text{Entail}(i, j) \wedge a_i \Rightarrow a_j$ with $w = \text{logit}(p)$. This is *not* equivalent to Definition 1. Its induced pairwise factor is $[e^w, e^w, 1, e^w]$, which is not row-stochastic in the sense of Prop. 1: it rewards the vacuous $a_s = 0$ rows with e^w on *both* target states, so a clause is best satisfied by making its antecedent false. Brute force over the 2^{16} worlds of the recall report of §10.3 gives mean marginal support 0.455 under the naive MLN against 0.587 under the corresponding MRF, with per-claim marginals differing by up to 0.31. These are two different models, not two encodings of one.

The failure is sharpest at the mode. The all-false world satisfies every implication vacuously, so under the naive encoding the most probable configuration of any purely-implicational graph is the one in which the response asserts nothing — the maximum-a-posteriori reading of a response is that none of it holds. This is Prop. 2 appearing in a second guise, and it is the same lesson: an encoding that scores rule satisfaction rather than belief can be maximized by having nothing to satisfy.

C.2 A THREE-CLAUSE EXPANSION IS EXACTLY EQUIVALENT

Writing the log-potential in the general pairwise form $\ln \psi(a_s, a_t) = \alpha + \beta a_s + \gamma a_t + \delta a_s a_t$ and matching it against the with-priors tables of Table 2 gives, for π_s the source prior and $L(x) = \text{logit}(x)$:

τ	α (unit)	β (source)	γ (target)	δ (conjunction)
ENTAILMENT	$\ln(1 - \pi_s)$	$\ln \frac{1-p}{1-\pi_s}$	$L(\pi_s)$	$+L(p) - L(\pi_s)$
CONTRADICTION	$\ln(1 - \pi_s)$	$\ln \frac{p}{1-\pi_s}$	$L(\pi_s)$	$-L(p) - L(\pi_s)$
EQUIVALENCE	$\ln p$	$\ln \frac{1-p}{p}$	$\ln \frac{1-p}{p}$	$+2 L(p)$

At the atom prior $\pi_s = \frac{1}{2}$ the $L(\pi_s)$ terms vanish, and the table reduces to $\gamma = 0$ with $\delta = \pm L(p)$: in that setting the conjunction weight *is* exactly $\pm \text{logit}(p)$, the interaction the implication was meant to carry all along, and the unit and source clauses are the correction that makes each row a conditional. Away from $\frac{1}{2}$ the $L(\pi_s)$ terms are load-bearing, which matters for the two-stage composition of §6, where the priors are factuality posteriors rather than uniform. With this expansion the ground Markov logic network reproduces the MRF exactly: at $\pi_s = \frac{1}{2}$, mean support 0.5866 and $\log Z = -9.751$ for both, with zero per-claim difference to machine precision; the agreement is verified against the brute-force oracle at priors on and off $\frac{1}{2}$.

We establish the equivalence for the three original couplings only. EXCLUSIVE and CO_NECESSITY are not tabulated here, so the claim “the MLN encoding reproduces the MRF” should be read as scoped to the three-coupling fragment. That is a real limitation for the recall report of §10.3, whose

conflicts are EXCLUSIVE; the five-claim pair of §10.1 uses CONTRADICTION throughout and so falls inside the tabulated fragment.

C.3 WHICH VIEW TO KEEP, AND WHY BOTH

The reason to prefer the MRF for the measure itself is Remark 1: its pairwise factor family is closed at three parameters, so the inferential vocabulary cannot quietly grow as senses are added, and every relation type is guaranteed to be a 2×2 potential. The reason to keep the Markov logic view is that the rules our model cannot express are naturally written there — a transitivity rule $Entail(i, j) \wedge Entail(j, k) \Rightarrow Entail(i, k)$, a concession-cancellation rule $Resolves(h, i, j) \wedge a_h \Rightarrow \neg Penalize(i, j)$ replacing the fixed discount of Eq. 3, and a double-conflict rule — and their weights could be *learned* from labelled ladders rather than set by hand. That is the route past the two limitations §11 names, and it requires an MLN solver (MC-SAT for marginals, MaxWalkSAT for the mode) rather than the bucket-elimination stack used here.

D SERIALIZATION AND THE VARIABLE-ORDERING TRAP

One implementation detail is worth recording because it is silent when wrong. Networks are serialized in UAI format, whose variables are identified by integer index; the index order is the sort order of (cardinality, name), and with uniformly binary claim variables that reduces to *lexicographic order by name*. So a_{10} sorts before a_2 . Inference returns marginals by index, and mapping them back by any other convention — source order, numeric order — silently permutes marginals across claims. The score is unchanged, since Eq. 5 is a mean, but every per-claim diagnostic is then attributed to the wrong claim, and the “claim dragged below its prior” readout points at an innocent claim. The auxiliary variables of §5.2 are named with a prefix that sorts after all claim names for the same reason, so that adding them cannot shift the claim indices.