LoCoBench: A Benchmark of Coherence Ladders for Evaluating Logical Coherence Measures

Anonymous authors
Paper under double-blind review

Abstract

Measuring whether a long-form response *hangs together* — whether the claims it asserts form a consistent argument rather than a bag of individually checkable facts — requires evaluation data that no existing resource provides. Two documents that differ in coherence almost always differ in topic, length and content too, so a score difference between them cannot be attributed to coherence; human ratings supply a target but not a controlled contrast; and no resource supplies the *relation graph* that would let a measure's scoring errors be separated from its relation-extraction errors. We introduce **LoCoBench**, built from *coherence ladders*: families of five responses over one claim set that differ in one controlled structural respect, whose relative ordering is declared before any measurement, and whose every rung carries its own machine-readable gold relation graph. Four ladder types vary conflict structure, entailment-chain integrity and assertion order, and include a control whose rungs must all score *identically*. Families are generated by a five-stage prompted pipeline and admitted only after an independent cross-vendor committee, which excludes the generator, certifies that the prose realizes its planned relations, reads as expert prose, and asserts every planned claim. The released corpus comprises 100 families and 500 responses over 36 topics, with $4{,}761$ gold edges and gold data for six discourse senses — including exhaustive alternatives and resolved concessions — that no existing dataset labels. Scored from its gold graphs, a probabilistic coherence measure satisfies 92.0% of $1{,}408$ declared constraints, and both invariance ladders are exactly flat. Scored from graphs mined by two open models, the same measure satisfies 32.4–34.9% and fails almost every invariance check, and the best relation extractor reaches 54.6% coupling-level F1: the bottleneck is relation extraction, not scoring. A pilot human study finds that readers agree with the declared ordering on $8/10$ increase comparisons ($\alpha = 0.75$), and identifies one ladder type whose declared invariance readers reject; we specify a human validation study that keeps every annotator under three hours, release all prompts verbatim, and release every generation attempt with its outcome.

1 Introduction

Long-form factuality evaluation decomposes a response into atomic claims and verifies each against retrieved evidence (Min et al., 2023; Wei et al., 2024; Song et al., 2024). The paradigm cannot see whether the claims form an argument. A response asserting both that no one was harmed in an incident and that three people died contains one true claim and one false one, and per-claim verification reports exactly that, without registering that the response contradicts itself; a chain of true statements joined by a spurious "therefore" has no unsupported atom in it at all. Xie et al. (2026) quantify the omission: on logic-related samples, SAFE, VeriScore and FactCheck-GPT miss the logical dimension on 60%, 78% and 48% of cases respectively.

Measures of *logical coherence* that address this exist or can be built — from entity grids (Barzilay & Lapata, 2008) and their neural successors (Li & Jurafsky, 2017) through discourse-representation scores (Zhao et al., 2023) and reasoning-chain metrics (Golovneva et al., 2023) to LLM judges (Liu et al., 2023) and probabilistic models over claim-relation graphs. What does not exist is the data to evaluate them on. This paper builds it.

Why ordinary corpora will not do. The natural experiment — score a set of documents and correlate against human coherence ratings — has three compounding defects. *Confounding*: two documents that differ in coherence differ in topic, length, register and claim content as well, so a measure that tracks length or claim count can correlate respectably with ratings for reasons unrelated to coherence. *No ground truth for the mechanism*: a rating says a document is incoherent, not *which claims* are in tension, so a measure that is right for the wrong reason is indistinguishable from one that is right. *No separation of error sources*: a measure that builds an explicit relation graph both extracts relations and scores them, and a rating cannot say which of the two failed.

What LoCoBench provides instead. We build families of responses — *ladders* — over a single claim set, where the rungs differ only in a controlled structural respect and their relative ordering is *declared before any measurement*. A conflict ladder, for example, holds topic, claims and every unperturbed relation fixed while varying only the conflict structure. Because the rungs share a claim set, satisfying the declared ordering is attributable to coherence and not to content; and because two ladder types declare that every rung must score *identically*, the benchmark can catch a measure that tracks the fact of having been edited rather than coherence. Every rung ships with a gold relation graph — which claims are related, by which discourse sense, in which direction, at what strength, and whether the relation genuinely holds or is a planted error — so a measure can be run once on its own extracted relations and once on the gold graph, and the difference is the extractor's contribution, isolated by construction.

A documented gap in existing annotation. Six of the twelve discourse senses in our scheme — Instantiation, Restatement, Contrast, Concession, Alternative and Disjunction — have no gold data in any existing dataset we are aware of. Two of them are what adjudicative prose turns on: an Alternative is a pair of *exhaustive* competing claims of which exactly one holds, and a Concession is a tension the text itself resolves. A measure that conflates mere opposition with exhaustive opposition mis-scores every incident investigation and contested judgement, and it cannot be caught doing so without labelled examples of both. Supplying them is a primary purpose of the corpus.

Contributions. We make four contributions, each stated with the evidence offered for it.

1. **Coherence ladders as an evaluation design** (§3–§4). A formal definition of a ladder; a two-level annotation scheme whose discourse senses compile deterministically to truth-functional couplings; four ladder types that vary conflict structure, entailment-chain integrity and assertion order; and three classes of declared ordering constraint, including predicted inversions and exact invariances, which let the benchmark detect a measure that counts edits rather than measuring coherence.
2. **The LoCoBench corpus** (§7–§8). 100 families and 500 responses over 36 topics in four domains, with per-rung gold relation graphs, 1,928 planted invalid relations, declared non-relations as a genuine negative set, and gold data for six discourse senses no existing dataset labels.
3. **A reproducible generation and committee-validation pipeline** (§5–§6, Appendix A). Five prompted generation stages, eleven perturbation calls with declared effects on the gold graph, and three validators run by a cross-vendor committee that excludes the generator. Every acceptance threshold is in one table, every prompt — generation, validation, retry feedback, relation mining and judge baselines — is reproduced verbatim and checked against the code by a test, and every generation attempt is released with its outcome and the gate that rejected it.
4. **Validation of the benchmark, by measurement and by people** (§10–§11). Evidence that the ladders discriminate between relation extractors and between coherence measures; a localization of current measures' failures in relation extraction rather than in scoring; a pilot human study; and a pre-registered validation study, under three hours per annotator, testing label validity, edit fidelity and perceived ordering.

2 Related work

Factuality benchmarks and their blind spot. FActScore (Min et al., 2023) pairs biography prompts with a knowledge source and scores per-atom precision; SAFE (Wei et al., 2024) contributes LongFact-Objects and search-augmented verification; VeriScore (Song et al., 2024) restricts to

verifiable claims. Each supplies claims and evidence, and none supplies inter-claim structure, so none can serve as evaluation data for a coherence measure. Xie et al. (2026) is the closest prior art; it constructs data for *logic-related* fact-checking and labels dependencies between claims, but as a classification target rather than as a graded ordering over controlled variants.

Perturbation-based coherence evaluation. Evaluating coherence models against synthetic perturbations is long established: entity-grid work (Barzilay & Lapata, 2008) discriminates a document from its sentence permutations, and neural coherence models (Li & Jurafsky, 2017) add insertion and replacement perturbations. LoCoBench inherits the methodology and changes one thing: the perturbation is an operation on a *declared relation graph*, whose effect on the graph is computed and recorded per rung, rather than an operation on text whose effect on coherence is assumed. The benchmark can therefore say not only that rung $k{+}1$ should score above rung k but which relation changed, and — for meaning-preserving edits — that two rungs must score *identically*. Shuffling alone cannot express that distinction; indeed our pilot study (§11) suggests that the classic shuffle test measures something other than logical coherence.

Discourse annotation as precedent. Our sense inventory is grounded in PDTB 3.0 (Prasad et al., 2019) and RST (Mann & Thompson, 1988). Both establish that discourse relations can be annotated from surface cues, which is what makes the relation-recovery gate of §6 a reasonable demand. PDTB annotates naturally occurring text, where the distribution of relations is whatever the corpus happens to contain — which is why six of our senses lack gold data. A generated corpus can *require* the rare relations, and ours does.

LLM-generated benchmarks, LLM judges, and their validation. Generating evaluation data with a language model is now common, and its known risks are circularity (the generator's biases become ground truth), leakage (artifacts revealing the construction), and unverifiable labels. We address these structurally — a committee that excludes the generator, a leakage validator, and a gate requiring an independent model to *recover* the planned relations from the prose — and empirically, with human studies (§11) whose design follows standard practice for annotation reliability (Krippendorff, 2018) and paired-comparison modelling (Bradley & Terry, 1952). Justifying the committee as a stand-in for human annotators, by the alternative-annotator test (Calderon et al., 2025), is left to future work (§11.2).

3 THE ANNOTATION SCHEME

An item's gold annotation is a graph over its claims. Edge labels come from a two-level scheme: an interpretable *discourse sense* that an annotator can read off surface cues, and an inferential *coupling* that says what the relation asserts about the joint truth of its endpoints. A deterministic map connects the two levels, and that map — not the annotator and not the generator — is the authority on what a sense means.

3.1 LEVEL ONE: INFERENTIAL COUPLINGS

Definition 1 (Inferential couplings)**.** *A coupling τ is specified by the set $V(\tau) \subseteq \{0,1\}^2$ of joint truth assignments (a_s, a_t) of its source and target claims that it* penalizes*:*

Coupling	*Reading*	*Directed*	*Penalized set $V(\tau)$*
ENTAILMENT	$a_s \Rightarrow a_t$	*yes*	$\{(1,0)\}$
CONTRADICTION	$a_s \Rightarrow \neg a_t$	*yes*	$\{(1,1)\}$
EQUIVALENCE	$a_s \Leftrightarrow a_t$	*no*	$\{(0,1),(1,0)\}$
EXCLUSIVE	*exactly one of a_s, a_t holds*	*no*	$\{(0,0),(1,1)\}$
CO_NECESSITY	*at least one of a_s, a_t holds*	*no*	$\{(0,0)\}$
NONE	*no truth coupling*	—	$\varnothing$

The two couplings whose penalized set contains the both-true world, {CONTRADICTION, EXCLUSIVE}*, are the* conflict *couplings.*

Table 1: The compile map $\mathcal{C}$ of Def. 2, with the number of gold edges of each sense over all rungs of the released corpus. *Ordering-only* senses record assertion order but carry **no truth coupling**, so they produce no factor in any scorer that consumes the scheme. **Rev.** marks the senses whose edge can be reversed without changing any factor the scorer builds (§5.6); this is a property of the factor table, not of the *dir.* flag, so the symmetric sense Contrast does not qualify. **Bold** marks the six senses for which we are aware of no gold data in any existing dataset.

Level-2 sense	$\mathcal{C}$ coupling	prior	dir.	flag	rev.	gloss	edges
Cause-Effect	ENTAILMENT	$\perp$	yes			source brings target about	667
Effect-Cause	ENTAILMENT	$\perp$	yes			source is effect, target cause	51
Evidence	ENTAILMENT	$\perp$	yes			source is evidence for target	660
Condition	ENTAILMENT	$\perp$	yes			source conditions target	36
Instantiation	ENTAILMENT	$\perp$	yes			target instantiates source	70
Restatement	EQUIVALENCE	0.90	no		✓	definitional identity	560
Contrast	CONTRADICTION	$\perp$	no			opposed, *not* exhaustive	716
Concession	CONTRADICTION	$\perp$	yes	conc.		a tension the text resolves	500
Alternative	EXCLUSIVE	$\perp$	no		✓	exhaustive: exactly one holds	491
Disjunction	CO_NECESSITY	$\perp$	no		✓	at least one holds	500
Precedence	NONE	—	yes	ord.	✓	temporal order only	459
Succession	NONE	—	yes	ord.	✓	temporal order only	51
Total							4,761

The distinction between CONTRADICTION and EXCLUSIVE is the one the scheme exists to draw. CONTRADICTION forbids only the both-true world: two claims that cannot both hold but might both fail. EXCLUSIVE additionally forbids the both-false world, which is what an exhaustive pair asserts. CO_NECESSITY is the dual, forbidding only the both-false world, which is what a disjunction asserts. A scheme offering only CONTRADICTION can express neither, and a measure evaluated against such a scheme is untested on precisely the relations adjudicated disputes turn on.

3.2 Level two: discourse senses, and the compile map

Definition 2 (Compile map)**.** *The compile map $\mathcal{C}$ assigns to each discourse sense a coupling, a strength prior ($\perp$ meaning "use the estimated value"), and the flags* directed*,* concession *and* ordering-only*. It is total and deterministic, and is given in Table 1.*

Five senses collapse onto ENTAILMENT, a deliberate loss of information at the inferential level: whether one claim causes another or merely evidences it makes no difference to what the pair asserts about joint truth. The sense is retained alongside the coupling so that error analysis can distinguish "related, but mislabelled" from "not related", and so that the corpus serves work on the discourse layer for its own sake.

The compile map is enforced. A gold relation carries both a sense and a coupling, which invites inconsistency. The schema validator rejects any item in which coupling $\neq \mathcal{C}(\text{sense})$, and re-derives directedness, ordering-only status and exhaustiveness from the sense rather than trusting the generator's assertion of them. This matters for a generated corpus: an LLM asked for both a sense and a coupling will occasionally emit a pair no reading of the scheme licenses, and keeping such an edge would place a factor in the graph that its sense does not justify. Every gold edge in the corpus therefore uses one of the twelve $(\text{sense}, \text{coupling})$ combinations of Table 1.

4 Coherence ladders

Definition 3 (Coherence ladder)**.** *A* ladder *is a tuple $(\mathcal{P}, \langle \rho_0, \ldots, \rho_4 \rangle, \Gamma)$ where $\mathcal{P}$ is a base* relation plan *— a selected claim set with a declared relation graph over it — each rung ρ_k is a sequence of perturbation operator calls applied to $\mathcal{P}$, yielding a response text together with its* own *gold relation graph, and Γ is a set of constraints on the relative ordering of the rungs' scores. A ladder is* admitted *only as a whole (§6).*

Table 2: The four ladder types. Each holds the claim set, the topic and every unperturbed relation fixed, varying one structural dimension. n is the number of families in the released corpus, the generator's proportional allocation at 100 families. The ORDER and CONTROL patterns are declared with respect to *logical* coherence: their edits change no truth coupling, so the gold graph, and any score computed from it, must not move. Whether *readers* agree is an empirical question that the pilot study answers differently for the two types (§11).

Type	Dimension varied	Rungs, $\rho_0 \to \rho_4$	Declared pattern	n
CONFLICT	conflict structure	`worse`, `base`, `concession_resolved`, `fix_one_conflict`, `coherent`	increasing, with one predicted dip	46
CHAIN	entailment-chain integrity	`break_3`, `break_2`, `break_1`, `base`, `coherent`	increasing; $\log Z$ decreasing until `base`	21
ORDER	assertion order	`shuffle_full`, `shuffle_block`, `shuffle_adjacent`, `base`, `ordering_only`	**flat throughout**	21
CONTROL	nothing (meaning-preserving)	`ordering_only`, `direction_reversal`, `base`, `direction_reversal`, `ordering_only`	**flat throughout**	12

Whole-family admission. A ladder is admitted entire or not at all. Four of the five rungs carry no standalone claim, only a claim relative to their neighbours, so a ladder missing a rung is not a partial success but a broken comparison. Every gate therefore rejects at family granularity.

Table 2 gives the four types. CONFLICT is the reference ladder: only the conflict structure varies. A spurious relation makes the lowest rung worse than the base; the upper rungs resolve the conceded tension and then drop one and two conflict edges. CHAIN breaks $N \in \{3, 2, 1\}$ links of a chain of *valid* entailment relations that its plan is required to contain, so the defect is a *missing* relation rather than a present conflict — a distinction that matters because a measure built on contradiction detection alone will do well on the first ladder and fail the second. Its top rung drops one planted-invalid conflict.

ORDER and CONTROL are *invariance* ladders. ORDER permutes sentences at three granularities and then exchanges a Precedence for a Succession; CONTROL exchanges Precedence for Succession on its outer rungs and, on its inner rungs, reverses a *coupled* relation whose factor is symmetric in its endpoints (Restatement, Alternative or Disjunction; Table 1). None of these edits changes a factor, so the declared pattern is a flat line and a measure that trends across either has demonstrated that it responds to *having been edited*. The two types are kept separate on purpose. CONTROL edits are meaning-preserving for any reader. ORDER edits preserve every claim and every coupling but can strand an anaphor or put evidence after the conclusion it supports — they preserve *logical* coherence while degrading *discourse* coherence. Keeping the types apart lets the benchmark, and the human study, ask whether a measure (or a reader) can tell those two properties apart.

4.1 ORDERING CONSTRAINTS

A family declares its constraints in three classes, each asserted per *readout* of the measure under test, so that a family can declare that one readout must increase between two rungs while another is unconstrained there.

C1, strict increase. Consecutive rungs must be strictly ordered.

C2, predicted inversion or invariance. A score must *decrease*, or must be *exactly invariant*. This is what distinguishes a declared-ordering benchmark from a monotonicity test. Some readouts are expected to dip when a conceded tension is resolved, because softening a conflict raises the probability of the very configuration the conflict forbids; asserting the dip turns a known peculiarity into a falsifiable prediction. The same applies along CHAIN: restoring an entailment link adds a factor that concentrates the joint's mass, so the normalized log partition is predicted to *fall* at each breaking

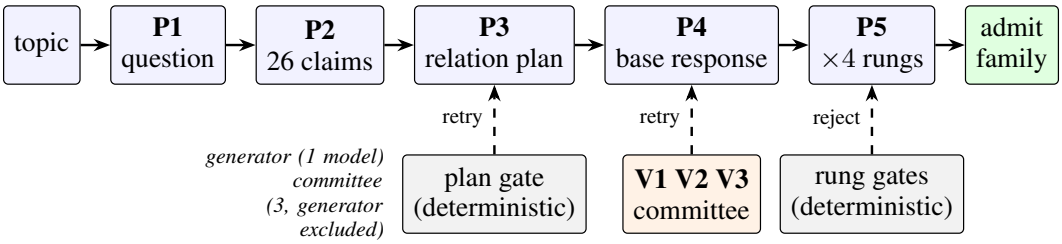


Figure 1: The pipeline. Five generation stages produce a family of five rungs from one topic. The plan is checked by a deterministic gate, the base response by the three-model committee, and each perturbed rung by deterministic gates; gate failures are fed back to the generating stage as a retry, and a family that exhausts its retries is rejected whole and stored with its reason.

step while the belief readouts rise. This declaration was revised on measurement, and §10 reports the evidence. And a coupling-free edit *must* be exactly flat: C2 then asserts an equality, and a measure that moves at all has been caught reading surface form.

C3, endpoint separation. The extreme rungs must separate (or, on the invariance ladders, must not), for all readouts. This is the robust headline claim, insensitive to the fine structure of intermediate rungs.

For the three-readout measure of §9, a CONFLICT family declares 13 constraints (9 C1, 1 C2, 3 C3) and every other family 15: CHAIN has 9 C1, 3 C2 and 3 C3, and ORDER and CONTROL have 12 C2 invariances and 3 C3. The released corpus declares 1,408, and a manifest's constraints are re-derived from the code and checked against it on release. A single-score measure inherits only the readout-independent content: five increase assertions per CONFLICT or CHAIN family (four C1 pairs and the C3 endpoints) and five invariance assertions per ORDER or CONTROL family. A measure is scored on the fraction of declared constraints it satisfies.

5 THE GENERATION PIPELINE

A family is produced by five prompted stages, P1–P5, whose outputs are checked before the family is admitted (Figure 1). Stages P1–P4 build a single *base* response: P1 turns a topic into a question, P2 writes a pool of candidate claims that bear on it, P3 selects claims from the pool and plans the relations between them, and P4 writes prose that realizes the plan. P5 then edits the base response, one operator call at a time, to produce the ladder's four other rungs. The stages are kept separate so that a defect can be attributed to one of them and repaired there. Every prompt is fixed and versioned and is reproduced verbatim in Appendix A, one subsection per prompt; the appendix files are written from the source by a script, and a test fails if any of them differs from the prompt the code issues. Algorithm 1 gives the procedure end to end. The subsections below take the stages in order, stating for each what it receives, what it must return, and why its instructions are as they are.

5.1 P1: TOPIC TO QUESTION

P1 (Appendix A.1) receives a topic from the inventory of §7 and returns one open-ended question. Three of its nine instructions do real work. The question must *elicit* diverse logical structure, and the prompt enumerates the target relations so that the question is chosen with the downstream plan in mind. It must prefer adjudicative framings — incident investigations, litigation, diagnostic reasoning, attribution disputes, competing scientific hypotheses — because these are where exhaustive alternatives and resolving holdings arise naturally. And it must *not* contain the words "cause", "condition", "alternative" or "contrast", so that the relations of the eventual answer are not primed lexically by the question.

5.2 P2: QUESTION TO TAGGED CLAIMS

P2 (Appendix A.2) receives the question and returns exactly 26 atomic claims, each a single declarative sentence with no connective joining it to another. Relations are deliberately left to the next stage,

Algorithm 1 Generating one family of type T on topic t. $\textsc{Ask}(P, \cdot, g)$ calls prompt P, parses the output, applies check g, and on a parse or check failure re-issues the prompt with the failure reason appended (Appendix A.9), up to 3 attempts at temperatures $(\text{default}, 0.3, 0.7)$.

$q \leftarrow \textsc{Ask}(\text{P1}, t)$ ▷ open question
$C \leftarrow \textsc{Ask}(\text{P2}, q)$ ▷ 26 tagged claims
$\mathcal{P} \leftarrow \textsc{Ask}(\text{P3}, q, C, T,\ g{=}\textsc{PlanGate}_T)$ ▷ claims, 10 relations, non-relations
$x_{\text{base}} \leftarrow \textsc{Ask}(\text{P4}, q, \mathcal{P},\ g{=}\textsc{Committee})$ ▷ V1, V2, V3
$E_{\text{base}} \leftarrow$ gold graph of $\mathcal{P}$, fields re-derived through $\mathcal{C}$
assign each operator call of each rung of T a distinct eligible target edge; **reject** if too few
for each non-base rung ρ_k of T, in ladder order **do**
  for each call o of ρ_k **do**
    $x \leftarrow \textsc{Ask}(\text{P5}, x, \mathcal{P}, o)$; **reject** unless x changed and drift $\leq 15\%$
  $E_k \leftarrow$ apply the declared edge effects of ρ_k's calls to E_{base}
reject if two adjacent rungs have equal edge sets and the upper one claims a change
reject if any two rungs have identical text but different edge sets
reject unless every item passes the schema validator; otherwise **admit** all five rungs

so that the claim set and the relation graph are controlled independently. The composition of the pool is fixed: 14 factually correct claims; 4 plausible-but-false claims; an *exhaustive alternative* pair, of which exactly one must be true; a *disjunctive* pair, of which at least one must be true; an *equivalent* pair; and two *holding* claims, in which an authority settles a dispute. Each claim carries a tag recording its role, and the parser checks the count and the tag composition. Two further instructions protect the scarce relations. The disjunctive pair must be phrased as states of the world rather than as findings, since a pair of *reports* is trivially compatible and yields no real disjunction; and the false claims must not be signalled by hedging, so that falsity cannot be read off the surface.

5.3 P3: CLAIMS TO RELATION PLAN

P3 (Appendix A.3) is the most constrained stage. It receives the question, the tagged claims and the ladder type, and returns a *relation plan*: 14–16 claims selected from the pool and renumbered to consecutive assertion positions; exactly ten relations over distinct claim pairs, each with a sense from Table 1, an intended strength band and a *validity* label; and a short list of declared non-relations. The prompt states the requirements below as exact targets. Two deterministic checks enforce them, in most cases with some tolerance. The *parser* checks the plan's structure: the claim, relation and non-relation counts; legal senses, with each coupling equal to $\mathcal{C}$ of its sense; an error kind on every invalid relation; at most one relation per claim pair; and no pair that is both a relation and a non-relation. The *plan gate* then checks the quotas. A failure of either is fed back to P3 as a retry (§5.7), and the thresholds are the P3 rows of Table 4. Algorithm 2 gives the plan gate in full. It evaluates every check rather than stopping at the first failure, so the retry note lists every violated quota at once. Each item below says what the prompt asks for and what is checked.

At least two false claims are selected, so the response asserts false content alongside true. Without this quota a planner can select none, and coherence and factuality would be perfectly correlated in the data. Two sits just below the ≈ 2.3 expected when 15 of 26 claims, 4 of them false, are selected indifferently. The plan gate checks this quota as stated.

Five senses are required: at least one Alternative, Disjunction, Restatement, Concession and temporal ordering (Precedence or Succession). This is what supplies the scarce senses, and the plan gate checks it as stated.

A ladder-specific quota reserves the edges that the ladder's own edits will consume. It is the only part of P3 that differs between ladders, and is substituted through a placeholder. For CONFLICT, ORDER and CONTROL, the prompt asks for at least four relations with a conflicting sense (Alternative, Concession or Contrast), at most one of them a resolved Concession. What the CONFLICT ladder actually needs is three *unresolved* conflict edges, enough to resolve one and drop two others, and that is what the gate checks. For CHAIN, the prompt asks for at least three valid entailment relations forming a chain, each link's target the next link's source; these are the links the ladder breaks. The gate checks that at least three valid relations have a sense that compiles to

Algorithm 2 The plan gate PLANGATE_T of Algorithm 1, applied to a plan $\mathcal{P}$ that has already passed the parser, for ladder type T. R is the plan's relation list, $\mathcal{C}$ the compile map, and F the set of claims P2 tagged as false; r.src and r.tgt are a relation's endpoint positions, and NORM normalizes whitespace, case and a final period. A non-empty list of failures is joined into the retry note sent back to P3.

fail $\leftarrow [\,]$
if $T =$ CHAIN **then** ▷ ladder-specific quota
  $L \leftarrow \{r \in R : \mathcal{C}(r.\text{sense}) = \text{ENTAILMENT} \wedge r.\text{validity} = \text{valid}\}$
  if $|L| < 3$ **then** append chain_links to *fail*
else ▷ CONFLICT, ORDER, CONTROL
  $K \leftarrow \{r \in R : \mathcal{C}(r.\text{sense}) \in \{\text{CONTRADICTION}, \text{EXCLUSIVE}\} \wedge \neg r.\text{resolved}\}$
  if $|K| < 3$ **then** append conflict_edges to *fail*
$S \leftarrow \{r.\text{sense} : r \in R\}$ ▷ required senses
$Q \leftarrow \{\text{Alternative}, \text{Disjunction}, \text{Restatement}, \text{Concession}\}$
if $Q \not\subseteq S$ **or** $\{\text{Precedence}, \text{Succession}\} \cap S = \varnothing$ **then**
  append rare_facets to *fail*
if $\left|\, |\{r \in R : r.\text{validity} = \text{valid}\}|/|R| - 0.55 \,\right| > 0.15$ **then** ▷ validity split
  append validity_split to *fail*
if $|\{a \in \mathcal{P}.\text{claims} : \text{NORM}(a.\text{text}) \in \text{NORM}(F)\}| < 2$ **then** ▷ false claims
  append factuality to *fail*
record $\{r \in R : |r.\text{src} - r.\text{tgt}| > 4 \wedge \neg \text{SHARESENTITY}(r)\}$ ▷ locality: logged only
return pass if *fail* is empty, else **retry** with the reasons in *fail*

ENTAILMENT, but not that they are linked end to end. The two quotas cannot be imposed on one plan: with ten relations, six of them valid, and five required senses, three valid entailments leave no room for four conflicts.

Exactly six relations are valid and four invalid. Each invalid relation carries an error kind — wrong_sense, wrong_direction, false_endpoint or spurious — and is asserted in the prose in the same register as the valid ones. The planted errors make the corpus a test of *precision*: an extractor that recovers every gold edge, planted ones included, has not exercised judgement. The prompt states the quota as a count, because a planner asked for "about 60%" rounds unpredictably. The plan gate is more tolerant: it accepts 8–12 relations with a valid fraction of 0.55 ± 0.15. Every released base plan has exactly ten relations, six of them valid (§7).

Four to six non-relations are declared: pairs of selected claims, within four positions of each other, between which the response must draw no dependence. Precision against all unlabelled pairs is close to meaningless, since most such pairs are simply unannotated; a violated *declared* non-relation is a genuine false positive against a real negative set. The parser checks the count and that no non-relation is also a planned relation; the distance of at most four positions is asked for but not checked.

Relations are local by default. Pairs within four positions are preferred, and long-range pairs are permitted when the two claims share a named entity. This biases the graphs toward ones a reader can recover, so that a failure of the relation-recovery gate (§6.1) indicates a defective realization rather than an unreasonable demand. The plan gate records the relations that fall outside this window but does not reject a plan for them.

5.4 P4: PLAN TO RESPONSE

P4 (Appendix A.4) receives the question and the plan and returns 550–650 words of prose (floor 500) that realizes the plan exactly: it asserts every planned claim once, in position order, and realizes every planned relation so that a careful reader can recover it from the text alone. The response must at the same time read as publishable expert prose rather than as an annotation walked through sentence by sentence, and most of P4's instructions resolve that tension:

- a back-reference uses a pronoun, or a noun phrase named after *the thing the claim is about* rather than after its status as a finding;

Table 3: The six perturbation operators and the eleven parameterized calls realizing them. Each call has a declared effect on the rung's gold relation graph and a declared set of eligible targets (§5.6). The two O6 calls, and direction_reversal on a reversal-invariant relation, change no factor: they are what the ORDER and CONTROL ladders exist to check.

Op	Calls	Effect on the text
O1 relabel	wrong_sense, direction_reversal, exhaustiveness_flip	rewrite the connective so the text draws a different sense, reverses the relation, or flips Contrast↔Alternative
O2 add conflict	inject_contradiction, spurious_relation	add one sentence directly contradicting a named claim at least two sentences later; or assert a link between two claims the response kept separate
O3 break chain	break_chain	remove or negate the link realizing a valid entailment so the chain no longer connects, keeping both endpoints
O4 drop relation	drop_relation	remove the connective realizing an edge, leaving the two claims as adjacent assertions
O5 resolve	add_resolution, remove_resolution	add or delete the holding sentence that settles a conceded tension
O6 reorder	shuffle_order, ordering_only	permute sentences at full, block or adjacent granularity; or exchange Precedence↔Succession preserving chronology — *no coupling changes*

- quantifiers, numbers, dates, polarity and modals are *content*, so a claim may be reworded but not weakened;
- exhaustive and disjunctive pairs must make their logical strength *explicit*, since a bare "X, or Y" leaves a reader unable to tell Alternative from Disjunction;
- a temporal ordering is stated within one sentence;
- non-relations are kept apart by distance and paragraphing rather than by an explicit marker, which would both leak the plan and itself draw a relation;
- the response must not mention the plan, its labels, or the fact that some of its content is incorrect.

The base response is the one output the committee validates (§6); a committee rejection is fed back to P4 as a retry.

5.5 P5: RESPONSE TO RUNG

P5 (Appendix A.5) receives the current text, the plan and one *operator call* with its target, and returns the edited text together with a structured diff. It must apply exactly that one perturbation, changing the minimum number of words and preserving everything else verbatim. The perturbations are organized as six *operators*, realized by eleven parameterized calls (Table 3); a rung is a fixed sequence of calls, applied one after another. Each call has a declared effect on the gold relation graph, which is how the rung's own graph is obtained (§5.6). Two deterministic gates check each call's output: the text must have changed, and its length may drift by at most 15%, with one sentence's further allowance for the four calls that add or remove a sentence by definition.

5.6 DERIVING EACH RUNG'S GOLD RELATION GRAPH

A rung's gold graph is the base plan's graph with that rung's operator calls applied to it. Four mechanisms make this well defined.

Target assignment. Each call in a rung is assigned a distinct eligible edge, so a rung composing two drop_relation calls drops two *different* edges; a plan with too few eligible edges is rejected rather than allowed to emit a rung that duplicates its neighbour. Eligibility depends on the call. break_chain targets only *valid* entailment edges, since breaking a planted-invalid one would remove an error and make the rung more coherent than the base. direction_reversal targets only senses whose factor table is symmetric in its endpoints, so that the reversal changes no factor. This set is *derived* from the scorer's tables rather than listed by hand, so it cannot drift from them, and a test reverses one edge of

each sense through the scorer and checks that the score moves exactly when the sense is ineligible. The derived set excludes Contrast: the sense is symmetric, but its CONTRADICTION table with priors carries the source's prior on one row. Coupled senses are preferred over ordering-only ones, so that CONTROL's two edits never land on the same edge.

Re-derivation of derived fields. When a call changes an edge's sense, the edge's coupling, directedness, ordering-only status and exhaustiveness are re-derived through $\mathcal{C}$.

An adjacency gate. A family is rejected when two *adjacent* rungs have the same edge set while the upper rung's calls claim a change. Adjacency, not parentage, is the right comparison: every CONFLICT rung declares the base as its parent, so a check against the parent alone would pass a rung that duplicates the rung immediately below it.

A text–label gate. Over *all* pairs of rungs, two rungs with byte-identical prose must carry identical gold edge sets; otherwise at least one rung's labels describe an edit its text does not contain.

5.7 THE FAILURE POLICY

What happens when a stage's output is defective is part of the specification, and there are two responses. A **retry** re-issues the prompt with the failure reason appended, using one of two fixed notes (Appendix A.9): a *format* note after a parse failure, and a *content* note after a gate failure, which tells the model that its format was correct and its content must change. The two are kept distinct because a model told only to fix its format tends to resubmit the same content. Retries are issued at rising temperature, since a byte-identical prompt re-sent at temperature zero cannot produce a different output. Parse failures at every stage, plan-gate failures and committee failures are retried in this way. A **reject** applies when a family cannot carry its ranking claim — retries exhausted, too few eligible targets, a rung that did not change or drifted too far, or a degenerate pair of rungs — and stores the family with the failing gate, its threshold, the observed value and the raw output. No stage patches a defective artifact into an acceptable one. One further guard concerns length: P4's target exceeds the default completion-token limit of some gateways, and a truncated response fails every parser for the wrong stated reason, so the output limit is raised explicitly.

6 COMMITTEE VALIDATION

The base response of every family is checked by a committee of three models before the family is admitted. Each committee member runs three validators, each a fixed prompt reproduced in Appendix A:

- **V1, relation recovery** (§6.1, Appendix A.6): can a reader who has not seen the plan recover the planned relations from the prose?
- **V2, prose audit** (§6.2, Appendix A.7): does the response read as expert prose, without leaking the plan or hedging its claims?
- **V3, claim coverage** (§6.3, Appendix A.8): does the response assert every planned claim?

A family is admitted only if its base response passes all three, its plan passed the plan gate, and every rung passed the deterministic gates of §5. Table 4 collects every acceptance threshold in the pipeline. The rest of this section describes who sits on the committee, how their votes are combined, and what each validator asks.

Who votes. Any committee member that shares the generating model's identifier is removed before voting, and the configuration is refused at load time unless at least three voters remain after that exclusion and they span at least three model families, so that agreement cannot be an artifact of shared training. The released corpus was generated by `claude-opus-5-5` and validated by `gpt-5.6`, `gemini-3.1-pro` and `claude-sonnet-5-5`. The generator and one committee member therefore share a vendor; we return to this dependence in §12.

How votes are combined. The three validators measure different things, so each has its own decision rule (rightmost column of Table 4). **V1 is any-of**: it certifies that a careful reader *can* recover the plan, so one rater who clears both thresholds is sufficient evidence; every rater's rates are recorded

Table 4: Every acceptance threshold, defined in one table in the source and nowhere else. V1 rates are computed over the *valid* planned relations only. The rightmost column is the committee decision rule.

Gate	Quantity	Threshold	Rule
V1	valid planned relations recovered, correct coupling	≥ 0.80	any-of
V1	… and with correct sense	≥ 0.70	any-of
V2	fluency, formality, organization (each, of 5)	≥ 4	majority/facet
V2	leakage and hedging span lists	empty	majority/facet
V3	claim coverage over gating statuses	$= 1.00$	majority/claim
V3	gating statuses	missing, merged	—
V3	… and on a planted-false claim	+altered	majority/claim
P3	claims selected (parser)	14–16	—
P3	relations planned (parser; prompt asks for 10)	8–12	—
P3	declared non-relations (parser)	4–6	—
P3	fraction valid (prompt asks for $6/10$)	0.55 ± 0.15	—
P3	false claims selected	≥ 2	—
P3	each required sense, temporal ordering by either sense	≥ 1	—
P3	unresolved conflict relations (CONFLICT, ORDER, CONTROL)	≥ 3	—
P3	valid relations compiling to ENTAILMENT (CHAIN)	≥ 3	—
P3	relation locality (recorded, not enforced)	≤ 4 positions	—
P5	length drift	$\leq 15\%$	—
Panel	voters after generator exclusion	≥ 3	—
Panel	distinct model families	≥ 3	—

alongside the decisive one. **V2 is a majority per facet**: a prose defect counts only if a majority of raters report it, so a single unusually strict auditor cannot reject a sound response. **V3 is a majority per claim**: a claim counts as lost only if a majority of raters call it lost.

What the committee sees. The committee runs on the base response only. The four perturbed rungs are checked by the deterministic gates of §5.5 and §5.6 — the text changed, its length drift is within bound, adjacent rungs differ in their edges, identical texts carry identical labels, and every item satisfies the schema — and their gold graphs are derived from the base graph by the declared edge effects rather than re-certified from their prose. That each P5 edit realizes its declared effect, and nothing more, is therefore an assumption of the pipeline, which the validation study tests (§11.2, Claim B).

6.1 V1: RELATION RECOVERY — THE LOAD-BEARING GATE

V1 (Appendix A.6) gives a committee member the response and its numbered claims — but *not* the plan — and asks for the relations the response itself draws, each with a sense, a coupling, a strength band and a resolved flag. Admission requires that at least 0.80 of the *valid* planned relations be recovered with the correct coupling, and at least 0.70 with the correct sense. This is the gate that gives the labels their meaning: the rest of the pipeline produces a plan, and V1 certifies that *the prose realizes it*. Without it, a relation that the planner intended and the writer failed to express would remain in the gold graph, and an extractor would be penalized for correctly not finding it.

The rates are computed over valid relations only. Planted-invalid relations are deliberately wrong, so requiring their recovery would reward exactly the over-assertion the corpus is built to detect; it would also make the gate unsatisfiable, since at most 6 of the 10 planned relations are valid. With six valid relations, both thresholds require five of six, so a single relation decides admission, and a coupling threshold above 0.834 would silently demand all six. An admitted item may therefore carry at most one valid gold relation that the decisive rater did not recover; the validation study's claim-pair check (§11.2) measures how often a human reader also fails to see it.

Recovery is graded at three levels. A recovered relation is *identified* by its unordered claim pair, independently of its sense, and its direction, coupling and sense are graded separately. Keying identity on the sense instead would place a planned Concession$(7, 12)$ and a recovered Contrast$(7, 12)$ —

same endpoints, both CONTRADICTION — in different key spaces, and would score "recovered with the wrong sense" as "not recovered", which is precisely the distinction V1 exists to measure.

6.2 V2: PROSE AUDIT

V2 (Appendix A.7) scores the response for fluency, formality and organization, each from 1 to 5 with a floor of 4, and collects three lists of spans: *leakage*, *hedging* and *artifacts*. The leakage and hedging lists must be empty; artifacts are recorded but do not gate. Leakage is defined narrowly: naming a relation type or sense, naming the plan or its labels, referring to "claims" or "atoms" as the material the text was assembled from, or any bracketed marker or index. It explicitly does *not* cover the connectives that *realize* relations — "either X or Y", "at least one of these must hold" — because P4 is required to use them, and flagging them would reject correctly written responses. Hedging is a closed list of seven words. Every reported span must be quoted verbatim from the response, so that a hallucinated objection can be detected and discarded.

6.3 V3: CLAIM COVERAGE

V3 (Appendix A.8) asks, for each planned claim, whether the response asserts it, and returns asserted, altered, missing or merged with a quoted span. Coverage must be perfect over the *gating* statuses, missing and merged. altered does not gate in general: it is the judgement most sensitive to the rewording P4 licenses, and a permitted rewording is easily read as an alteration. The exception is a planted-*false* claim, on which a majority verdict of altered also rejects. There the characteristic alteration is the generator correcting the claim to the truth, which would leave a `factual: false` label on prose asserting the opposite. The prompt also states that two claims joined by a connective are *not* merged, since "either X or Y" asserts both; otherwise every realized Alternative would count as a merge.

6.4 ADMISSION, AND WHAT THE COMMITTEE DOES NOT ESTABLISH

Admission. All 100 released families were admitted. The generation log records 251 family attempts for the 100 slots, and the release includes every one with its outcome. Of these, 67 were rejected by a gate: the prose audit in 44 cases (43 on scores, 1 on a leakage span that all three raters quoted), an unchanged or over-long P5 rung in 20, and relation recovery in 3. The remaining attempts were either admitted or regenerated from an earlier stage and superseded by a later attempt.

How decisive the decision rules are. Each admitted family records every rater's verdict. On relation recovery, each rater alone would have admitted 89–91 of the 100 families, and all three independently clear both thresholds on 77; the any-of rule is therefore decisive on 23 families and redundant on the rest. On the prose audit the raters are not interchangeable: `claude-sonnet-5-5` scored some facet below the floor on 36 admitted families, and each of the other two raters on 3. The majority rule is what keeps one strict rater from deciding admission alone.

What admission does not show. A committee whose members can share biases with the generator, and with each other, can certify an item that a person would reject, and no admission rate can reveal that. The committee's gates are therefore themselves an object of validation. Testing a gate's agreement with human judgement needs responses that the gate rejected, which the release does not keep, so the validation study defers that test (§11.2).

7 CORPUS COMPOSITION

Table 5 gives the inventory. Topics and ladder types are assigned to family slots deterministically, so a resumed run assigns the same topic and type to the same slot. Ladder types are allocated in proportion to Table 2. Topics are cycled through the inventory, so that at 100 families every one of the 36 topics contributes at least two families and 28 contribute three; the generator also implements a floor-then-surplus policy (three families per topic, surplus to the five adjudicated topics) that applies from 108 families upward.

Table 5: The topic inventory: 36 topics across four domains. The five **bold** topics are the adjudicated subjects where exhaustive alternatives and resolving holdings occur naturally, and receive the surplus families at corpus sizes above the per-topic floor.

Domain	Topics	n
Natural Sciences	Astronomy, Biology, Botany, Chemistry, Geology, Human Biology, **Medicine**, Meteorology, Oceanography, Physics, Zoology	11
Humanities	Archaeology, Art, Culture, **Ethics**, Gender Studies, History, Linguistics, Literature, Philosophy, Religion	10
Engineering & Technology	Artificial Intelligence, **Civil Engineering**, Computer Science, Cybersecurity, Electrical Engineering, Mechanical Engineering, Renewable Energy Technologies, Robotics	8
Social Sciences	Anthropology, Economics, Education, **Law**, **Political Science**, Psychology, Sociology	7

The released corpus. The release comprises 100 families and 500 responses generated by `claude-opus-5-5`: 46 CONFLICT, 21 CHAIN, 21 ORDER and 12 CONTROL, over all 36 topics in all four domains (31 natural sciences, 29 humanities, 22 engineering and technology, 18 social sciences). Each base plan selects 15–16 claims, 2–3 of them false, and plans 10 relations (6 valid, 4 invalid) and 5 declared non-relations; responses run 567–731 words, mean 643. Across all rungs there are 4,761 gold edges, 4,251 of them edge-producing; 2,833 are valid and 1,928 planted-invalid (768 false_endpoint, 460 spurious, 365 wrong_sense, 335 wrong_direction). Every one of the six senses with no prior gold data is represented by at least 70 edges (Table 1). Every family carries five distinct response texts except the CONTROL families, whose ladder is symmetric by design — rungs 0 and 4 apply the same edit, as do rungs 1 and 3 — so each pair carries identical labels and usually identical text: 11 of the 12 have three distinct texts, and one has five, its paired rungs differing only in wording.

8 ITEM SCHEMA AND RELEASE FORMAT

Table 6 gives the schema. The release is four artifacts per corpus: `items.jsonl`, one item per line; `families.json`, the manifest carrying each family's declared ordering constraints; `state.json`, the generation log, including each admitted family's committee verdict; and a `rejected/` directory holding every family that failed a gate with the gate, threshold, observed value and raw output. The rejected set is part of the release because admission rates are auditable only if the denominator is available. Several fields reveal the declared answer (`rung_name`, `expected.perturbation`, the item-id suffix, `factual`); any human or model evaluation must be run on redacted copies, and the release includes the redaction script used in §11.

9 THE EVALUATION HARNESS

Arms: holding the scorer fixed. An *arm* determines where the relation graph comes from while holding the claims, priors and scoring procedure identical. GOLD takes each item's labelled relations; GOLD_VALID takes only those labelled valid, which quantifies what the planted errors cost; a MINED arm, written `mined:`$\langle$`model`$\rangle$`:`$\langle$`policy`$\rangle$, runs a relation extractor on the response text. Because the scorer is identical across arms, any shortfall of a mined arm below the gold arm is attributable to extraction. The gold arm needs no LLM, so it is deterministic and serves as a regression test on the corpus itself.

What a scorer must expose. The harness requires only that a measure map a response (and, for graph-based measures, a relation graph over its claims) to one or more scalar readouts. To validate the corpus we use a Markov random field over claim truth variables: one binary variable per claim with unary factor $[1-\pi_i, \pi_i]$, and one pairwise factor per relation penalizing the assignments $V(\tau)$ of Def. 1. Three readouts are graded: the mean posterior marginal; a consistency probability combining the probability that no conflict is jointly active with the mass of upheld support relations; and the partition function normalized between a conflict-free ceiling and the base network's most probable

Table 6: The item schema. Every field is machine-readable and validated on load; fields marked *derived* are re-computed from the sense through $\mathcal{C}$ rather than trusted as written.

Field	Type	Meaning
Item level		
`id`, `name`	string	stable identifier and human-readable label
`source`	string	generating model
`response`	string	the prose
`num_atoms`	int	claim count
`notes`	string	free-text provenance note
Per claim (`atoms`)		
`id`, `text`	string	claim identifier and its exact text
`label`	string	P2 tag: correct, incorrect, alt-pair, disj-pair, equiv-pair, holding
`factual`	bool	whether the claim is true — *false* for planted-false claims
`role`	string	structural role in the plan
`pos`	int	1-based assertion position
Per relation (`relations`)		
`id`	string	edge identifier, referenced by perturbation targets
`source_id`, `target_id`	string	endpoint claims, ordered
`level2_sense`	string	one of the twelve senses
`level1_coupling`	string	*derived*: $\mathcal{C}$(sense)
`directed`	bool	*derived*: whether endpoint order is meaningful
`ordering_only`	bool	*derived*: carries no truth coupling, produces no factor
`exhaustive`	bool	*derived*: true iff the coupling is EXCLUSIVE
`intended_strength_band`	string	strong, moderate, or weak
`strength_range`	[float, float]	the band's numeric interval
`validity`	string	valid or invalid
`error_kind`	string|null	for invalid edges: wrong_sense, wrong_direction, false_endpoint, spurious
`is_concession`	bool	whether the relation is a conceded tension
`is_resolved_concession`	bool	whether the text resolves it
`resolver_atom_id`	string|null	the holding claim that resolves it
`position_distance`	int	signed target-minus-source position gap
`window_admission`	string	whether the pair falls inside the locality window
`provenance`	object	which stage planned or transformed the edge
Negative set and expectations		
`non_relations`	list	declared-unrelated claim pairs with their position distance
`expected.family_id`, `expected.family`	string	ladder membership and type
`expected.rung_index`, `expected.rung_name`	int, string	position in the ladder
`expected.perturbation`	object	the calls applied, their targets, their parent rung, and the declared `edge_effects`
`meta`	object	canonical topic, domain, framing, split, word count, generator

world. Gold graphs are scored with $\pi_i = 0.9$ for factual claims and 0.1 otherwise. These readouts have *different* declared expectations under C2: the consistency readout is expected to dip when a conceded tension is resolved, and the manifest asserts that dip.

Extraction metrics. Mined relations are matched to gold at three nested levels: **pair** (were the claims related at all), **coupling** (the Level-1 coupling, with direction for the asymmetric ones), and **sense**. Coupling is the level to report, since it is the level at which a scorer selects a factor.

Table 7: Declared ordering constraints satisfied by the MRF measure scored from the corpus's own gold graphs, by ladder and constraint class. GOLD uses every gold edge; GOLD_VALID discards the planted-invalid ones. A dash marks a class the ladder does not declare.

Arm	Ladder	C1	C2	C3	total (%)
GOLD	CONFLICT	339/414	45/46	119/138	503/598 (84.1)
	CHAIN	174/189	62/63	62/63	298/315 (94.6)
	ORDER	—	252/252	63/63	315/315 (100)
	CONTROL	—	144/144	36/36	180/180 (100)
	all	513/603	503/505	280/300	**1296/1408 (92.0)**
GOLD_VALID	CONFLICT				253/598 (42.3)
	CHAIN				270/315 (85.7)
	all	277/603	503/505	238/300	1018/1408 (72.3)

Recall is stratified by direction in claim order, by directedness, by locality and by validity; on the invalid subset it reads inverted, since those relations are planted errors. Precision is reported against the declared non-relations separately from precision against all unlabelled pairs, and edges outside the scheme's $(\text{sense}, \text{coupling})$ combinations are counted as unscorable. The ordering checker evaluates every declared constraint and records the observed direction and magnitude, so a near-miss is distinguishable from an inversion.

10 VALIDATING THE BENCHMARK BY MEASUREMENT

This section asks whether the released corpus behaves as an evaluation instrument should. It first checks that the corpus is internally consistent, by scoring it from its own gold graphs (§10.1). It then shows that it separates relation extractors (§10.2), that it locates where a graph-based measure fails (§10.3), and that it separates coherence measures of different kinds (§10.4). Every number is on the released 100-family corpus: 500 responses and 1,408 declared constraints.

Setup. The measure under test throughout is the MRF of §9, with its three readouts. It is scored under three kinds of arm. The GOLD and GOLD_VALID arms need no LLM and are deterministic. The MINED arms take relations from two open extractors, `llama-3.3-70b-instruct` and `gpt-oss-120b`, each under two candidate-pair policies (prompts in Appendix B): *windowed*, which takes forward pairs within four positions refined by what the response links, and *bidirectional*, which takes both directions of every pair within the window. Mining is sampled once per configuration; repeated runs of an identical configuration moved coupling F1 by up to 0.06, so smaller differences between mined arms are not interpretable. The single-score baselines of §10.4 are the DiscoScore intrinsic metrics (Zhao et al., 2023), an entity-grid score, ROSCOE self-consistency (Golovneva et al., 2023), an NLI contradiction score, a G-Eval coherence judge with the SummEval rubric (Liu et al., 2023; Fabbri et al., 2021), a direct coherence judge, and two controls, response length and claim count.

10.1 THE GOLD GRAPHS SATISFY THE DECLARED ORDERINGS

Scored from its gold graphs, the measure satisfies 92.0% of the declared constraints (Table 7). Because this arm involves no model, it is a deterministic check that each ladder's rungs relate to one another as declared, and it is rerun as a regression test whenever the corpus changes. Three results stand out.

Both invariance ladders are exactly flat. All 495 invariance and endpoint assertions on ORDER and CONTROL hold, which certifies that their edits change no factor.

CHAIN *holds* 94.6% *and* CONFLICT 84.1%*, and the failures are concentrated.* On CONFLICT, the largest block of failures is C1 increases on the mean-marginal readout at the two upper steps (53 of the 95 failures), and 60 of the 95 are changes smaller than 10^{-3}: near-ties rather than inversions. A benchmark on which the gold arm scored 100% would have constraints too weak to be informative.

Table 8: Relation extraction against gold, over the 4,251 edge-producing gold relations, micro-averaged. **Edges** is the number of relations mined; **Pair R** the recall of related claim pairs; **Coup** precision, recall and F1 at the coupling level. **Fwd/Bwd** split coupling recall by relation direction in claim order. **Non-rel** counts violations of the 2,454 declared non-relations.

Extractor	Policy	Edges	Pair R	Coup P	Coup R	Coup F1	Fwd	Bwd	Non-rel
`llama-3.3-70b`	windowed	6,587	71.9	36.2	56.1	44.0	63.4	8.5	75
`llama-3.3-70b`	bidirectional	22,448	89.4	16.9	75.3	27.6	77.1	63.6	775
`gpt-oss-120b`	windowed	6,045	74.8	46.5	66.1	**54.6**	74.4	12.9	61
`gpt-oss-120b`	bidirectional	19,197	95.4	24.2	88.8	38.1	90.2	80.3	600

Table 9: Declared ordering constraints satisfied by the MRF measure under each mined arm, by ladder type (1,408 assertions). The scorer is the one of Table 7; only the relation graph differs.

Arm	CONFLICT	CHAIN	ORDER	CONTROL	total (%)
GOLD (Table 7)	503/598	298/315	315/315	180/180	1296/1408 (**92.0**)
MINED, `llama`, windowed	303/598	174/315	2/315	1/180	480/1408 (34.1)
MINED, `llama`, bidirectional	303/598	154/315	0/315	0/180	457/1408 (32.5)
MINED, `gpt-oss`, windowed	308/598	169/315	13/315	2/180	492/1408 (34.9)
MINED, `gpt-oss`, bidirectional	290/598	158/315	5/315	3/180	456/1408 (32.4)

The planted errors carry the CONFLICT *signal.* GOLD_VALID falls to 72.3%, almost entirely on CONFLICT (42.3%). The upper CONFLICT rungs improve coherence by removing planted-invalid conflicts, and an arm that has already discarded those edges cannot see the improvement.

Two log-partition declarations were revised on measurement. We report both because a declaration revised on the data it scores is weaker evidence than one made in advance. First, along CHAIN's three breaking steps the log-partition readout fell on every one of the first 16 families generated, while the mean-marginal readout rose on all of them. The fall follows from the readout's construction (§4.1), and we moved those three pairs per family from C1 to a C2 predicted decrease. Second, at CONFLICT's `concession_resolved` step the log partition had been declared exactly invariant, because that edit only reweights a factor; on the release it failed on all 46 families, as a small decrease on 40 and a small increase on 6. This is not a uniform reversal, so instead of predicting a direction we left the pair unconstrained, and the corpus declares 1,408 assertions rather than 1,454. Since that pair failed on every family, the revision raises the gold arm from 89.1% to 92.0% without changing a single passing assertion. The mean-marginal and consistency declarations, and every endpoint separation, are as first declared.

10.2 THE LADDERS DISCRIMINATE BETWEEN RELATION EXTRACTORS

The corpus spreads the four extraction configurations over 27.0 points of coupling F1, from 27.6 to 54.6 (Table 8). Three of its findings are properties of the corpus rather than facts about one extractor.

The declared non-relations work as a negative set. Violations range from 61 to 775 of 2,454 and rank the four configurations in exactly the order of their coupling precision, while being interpretable in a way that precision against unlabelled pairs is not.

Locality metadata exposes a structural limitation. Backward-running gold relations are recovered at 9–13% under the windowed policy and 64–80% under the bidirectional one, because a forward-only candidate set cannot emit a backward directed arc. Without each relation's signed position distance, this would be indistinguishable from a failure of the model.

Recall is cheap and precision is not. The bidirectional policy raises coupling recall by 19–23 points, but more than triples the number of mined edges and multiplies non-relation violations by ten.

Table 10: Single-score measures on the readout-independent constraints. **Increase**: the 335 increase assertions of the 67 CONFLICT and CHAIN families (adjacent C1 pairs and the C3 endpoint pair). **Invariance**: the 165 rung pairs of the 33 ORDER and CONTROL families, satisfied when the score does not move. Only the two judges call the backend model, so every other row is the same in both column pairs. The mined MRF row reports, per backend, the mined arm and readout with the most increase assertions.

	`gpt-oss-120b`		`llama-3.3-70b`	
Measure	Increase	Invariance	Increase	Invariance
MRF, GOLD, mean marginal	273/335 (81.5)	165/165 (100)	(no LLM)	
MRF, GOLD, consistency	242/335 (72.2)	165/165 (100)	(no LLM)	
MRF, GOLD, log partition[†]	233/335 (69.6)	165/165 (100)	(no LLM)	
MRF, mined (best arm)	175 (52.2)	4 (2.4)	188 (56.1)	1 (0.6)
DiscoScore RC	199 (59.4)	110 (66.7)	199 (59.4)	110 (66.7)
DiscoScore LC	186 (55.5)	112 (67.9)	186 (55.5)	112 (67.9)
G-Eval judge	123 (36.7)	62 (37.6)	11 (3.3)	133 (80.6)
Direct judge	115 (34.3)	48 (29.1)	32 (9.6)	113 (68.5)
Entity grid	3 (0.9)	153 (92.7)	3 (0.9)	153 (92.7)
NLI contradiction	0 (0.0)	165 (100)	0 (0.0)	165 (100)
ROSCOE self-consistency (max, mean)	0 (0.0)	165 (100)	0 (0.0)	165 (100)
Response length (control)	140 (41.8)	110 (66.7)	140 (41.8)	110 (66.7)
Claim count (control)	0 (0.0)	165 (100)	0 (0.0)	165 (100)

[†]The Increase column holds CHAIN's three breaking steps to an increase, where C2 predicts a decrease for this readout (§10.1); that accounts for most of the row's shortfall.

10.3 THE LADDERS LOCATE THE BOTTLENECK IN EXTRACTION

The same scorer satisfies 92.0% of the constraints from gold graphs and 32.4–34.9% from mined graphs (Table 9). Since the arms differ only in where the relation graph comes from, the gap is the contribution of extraction, isolated by construction: the bottleneck is relation extraction, not scoring. The four mined arms lie within 2.5 points of each other, inside what a single mining sample can resolve, so we do not rank them on this table.

The breakdown by ladder shows *how* mined graphs fail. They satisfy about half of the CONFLICT and CHAIN constraints, but almost none of the invariance constraints: at most 15 of the 495 ORDER and CONTROL assertions, against all 495 from gold graphs, and each that passes does so only because two independently mined graphs happened to score identically. Re-mining a reordered or re-directed response returns a slightly different graph, so the score moves. The invariance ladders therefore expose extractor *instability under meaning-preserving edits*, which a benchmark that tests only for increases could not detect.

10.4 THE LADDERS DISCRIMINATE BETWEEN COHERENCE MEASURES

Table 10 scores single-score measures on the constraints that do not depend on a readout. A measure must do well on *both* columns, because on either column alone a degenerate measure wins: the claim-count control is perfect on invariance by never moving, and a judge that moves on every edit scores well on increase.

Only the gold-graph measure does well on both. Scored from gold graphs, the MRF has the highest increase score in the table (69.6–81.5%, depending on the readout) and perfect invariance.

No evidence-free measure comes close. The DiscoScore cohesion metrics do best, at 55–59% on increase, and are the only measures above the response-length control (41.8%). They measure noun repetition, which the CONFLICT and CHAIN edits happen to change, and on invariance they are no better than length.

The LLM judges fall below the length control on increase, and fail in opposite directions with the two backends. With `gpt-oss-120b` they move on most edits, meaning-preserving or not (34–37%

Table 11: Which study tests which claim, and how.

Claim	Tested by	Measurement
A	validation study, core sessions	claim-pair checks (A1–A3) and claims-as-stated checks (A4)
B	validation study, core sessions	per-version check for collateral change (B)
C	pilot; validation study	pilot: whole-text pairwise judgements; study: per-version judgements (H1–H3) and the instruction test (H4)
D	—	deferred (§11.2)

increase, 29–38% invariance); with `llama-3.3-70b` they barely move (3–10% increase, 69–81% invariance).

Measures that read only the claims cannot see a ladder. Every rung of a family carries the same claims; only the prose, and the relations it draws, differ. The NLI contradiction score and ROSCOE self-consistency read only the claims, so they receive the same input at every rung and, like the claim-count control, return one value per family: 0% on increase and, trivially, 100% on invariance. Giving them the response as NLI context would let them see the ladder, but that is a different measure from the one the literature reports.

Extraction costs the MRF its invariance rather than its sensitivity. Scored from mined graphs, the MRF reaches 52–56% on increase, level with the cohesion metrics, but fails invariance almost entirely (0.6–2.4%), the instability identified in §10.3.

11 Validating the benchmark with people

The measurements of §10 show that the benchmark discriminates. They cannot show that its labels are correct, because every label and every acceptance decision is model output. The benchmark rests on four claims that could reasonably be doubted:

A. Label validity. The gold graph of a base response describes that response: its valid relations are drawn by the text, its declared non-relations are not, and its planted-false claims are asserted as written.

B. Edit fidelity. Each rung's edit does what its record says, and nothing else.

C. Ordering validity. Readers perceive the declared orderings, and the invariances where they are declared.

D. Committee validity. The committee's gates agree with human judgement well enough to stand in for it.

This section reports a pilot study, which tested Claim C, and specifies a validation study that tests Claims A, B and C on the release (Table 11). Claim D is out of scope for reasons given in §11.2. In both studies every screen is built from redacted items, with every answer-revealing field removed and the side of each pair randomized; the redaction is asserted against the serialized screens rather than assumed. The annotator briefs and interfaces are in Appendix D.

11.1 Pilot study

Design. The pilot was run on a preliminary 14-family corpus built with an earlier version of the pipeline, before the release was generated. Four volunteer annotators judged 14 pairwise screens drawn from seven of its families (Appendix D.3). The families were chosen where the gold-graph measure and the strongest evidence-free baseline disagree most. For CONFLICT and CHAIN, the screens compared the hardest adjacent step (ρ_0, ρ_1) and the endpoints (ρ_0, ρ_4); one ORDER and one CONTROL family contributed two screens each. Each screen showed two whole responses and asked which is more *logically coherent* — whether everything it says can hold together — with an explicit *about equal* option, a confidence rating and a one-sentence reason. The brief (Appendix D.1) stated that length, fluency, vocabulary, topic knowledge and sentence order in itself are not the target.

Table 12: Pilot study: 4 annotators, 14 screens, 56 ratings. Agreement is Krippendorff's α (nominal) over the choices $\{A, B, \text{equal}\}$. Two of the ten increase screens had no majority.

Screen kind	Screens	Ratings	Outcome
Increase (CONFLICT, CHAIN)	10	40	majority matches declared on $8/10$; ratings $32/40$
Invariance, CONTROL	2	8	"equal" on $6/8$ ratings
Invariance, ORDER	2	8	"equal" on $0/8$; unshuffled preferred on $8/8$
Agreement	14	56	$\alpha = 0.75$ overall, 0.74 increase, 0.67 invariance

Readers agree with the declared increases. On the increase screens, 32 of 40 ratings matched the declared ordering (one-sided binomial test against a fair coin, $p < 10^{-4}$), and every annotator individually matched it on at least 7 of 10 (Table 12). Readers resolved both kinds of pair, adjacent and endpoint. The one screen on which all four answered *equal* was an adjacent CONFLICT pair (ρ_0, ρ_1), whose spurious relation none of them noticed; that step is where the benchmark's subtlest contrast is expected. Of the 28 reasons on increase screens that could be checked against the perturbation record, 13 named a claim that the edit had changed.

Readers accept one invariance and reject the other. The two invariance ladders came out differently (Fisher's exact test on the 2×2 table, $p = 0.007$). Readers called the CONTROL rungs equal, as declared. They unanimously preferred the unshuffled ORDER rung, and their reasons cite stranded references ("cites *that deficit* before any deficit has been introduced") — a failure of *discourse* coherence, which the brief told them to disregard but which evidently cannot be disregarded. The result has two readings. Either readers cannot separate logical from discourse coherence once the latter is degraded, in which case ORDER's invariance is a property of logical coherence that human judgement does not directly measure; or the shuffle genuinely damages the argument, in which case the declared pattern is wrong. The validation study's instruction test is designed to choose between them. Until it does, we report ORDER results separately from CONTROL and never pool them as "invariance".

Limits. The pilot used seven families, two screens per family, four annotators and a single instruction, and tested neither Claim A nor Claim B. The validation study addresses all of these except Claim D.

11.2 VALIDATION STUDY

Design at a glance. The study has two parts, run by different annotators (Appendix D.2, Figure 2):

- *Core sessions* test Claims A, B and C on 22 CONFLICT, CHAIN and CONTROL families. Each annotator reads a base response once, then answers every question about its family while the text is fresh: the claim pairs, the claims as stated, and each edited version shown as a highlighted change.
- *An instruction test* on 7 ORDER families decides between the two readings of the pilot's ORDER result, by varying the question readers are asked.

Two measurements of the release make the core design possible. First, an edited rung is a small change to its base: a CONFLICT, CHAIN or CONTROL rung differs from its parent by a median of 37 words (90th percentile 75), against 643 words for a whole response. A reader who has just read the base can therefore judge each rung from its highlighted change without rereading it. Second, 62 of the 63 ORDER shuffle rungs are exact sentence permutations of their base, which a script certifies without a reader. (The exception, rung 0 of family f063, rewrites two sentences and is checked by hand.) Reading every response once in this way costs about 31 annotator-hours in total; a design with a separate study per claim, reading the same responses once for labels, once for edits and once for pairwise comparison, would cost about 550 at the release's scale. No annotator works for more than three hours.

Sample. 29 families are drawn per ladder type with a fixed seed: 11 CONFLICT and 6 CHAIN, in proportion to the release; 5 CONTROL, oversampled because the ladder is small and Hypothesis H3

depends on it; and 7 ORDER, used only in the instruction test. The 22 core base responses carry 132 valid and 88 planted-invalid gold relations, 110 declared non-relations and 45 planted-false claims. Every sense the scheme exists to separate is represented — Contrast 33, Alternative 23, Disjunction 22, Concession 23 — but Instantiation (5) and Condition (3) are too rare to estimate alone, and are reported pooled with the other senses of their coupling.

Core sessions. Seven annotators with a linguistics or NLP background each judge 7 or 8 CONFLICT or CHAIN families and 1 to 3 CONTROL families. Each of the 22 families is judged by exactly three annotators, and no annotator judges a family twice. Families on the same topic are spread across annotators; the three Botany families force two annotators to hold two of them, which leaks nothing, since different families pose different questions over different claims. Annotators never see validity labels, rung names, or which version is the base. For each family they work through four steps:

1. *Read* the base response, with its claims numbered in the margin.
2. *Check* 15 *claim pairs* (Claim A): the 10 gold relations and 5 declared non-relations, shuffled and unlabelled. For each pair: does the response draw a link between the two claims (yes / no)? And, whatever the answer, could both claims be true together, and could both be false together? The two truth questions recover the coupling from first principles rather than from a sense name — a "yes" to both-false rules out EXCLUSIVE and CO_NECESSITY, and a "yes" to both-true rules out CONTRADICTION and EXCLUSIVE — so a reader who does not know the sense inventory is still scored correctly.
3. *Check the claims as stated* (Claim A): for each numbered claim, does the response assert it as written? The planted-false claims are among them, unmarked, so this checks that the response asserts its false claims rather than silently correcting them. It takes about a minute, because the claims were just read.
4. *Judge each edited version* (Claims B and C) — four for CONFLICT and CHAIN, two for CONTROL — shown as the base with its change highlighted. Is the response now more coherent, less coherent, or about as coherent? And did the change alter anything other than a link between claims (none / minor / material, with a quoted span)?

Each session begins with the brief and one practice family, excluded from analysis, whose gold answers are adjudicated in advance.

Instruction test. Choosing between the two readings of the pilot's ORDER result needs whole texts but no label work, so it runs as a separate block with three further annotators — educated readers, screened for reading comprehension. Each judges 21 whole-text pairs: the three shuffled rungs (full, block and adjacent) of each of the 7 ORDER families, each against its base. (The fourth non-base ORDER rung relabels a temporal relation without moving any sentence, so it is not a shuffle and is left out.) Half of each annotator's screens ask which response is more *logically consistent* ("if you believed everything each response says, which would leave you believing something consistent?"), and half ask which *hangs together better as an argument*. Varying the instruction within rather than between annotators makes each annotator their own control, so differences in strictness cancel. The instructions are assigned by a cyclic Latin design across the three annotators, so every pair is judged under both instructions and pair difficulty cannot masquerade as an instruction effect; this yields 31 and 32 judgements per instruction. Which side holds the original is balanced exactly within each annotator, and no two consecutive screens come from the same family.

Timing and quality control. At 200 words per minute, 27 seconds per claim pair and 45 seconds per edited version, a CONFLICT or CHAIN family takes about 15 minutes and a CONTROL family about 13. Core sessions are modelled at 156–169 minutes including a ten-minute break, and an instruction-test session at 171 minutes, at about seven minutes per pair of whole responses. These are estimates, not measurements, and the pilot cannot calibrate them, because its annotation tool logged the time a task was open rather than the time spent on it. Two annotators are therefore timed on the practice family before the study runs; if it takes more than 17 minutes, an eighth core annotator is added and the same 66 annotator–family slots are spread more thinly, which keeps three annotators on every family. A session that reaches three hours stops at a family boundary, and its remaining families move to another annotator. Each session also contains two attention checks at no extra cost

Table 13: Pre-registered hypotheses of the validation study. *Judged* counts judgements over all three raters; *Units* counts the items each test is run on, one majority per item; *Precision* is the 95% Wilson half-width (in percentage points) on those units at the expected rate, or the number of units the test needs for power 0.8.

Hypothesis	Judged	Units	Precision	Pass criterion
A1. Valid gold relations are confirmed	396	132	±5.2	≥ 90% at the coupling level
A2. Declared non-relations are confirmed	330	110	±7.4	≥ 80%
A3. How often planted-invalid relations read as valid	264	88	±8.3	descriptive, by error kind
A4. Planted-false claims are asserted	135	45	±7.1	≥ 95%, every exception listed
B. Edits realize their declared effect only	234	78	±6.8	≥ 90%, no material collateral change
H1. Readers agree with the declared increase	204	68	needs 23	majority above chance, one-sided binomial
H2. Agreement is lowest at CONFLICT's first step	33	11	—	descriptive, with its interval
H3. CONTROL versions read as equal	30	10	—	descriptive
H4. The instruction moves ORDER toward equal	63	∼32/arm	needs 32/arm	higher "equal" rate under logical consistency

— one claim pair repeats a claim, and one edited version restores its base word for word — and annotators who fail either are excluded before analysis.

Hypotheses and analysis. Table 13 lists the hypotheses, which are registered before data collection together with the analysis models, exclusion rules and pass criteria. Each test runs on majorities, one per item, so its sample is a third of the judgements.

- *Rates (A1, A2, A4, B)* are reported with 95% Wilson intervals, 5.2 to 7.4 points wide in each direction at the expected rates. That distinguishes a true rate of 80% from the 90% threshold, but not one of 85%.
- *A3 has no threshold*, because it is not a test the labels can fail. Planted-invalid relations are asserted in the prose like valid ones (§5.3), so a careful reader links them. The truth questions can expose a planted wrong_sense, whose coupling differs from the label, but not a wrong_direction, false_endpoint or spurious link, whose coupling reads correctly. A3 therefore reports, per error kind, how often a planted error reads as valid, which measures how convincingly the errors were written.
- *H1* needs 23 items for power 0.8 at $\alpha = 0.05$, at the pilot's measured agreement of 0.75, and has 68. *H4* needs 32 judgements per instruction for power 0.8 to detect a rate of "about equal" of 0.40 against 0.10, and has 31 to 32; this is why the test draws 7 ORDER families rather than 5, which would give power 0.65. *H2 and H3* are too small to test and are reported descriptively.
- *Agreement* is Krippendorff's α per question (Krippendorff, 2018), with $\alpha \geq 0.67$ on the coupling as the reliability target.
- *Orderings*: rung strengths per family are fitted with a Bradley–Terry model with ties (Bradley & Terry, 1952) and compared with the declared order by Kendall's τ, and every measure of §10 is rescored against the human majority as well as against the declared answer.

A rung whose edit is judged unrealized, or whose collateral change is judged material, has a gold graph that does not describe its text. Such rungs are flagged in the release, and every result is re-reported without their families.

Both outcomes of H4 have consequences fixed in advance. If H4 holds, the ORDER ladder is validated as an invariance test of *logical* coherence, and the benchmark documents that human "coherence" ratings conflate logical and discourse coherence. If it fails, ORDER is re-designated a

discourse-coherence probe: its declared pattern becomes increasing toward the unshuffled base, it is removed from the invariance score, and the change is reported. Either way, ORDER and CONTROL are never pooled.

Scope. Claim D, that the committee's gates agree with human judgement, is not tested. Testing an admission gate needs responses the gate rejected, and the release keeps the rejection records but only two rejected responses; the rest were overwritten by the retries that replaced them. A test of Claim D therefore needs a separate generation run with the gates logged but not enforced. That would validate the generator rather than the released labels, so we leave it to future work. Until then, the committee's verdicts are evidence that the prose realizes its plan, not a substitute for the human checks above.

Cost and procedure. The core sessions total about 19 annotator-hours, the instruction test about 9, and adjudicating the practice families about 3: about 31 hours in all, paid at or above the local minimum wage for the time worked. Screens are exported from the redacted release by the script of §8, which adds a highlighted-change view to the existing pairwise screen; its leak checks apply unchanged. The pilot's four annotators do not take part. [Ethics review: obtain and report the institutional determination before data collection.]

12 LIMITATIONS

The corpus is LLM-generated, including its labels. Questions, claims, plans, responses, perturbations and gold labels are model output, validated by a model committee. The structural mitigations of §6 address the most obvious circularity, but not a bias shared between generator and committee; current frontier models are not independent in the relevant sense, and in our release one committee member shares the generator's vendor. Until the validation study has tested Claims A–C, and Claim D has been tested, the gold graphs should be treated as model-certified rather than human-verified.

Perturbed rungs are validated by construction. The committee checks base responses; rung gold graphs are derived by declared edge effects, and whether each P5 edit realizes its effect and nothing more is checked only by deterministic gates until the validation study's edit check (§11.2) is complete.

Declared orderings are premises, and two were revised. A ladder's constraints encode our judgement that, for example, resolving a conceded tension increases coherence. The pilot shows that at least one declared pattern — ORDER invariance — is not what readers perceive. Two log-partition declarations were revised after measuring the families they govern (§10.1): CHAIN's breaking steps, from an increase to a predicted decrease, and CONFLICT's `concession_resolved` step, from invariance to unconstrained, after it failed on all 46 families. The second revision only removes a prediction, so a measure's score rises if it failed that pair more often than it passed it. Every arm in this paper did: the gold arm loses 46 failures and no passes, and the mined arms lose at most 3 passes each. A declaration revised on the data it scores is weaker evidence than one made in advance, and we report the unrevised figure (89.1% of 1,454) alongside the revised one. The validation study tests the premises (Hypotheses H1–H4 of §11.2), and we fix in advance how its outcomes change them.

Scale. One hundred families is enough to establish that the instrument discriminates, and not enough to rank closely matched measures. Mining is single-sample, with measured run-to-run variation of up to 0.06 coupling F1. The four mined arms' ordering satisfaction spans only 2.5 points (32.4–34.9%, Table 9), and a single sample per arm cannot rank arms that close, so we do not rank them on it; the gap that matters, 92.0% from gold graphs against at most 34.9% mined, is more than twenty times that spread. CONTROL is the smallest ladder type by design (12 families), and the type a reader should most want enlarged.

Genre, length, and provenance. Every item is 567–731 words of formal expository prose with 15–16 claims, and the locality bias of P3 under-represents long-range dependencies; conclusions do not transfer automatically to dialogue, summarization, much longer arguments, or informal registers. The release comes from a single generator and the mined arms from two extraction models, so generator-specific artifacts are controlled only by the V2 leakage gate. The model gateway caches

identical requests, so a repeated temperature-zero run through it can return a cached response rather than a new sample; one release family takes its question and claims from such a cached response (its plan and prose were generated afresh).

13 CONCLUSION

Evaluating a coherence measure requires contrasts that vary coherence while holding content fixed. LoCoBench supplies them as ladders: families of responses over one claim set, with declared orderings, per-rung gold relation graphs, a labelled negative set, and invariance ladders on which a sound measure must not move. The construction attributes a score difference to coherence rather than to content, and — by scoring a measure once on gold graphs and once on mined ones — isolates extraction error from scoring error by construction. On the released 100-family corpus the gold-graph measure satisfies 92.0% of $1{,}408$ declared constraints, with both invariance ladders exactly flat, and at most 34.9% from mined graphs, which are unstable under meaning-preserving edits: the bottleneck is extraction. Only two cohesion metrics beat the response-length control on increase, LLM judges fall below it with both backends, and measures that read only the claims cannot see a ladder at all.

Because the gold arm needs no model, it also serves as a deterministic regression test of the corpus itself. A pilot study finds that readers confirm the declared orderings and the meaning-preserving control, and reject the invariance of sentence shuffling; the validation study we specify will establish how far the labels and the edits can be trusted.

AI USE STATEMENT

This is a benchmark constructed *by* generative AI, and the disclosure is correspondingly extensive. Every artifact in the corpus — questions, claims, relation plans, response prose, perturbed rungs and gold relation labels — is the output of a large language model driven by the prompts of Appendix A, and every acceptance decision is made by a committee of three further language models. §11 and §12 state what this does and does not establish. We also used AI assistants for code implementation and editorial assistance on this manuscript. All LLM-generated code was reviewed by the authors and is covered by a test suite that, among other things, asserts every prompt in the appendix byte-for-byte against the code. The benchmark design — the ladder construction, the constraint classes, the operator inventory, the gate thresholds and their calibration, the schema, the metrics and the human-study protocol — is the authors' own. We take responsibility for the final content of this work.

ETHICS STATEMENT

Intended use. LoCoBench is for evaluating measures of logical coherence in long-form generated text. It is not a factuality benchmark: a response can be perfectly coherent and entirely false. It is not a training corpus for text generation.

Deliberately false content. Every item asserts at least two factually false claims, and every plan contains four deliberately invalid relations. This is essential to the design and makes the corpus hazardous if mishandled. False claims are labelled at the claim level and invalid relations at the edge level, and **the corpus must not be redistributed with those labels stripped**. Prose from these items should not be treated as a source of factual information.

Human participants. The pilot study of §11.1 involved four volunteer annotators who gave informed consent, judged generated text only, and provided no personal data beyond an anonymous identifier. The validation study of §11.2 will recruit ten paid annotators under an institutional ethics determination obtained before data collection, with payment at or above the local minimum wage for the time worked, and no session longer than three hours. No item concerns a real case or an identifiable individual.

Compute. Generating one family takes several dozen model calls across generation and validation, and a mined arm over the corpus several thousand more. This is modest but not free, and a benchmark that requires an LLM committee to extend is less accessible than one that does not.

Reproducibility statement

Appendix A reproduces every prompt the pipeline, the committee, the retry mechanism, the relation miner and the judge baselines issue; the files are written from the source by a script and a test fails if they diverge. Appendix C gives the prompt specification and substitution contract. Table 4 gives every acceptance threshold, Table 3 the perturbation operators, Table 6 the item schema, Algorithm 1 the procedure, and §5.6 the derivation of each rung's gold graph. §6 names the generating and committee models. The release comprises the items, the family manifests with their declared constraints, the generation log with each admitted family's committee verdict, and the rejected families with their reasons. The gold-arm evaluation requires no LLM access and is deterministic.

References

Regina Barzilay and Mirella Lapata. Modeling local coherence: An entity-based approach. *Computational Linguistics*, 34(1):1–34, 2008.

Ralph Allan Bradley and Milton E. Terry. Rank analysis of incomplete block designs: I. the method of paired comparisons. *Biometrika*, 39(3/4):324–345, 1952.

Nitay Calderon, Roi Reichart, and Rotem Dror. The alternative annotator test for LLM-as-a-judge: How to statistically justify replacing human annotators with LLMs. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics*, 2025.

Alexander R. Fabbri, Wojciech Kryściński, Bryan McCann, Caiming Xiong, Richard Socher, and Dragomir Radev. SummEval: Re-evaluating summarization evaluation. *Transactions of the Association for Computational Linguistics*, 9:391–409, 2021.

Olga Golovneva, Moya Chen, Spencer Poff, Martin Corredor, Luke Zettlemoyer, Maryam Fazel-Zarandi, and Asli Celikyilmaz. ROSCOE: A suite of metrics for scoring step-by-step reasoning. In *International Conference on Learning Representations*, 2023.

Klaus Krippendorff. *Content Analysis: An Introduction to Its Methodology*. SAGE Publications, 4 edition, 2018.

Jiwei Li and Dan Jurafsky. Neural net models of open-domain discourse coherence. In *Proceedings of EMNLP*, 2017.

Yang Liu, Dan Iter, Yichong Xu, Shuohang Wang, Ruochen Xu, and Chenguang Zhu. G-Eval: NLG evaluation using GPT-4 with better human alignment. In *Proceedings of EMNLP*, 2023. arXiv:2303.16634.

William C. Mann and Sandra A. Thompson. Rhetorical structure theory: Toward a functional theory of text organization. *Text*, 8:243–281, 1988.

Sewon Min, Kalpesh Krishna, Xinxi Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of EMNLP*, 2023.

Rashmi Prasad, Bonnie Webber, and Alan Lee. Penn discourse treebank version 3.0. LDC2019T05, Linguistic Data Consortium, 2019.

Yixiao Song, Yekyung Kim, and Mohit Iyyer. VeriScore: Evaluating the factuality of verifiable claims in long-form text generation. In *Findings of the Association for Computational Linguistics: EMNLP 2024*, 2024.

Jerry Wei, Chengrun Yang, Xinying Song, Yifeng Lu, Nathan Hu, Jie Huang, Dustin Tran, Daiyi Peng, Ruibo Liu, Da Huang, Cosmo Du, and Quoc V. Le. Long-form factuality in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

Qingyu Xie, Wenxuan Zheng, Xiaoyu Shen, and Rui Xia. Lorefact: Bridging the logic gap in fact-checking. In *Findings of the Association for Computational Linguistics: ACL 2026*, pp. 6974–6996, 2026.

Wei Zhao, Michael Strube, and Steffen Eger. DiscoScore: Evaluating text generation with BERT and discourse coherence. *arXiv preprint arXiv:2201.11176*, 2023.

A THE GENERATION AND VALIDATION PROMPTS, VERBATIM

Every prompt follows, exactly as issued, written to the files included here by `scripts/export_locobench_prompts.py` and checked against the code by a test. Placeholders have the form `[NAME_PLACEHOLDER]` and are substituted by the contract of Appendix C. Line numbers are added for reference. Each prompt is preceded by a short card: its inputs, what it must emit, and what checks its output. The source code and the released generation log keep the validators' original identifiers, under which the prose audit is `V3` and claim coverage `V4`; this paper numbers them V2 and V3.

A.1 P1 — TOPIC TO QUESTION

Input: a canonical topic. **Output:** one open-ended question. **Checked by:** the parser (a single question). 1,902 characters, 9 numbered instructions. Placeholder: `[TOPIC_PLACEHOLDER]`.

```
Instructions:
1. You are given a TOPIC. Generate one open-ended QUESTION about that topic.
2. The QUESTION must begin with Why, How, What, Does, or Which.
3. The QUESTION must elicit an answer containing MULTIPLE, DIVERSE logical
   relationships. Aim for an answer that would naturally contain at least four
   of the following: a causal chain; a condition; a temporal ordering; an
   evidential appeal; a general claim with a specific instance; a restatement
   or definitional equivalence; a non-exhaustive contrast; an EXHAUSTIVE pair
   of competing alternatives of which EXACTLY ONE can hold; a disjunctive pair
   of which AT LEAST ONE must hold; a tension that an authoritative finding
   later resolves.
4. Prefer topics where competing explanations are adjudicated: incident
   investigations, litigation, diagnostic reasoning, attribution disputes,
   competing scientific hypotheses. Such topics naturally produce exhaustive
   alternatives and resolving holdings.
5. The QUESTION must be answerable in 500-900 words of formal prose.
6. The QUESTION must not itself contain the words "cause", "condition",
   "alternative", or "contrast", so the answer's relations are not primed
   lexically.
7. Do NOT ask a multi-part question joined by "and" - one question only.
8. Wrap the QUESTION in square brackets.
9. Output the bracketed QUESTION only, with no preamble.

Use the following examples to better understand your task.

TOPIC: aviation component failure investigation
QUESTION: [How do investigators determine responsibility when a component
supplied by a third party fails in service and the operator and the
manufacturer offer competing explanations?]

TOPIC: antibiotic resistance in hospital settings
QUESTION: [Why does resistance emerge in some hospital units but not others,
and how is a specific outbreak traced to its origin?]

Your Task:
TOPIC: [TOPIC_PLACEHOLDER]
QUESTION:
```

A.2 P2 — QUESTION TO TAGGED CLAIMS

Input: the P1 question. **Output:** 26 atomic claims, each tagged correct, incorrect, alt-pair, disj-pair, equiv-pair or holding. **Checked by:** the parser (count and tag composition). 3,752 characters, 7 numbered instructions. Placeholder: `[QUESTION_PLACEHOLDER]`.

```
Instructions:
1. You are given a QUESTION. Generate a set of ATOMIC CLAIMS that a
   well-informed answer to it might contain.
2. An atomic claim is a sentence containing a singular piece of information,
   which should be basic, indivisible, and independently verifiable.
3. Generate exactly 26 atomic claims: the 24 below, plus the 2 holding claims
   required by instruction 4. The 24 break down as:
   (a) 14 factually correct claims;
   (b) 4 factually incorrect claims (plausible but false);
   (c) 2 EXHAUSTIVE ALTERNATIVE claims: a pair X, Y such that exactly one of X
       and Y must be true - they cannot both hold, and they cannot both fail.
       Example pair: "No one was harmed in the incident." / "Three people died
       in the incident."
   (d) 2 DISJUNCTIVE claims: a pair X, Y such that at least one must be true,
```

```
       but both may be true simultaneously. Make each one a STATE OF THE WORLD,
       not a finding or a check that reported something: "the alloy was
       substituted" and "the batch was mislabelled", not "assay A flagged it"
       and "assay B flagged it".
   (e) 2 EQUIVALENT claims: a pair X, Y asserting the same fact in different
       words, so that either one being true makes the other true.
4. Additionally generate 2 HOLDING claims: sentences in which an authority
   (a court, tribunal, regulator, final report, or review board) reaches a
   determination that settles a dispute between two other claims in your list.
   A holding claim must use adjudicating language - "held", "ruled",
   "concluded", "determined", "found that", "ultimately".
5. Each claim must be a single declarative sentence with no connective linking
   it to another claim. Do NOT use "because", "although", "therefore", "if",
   "while", "however". Relations are added at a later stage, not here.
6. Do NOT indicate which claims are factually incorrect in their wording; do
   NOT hedge with "assume", "might", "possibly", or "allegedly".
7. Output as a markdown code block: an unordered list with "-" bullets. Append
   a tag in square brackets to every claim, one of: [correct], [incorrect],
   [alt-pair-1], [alt-pair-2], [disj-pair-1], [disj-pair-2], [equiv-pair-1],
   [equiv-pair-2], [holding].

Use the following example to better understand your task.

QUESTION: [How do investigators determine responsibility when a component
supplied by a third party fails in service and the operator and the
manufacturer offer competing explanations?]

ATOMIC CLAIMS:
```
- AeroParts supplied turbine blades to the carrier. [correct]
- The turbine blades were installed in the carrier's A-320 fleet. [correct]
- Seventeen turbine blades failed in flight during service. [correct]
- The regulator opened a safety probe into the blade failures. [correct]
- Maintenance logs recorded abnormal vibration on the failed blades. [correct]
- The airline argued that the failures were caused by pilot error. [correct]
- The final report concluded the root cause was a metallurgical defect. [correct]
- The regulator issued an order grounding the fleet. [correct]
- Every affected aircraft was pulled from service. [equiv-pair-1]
- The regulator's order removed all affected aircraft from operation. [equiv-pair-2]
- The blades were certified under a 1998 airworthiness standard. [incorrect]
- No one was harmed in the incidents. [alt-pair-1]
- Three passengers died in one of the failures. [alt-pair-2]
- The alloy supplied for the blades was substituted. [disj-pair-1]
- The blade batch was mislabelled at the mill. [disj-pair-2]
- The tribunal held that AeroParts was liable for the blade failures. [holding]
- The accident report found the pilots not at fault. [holding]
```

Your Task:
QUESTION: [QUESTION_PLACEHOLDER]
ATOMIC CLAIMS:
```

A.3 P3 — CLAIMS TO RELATION PLAN

Input: the question and the tagged claims. **Output:** a JSON plan of 14–16 positioned claims, ten relations with sense, strength band, validity and error kind, and 4–6 non-relations. **Checked by:** the parser (counts, legal senses and couplings, error kinds, one relation per pair, disjointness of the relation and non-relation lists), whose complaint is fed back as a format retry, then the plan gate (the remaining P3 rows of Table 4), whose complaint is fed back as a content retry. Section 5.3 lists which of the prompt's requirements are checked, and with what tolerance. 7,778 characters, 10 numbered instructions. Placeholders: `[QUESTION_PLACEHOLDER]`, `[CLAIMS_PLACEHOLDER]`.

```
Instructions:
1. You are given a QUESTION and a tagged list of ATOMIC CLAIMS. Produce a
   RELATION PLAN: a machine-readable graph of the logical relations a response
   should draw between those claims.
2. Select 14-16 claims to use. Aim for 15. Number them with FRESH consecutive
   positions: the first claim you assert is pos 1, the next pos 2, and so on
   with NO gaps and NO repeats, ending at pos N where N is how many you
   selected. Do NOT carry over the numbering of the input list - if you pick the
   3rd, 7th and 12th claims they become pos 1, 2 and 3. The position is the
   order the response will assert the claim in. Every selected claim keeps its
   text EXACTLY as given. Do NOT introduce information not in the claims, and
   do NOT alter the claims. Your selection MUST INCLUDE AT LEAST 2 of the
   [incorrect] claims: the response has to assert false content alongside the
   true, so a selection drawn only from the [correct] claims is rejected. Carry
   each selected claim's tag through as a "factual" field - false for an
   [incorrect] claim, true for every other tag.
```

```
3. Plan EXACTLY 10 RELATIONS. Each relation connects exactly two selected claims
   and has a SENSE. Every relation must join a DIFFERENT pair of claims: once two
   claims are related, that pair is used up and must not appear again in the
   relations list, in either direction. The SENSE is one of: Cause-Effect,
   Effect-Cause, Evidence, Condition,
   Restatement, Instantiation, Contrast, Concession, Alternative, Disjunction,
   Precedence, Succession.
   - Cause-Effect: source causes target. Effect-Cause: source is the effect,
     target its cause.
   - Evidence: source is evidence for target. Condition: source is a condition
     for target.
   - Restatement: source and target assert the same thing.
   - Instantiation: source is a general claim, target a specific instance.
   - Contrast: opposition that is NOT exhaustive - both could be false.
   - Concession: a tension the text itself resolves.
   - Alternative: EXHAUSTIVE competing options - EXACTLY ONE holds. Not both,
     and not neither. Use your [alt-pair-*] claims here.
   - Disjunction: AT LEAST ONE holds; both may hold. Use your [disj-pair-*]
     claims here.
   - Precedence / Succession: ordered in time with NO truth dependence -
     neither claim being true or false says anything about the other.
4. Keep relations LOCAL: prefer pairs within 4 POSITIONS of each other, so the
   graph stays recoverable from the response by a local reader. A longer-range
   pair is acceptable when the two claim texts share a named entity, and a few
   long-range links (for instance a closing claim answering the opening one) are
   fine. Order the positions so that related claims sit near each other.
5. The plan must include AT LEAST ONE relation of each of these senses:
   Alternative, Disjunction, Restatement, Concession, and one of
   Precedence/Succession. Separately, [LADDER_PLACEHOLDER]
6. Every Concession must be marked resolved or unresolved. A resolved
   Concession must name a RESOLVER: a [holding] claim, also selected and
   positioned AFTER both endpoints.
7. Give each relation an intended STRENGTH BAND: strong (0.85-1.00), moderate
   (0.60-0.84), or weak (0.35-0.59) - how strongly the source determines the
   target under that sense. Restatement relations are always strong.
8. Give each relation an intended VALIDITY: valid if the relation genuinely
   holds between the claims as stated; invalid if it does not. Of your 10
   relations, mark EXACTLY 6 "valid" and EXACTLY 4 "invalid" - count them before
   you emit, and do not compute a percentage. For each invalid relation give an
   ERROR_KIND, one of: wrong_sense, wrong_direction, false_endpoint, spurious.
9. Also list 4-6 NON-RELATIONS (aim for 5): pairs of selected claims within 4
   positions of each other between which the response must draw NO logical
   dependence. These are the distractors a relation miner must correctly label
   None. A non-relation pair must NOT also appear in your relations list, in
   either direction - the two lists are disjoint.
10. Output a single JSON object in a code block and nothing else. Below is a
    COMPLETE, correctly-shaped example - copy its structure exactly, not its
    text. It has 14 atoms numbered 1..14 with no gaps, 2 of them factually
    incorrect, 10 relations (6 valid, 4 invalid), all five required senses, and 5
    non-relations disjoint from the relations:

```json
{
  "atoms": [
    {"pos": 1, "text": "<claim text>", "factual": true},
    {"pos": 2, "text": "<claim text>", "factual": true},
    {"pos": 3, "text": "<claim text>", "factual": true},
    {"pos": 4, "text": "<claim text>", "factual": false},
    {"pos": 5, "text": "<claim text>", "factual": true},
    {"pos": 6, "text": "<claim text>", "factual": true},
    {"pos": 7, "text": "<claim text>", "factual": true},
    {"pos": 8, "text": "<claim text>", "factual": true},
    {"pos": 9, "text": "<claim text>", "factual": false},
    {"pos": 10, "text": "<claim text>", "factual": true},
    {"pos": 11, "text": "<claim text>", "factual": true},
    {"pos": 12, "text": "<claim text>", "factual": true},
    {"pos": 13, "text": "<claim text>", "factual": true},
    {"pos": 14, "text": "<claim text>", "factual": true}
  ],
  "relations": [
    {"source_pos": 1, "target_pos": 2, "sense": "Cause-Effect",
     "strength_band": "strong", "validity": "valid",
     "error_kind": null, "resolved": null, "resolver_pos": null},
    {"source_pos": 2, "target_pos": 3, "sense": "Evidence",
     "strength_band": "moderate", "validity": "valid",
     "error_kind": null, "resolved": null, "resolver_pos": null},
    {"source_pos": 3, "target_pos": 4, "sense": "Restatement",
     "strength_band": "strong", "validity": "valid",
     "error_kind": null, "resolved": null, "resolver_pos": null},
    {"source_pos": 5, "target_pos": 6, "sense": "Alternative",
     "strength_band": "strong", "validity": "valid",
     "error_kind": null, "resolved": null, "resolver_pos": null},
    {"source_pos": 7, "target_pos": 8, "sense": "Disjunction",
```
```

```
     "strength_band": "moderate", "validity": "valid",
     "error_kind": null, "resolved": null, "resolver_pos": null},
    {"source_pos": 9, "target_pos": 10, "sense": "Concession",
     "strength_band": "moderate", "validity": "valid",
     "error_kind": null, "resolved": true, "resolver_pos": 11},
    {"source_pos": 11, "target_pos": 12, "sense": "Precedence",
     "strength_band": "weak", "validity": "invalid",
     "error_kind": "wrong_sense", "resolved": null, "resolver_pos": null},
    {"source_pos": 12, "target_pos": 13, "sense": "Cause-Effect",
     "strength_band": "moderate", "validity": "invalid",
     "error_kind": "wrong_direction", "resolved": null, "resolver_pos": null},
    {"source_pos": 13, "target_pos": 14, "sense": "Contrast",
     "strength_band": "weak", "validity": "invalid",
     "error_kind": "spurious", "resolved": null, "resolver_pos": null},
    {"source_pos": 4, "target_pos": 6, "sense": "Contrast",
     "strength_band": "weak", "validity": "invalid",
     "error_kind": "false_endpoint", "resolved": null, "resolver_pos": null}
  ],
  "non_relations": [
    {"source_pos": 1, "target_pos": 4}, {"source_pos": 2, "target_pos": 5},
    {"source_pos": 6, "target_pos": 9}, {"source_pos": 8, "target_pos": 11},
    {"source_pos": 10, "target_pos": 13}
  ]
}
```

```
Your Task:
QUESTION: [QUESTION_PLACEHOLDER]
ATOMIC CLAIMS: [CLAIMS_PLACEHOLDER]
RELATION PLAN:
```

A.4 P4 — PLAN TO RESPONSE

Input: the question and the plan. **Output:** 550–650 words of prose. **Checked by:** the parser (word floor), then V1, V2 and V3 on the committee, whose combined complaint is fed back as a content retry. 5,571 characters, 8 numbered instructions. Placeholders: `[QUESTION_PLACEHOLDER]`, `[PLAN_PLACEHOLDER]`.

```
Instructions:
1. You are given a QUESTION and a RELATION PLAN in JSON. Write a RESPONSE that
   answers the question and realizes the plan exactly.
2. Assert every planned atom EXACTLY ONCE, in the plan's POSITION order,
   preserving each atom's factual content. Do not restate an atom you have
   already asserted in order to attach a later relation to it - refer back with
   a pronoun or a short noun phrase instead. Name that phrase after the thing
   the atom is about, not after its status as a finding: "the primitive carpal
   morphology" or "that wrist evidence", not "the analysis's flagging of
   primitive traits"; "these two accounts" or "either possibility", not "the two
   claims in circulation". A summary phrase covering both endpoints of one
   connective is fine - "at least one of these must hold: X, or Y". You may
   adjust surface wording for fluency but must not change what is asserted.
   Quantifiers, determiners, numbers, dates, polarity and modals are CONTENT,
   not wording: keep "all", "only", "no", "some", "each", "never" and every
   figure exactly as the atom has them. "All tools were crafted by a single
   culture" may become "every one of the tools came from a single culture", but
   never "the tools were crafted by a single culture" - dropping "All" weakens
   the claim and counts as changing it.
   Never let one atom absorb another: each must stay separately true or false.
   Joining two atoms with a connective is fine and often required - "either X or
   Y" and "although X, Y" both assert two atoms and are exactly what instruction
   3 asks for.
3. Realize every planned RELATION so a careful reader could recover it from the
   text alone, using language appropriate to its sense:
   - Cause-Effect / Effect-Cause: "caused", "led to", "resulted from", "because"
   - Evidence: "indicates", "confirms", "as shown by"
   - Condition: "if", "provided that", "only when"
   - Restatement: "that is", "in other words", "equivalently"
   - Instantiation: "for example", "specifically", "in one case"
   - Contrast: "by contrast", "whereas", "on the other hand"
   - Concession: "although", "despite", "even though"
   - Alternative: make the exhaustiveness EXPLICIT - "either ... or", "one of
     these accounts must be wrong", "the two cannot both be true, and one of
     them must hold"
   - Disjunction: make the at-least-one reading EXPLICIT - "at least one of",
     "one or both", "perhaps both". Say "perhaps both", never "possibly both".
     A bare "X, or Y" is NOT enough for either of these two senses: on its own it
     leaves a reader unable to tell whether both claims could hold, which is the
     only thing separating Alternative from Disjunction. Every such pair needs its
     marker - "either X or Y" when exactly one holds, "at least one of X or Y,
```

```
     perhaps both" when both may.
   - Precedence / Succession: state the sequence EXPLICITLY in one sentence
     joining both claims - "before", "after", "subsequently", "earlier than",
     "which postdates" - while keeping the ordering purely chronological, so a
     reader sees only when things happened and infers nothing from it. Write
     "X, which predates Y" or "Y came after X", not "X and therefore Y".
     Asserting the two claims in separate sentences leaves this relation
     unrealized: an ordering a reader cannot see is the same as no ordering.
4. For each planned NON-RELATION, assert both atoms but draw no dependence
   between them. Separate them by DISTANCE and PARAGRAPHING, not by a marker:
   put other material between them, or place them in different paragraphs, and
   simply use no connective. Do NOT announce the absence of a link - phrases
   like "In an unrelated matter", "Separately", or "On a different note" read as
   stitched-together annotation rather than prose and will be marked down.
5. For a resolved Concession, assert the resolver atom after both endpoints and
   present it as settling the tension.
6. Write it as a continuous expository answer to the QUESTION - an argument a
   domain expert would publish, organized into 3-5 paragraphs, in which the
   planned relations are simply how the argument hangs together. Do NOT walk the
   plan relation by relation: a reader must not be able to tell that a plan
   exists. Restate a subject to carry an argument forward ("the radiocarbon
   date", "that attribution"), never the earlier sentence as a step ("the
   finding just reported", "the provision just stated", "the ordering just
   described"). Use precise and formal language, avoiding vague generalizations
   and rhetorical fillers. Aim for 550-650 words. The 500-word floor is a hard
   requirement checked by an automated reader, which rejects a shorter answer
   outright however good it is, so give each atom and each relation the space to
   be argued rather than compressing them into a summary.
7. Do NOT mention the relation plan, the senses, the strength bands, the
   validity labels, or the fact that some content is incorrect. Do NOT hedge
   with "assume", "might", "possibly", "allegedly", "supposedly",
   "reportedly", or "it is claimed" around planned-invalid relations - assert
   them in the same register as the valid ones. Do NOT add headings, bullet
   lists, or annotations.
8. Wrap the RESPONSE in a code block whose contents are the prose itself, as you
   would publish it. The block holds the answer, not a JSON object or any other
   payload with the prose inside a field.

Your Task:
QUESTION: [QUESTION_PLACEHOLDER]
RELATION PLAN: [PLAN_PLACEHOLDER]
RESPONSE:
```

A.5 P5 — RESPONSE TO PERTURBED RUNG

Input: the current text, the plan, and one operator call with its target. **Output:** the edited text and a structured diff. **Checked by:** the text-changed and length-drift gates, then the adjacency gate over the whole family. 2,592 characters, 5 numbered instructions; instruction 3 defines all eleven calls of Table 3. Placeholders: `[RESPONSE_PLACEHOLDER]`, `[PLAN_PLACEHOLDER]`, `[OPERATOR_PLACEHOLDER]`.

```
Instructions:
1. You are given a RESPONSE, its RELATION PLAN, and one PERTURBATION to apply.
2. Apply ONLY the named perturbation. Change the minimum number of words
   needed. Every other relation, atom, and sentence must be preserved verbatim.
3. The perturbations are defined as follows.
   - wrong_sense(edge, new_sense): rewrite the connective realizing that edge so
     the text draws new_sense instead. Keep both atoms' content unchanged.
   - direction_reversal(edge): rewrite so the relation runs from the target to
     the source. Keep both atoms' content unchanged.
   - exhaustiveness_flip(edge): if the edge is Alternative, rewrite so the two
     claims merely oppose without being exhaustive (both could be false). If it
     is Contrast, rewrite so exactly one must hold.
   - remove_resolution(edge): delete the resolver sentence and any reference to
     it, leaving the conceded tension unresolved. Change nothing else.
   - add_resolution(edge): the edge is a conflict the response leaves open. Add
     ONE sentence, after both of its atoms, in which an authority reaches a
     determination that settles the tension. Do not delete either atom or change
     the connective realizing the edge. Change nothing else.
   - inject_contradiction(atom): add ONE new sentence asserting a claim that
     directly contradicts the named atom. Place it at least two sentences later.
   - break_chain(edge): remove or negate the link realizing that edge so the
     causal chain no longer connects, keeping both endpoint atoms present.
   - shuffle_order(level): reorder the response's sentences without changing any
     wording. level=full permutes all sentences; level=block permutes paragraph
     blocks; level=adjacent swaps one adjacent pair.
   - spurious_relation(atom_a, atom_b): add a connective asserting a causal or
     contradictory link between two atoms the original response kept separate.
     Add no new factual content.
```

```
   - drop_relation(edge): remove the connective realizing that edge so the two
     atoms are merely adjacent assertions. Keep both atoms.
   - ordering_only(edge): rewrite a Precedence edge as Succession, or the
     reverse, preserving the actual chronology and adding NO causal or
     inferential language. The truth of neither atom may become dependent on
     the other.
4. Do not change the response's length by more than 15 percent, except for
   inject_contradiction / add_resolution (which add one sentence) and
   remove_resolution / break_chain (which remove one).
5. Output the perturbed RESPONSE in a code block, then a second code block
   containing a JSON diff:
   {"operator": "...", "target": ..., "sentences_changed": [...],
    "relations_added": [...], "relations_removed": [...],
    "relations_relabeled": [...]}

Your Task:
RESPONSE: [RESPONSE_PLACEHOLDER]
RELATION PLAN: [PLAN_PLACEHOLDER]
PERTURBATION: [OPERATOR_PLACEHOLDER]
```

A.6 V1 — RELATION RECOVERY

Input: the response and its claims keyed $1..N$ — *not* the plan. **Output:** the relations the response draws, with sense, coupling, strength band and resolved flag. **Decision:** any-of over the committee, thresholds of Table 4. 4,086 characters. Placeholders: `[RESPONSE_PLACEHOLDER]`, `[ATOMS_PLACEHOLDER]`. The validators use a function-signature idiom rather than numbered instructions. The final rule states the compile map, so a validator's coupling must be the one its sense compiles to.

```
Let's define a function named recover_relations(response: str,
atoms: dict[str, str]).
Given a RESPONSE and the ATOMS drawn from it, the function returns the
list of logical relations the response ITSELF draws between those atoms.

ATOMS maps an atom's NUMBER to its text, numbered from 1. Use those numbers, as
integers, to identify atoms in your output: if the relation holds between the
atoms numbered "3" and "7", emit source 3 and target 7. Never renumber them and
never count from zero.

For each ORDERED pair (i, j) that the response actually relates, emit one object:
  {"source": i, "target": j,
   "sense": one of ["Cause-Effect","Effect-Cause","Evidence","Condition",
     "Restatement","Instantiation","Contrast","Concession","Alternative",
     "Disjunction","Precedence","Succession"],
   "coupling": one of ["entailment","contradiction","equivalence","exclusive",
     "co_necessity","none"],
   "strength_band": one of ["strong","moderate","weak"],
   "resolved": true or false or null}

Rules:
- Emit a relation ONLY if the response itself draws that connection, as written
  or as a clear step in the author's argument. Do not emit a relation that is
  merely plausible in general.
- Pick the MOST SPECIFIC sense the text supports, never the most general one.
  Five senses compile to "entailment", and "Evidence" is the broadest of them,
  so it is the easy wrong answer: use it only for evidential support ("indicates",
  "confirms", "as shown by", "suggests"). Prefer "Cause-Effect" when the text says
  one thing brought the other about ("caused", "led to", "resulted in",
  "because"), "Effect-Cause" when that runs the other way ("resulted from",
  "is explained by"), "Condition" for "if", "provided that", "only when", and
  "Instantiation" when one claim is an INSTANCE of the other, i.e. a specific
  case of a general statement ("for example", "specifically", "such as").
- Use "Alternative" only when the two claims cannot both be true AND cannot both
  be false. Use "Contrast" when they merely cannot both be true - look for "by
  contrast", "whereas", "on the other hand", or two claims set against each other
  without either being conceded.
- Use "Concession" when the text grants a tension and then overrides it -
  "although", "despite", "even though", "notwithstanding", or a clause that admits
  one claim while asserting the other anyway. A concessive marker attaching two
  claims IS a relation between them, even when the concession sits in a
  subordinate clause rather than its own sentence.
- Use "Disjunction" only when the response rules out both claims being false
  ("at least one of", "one or both", "perhaps both").
- Use "Precedence" or "Succession" only when the ordering carries no truth
  dependence; set their coupling to "none". These ARE relations and must be
  emitted: "before", "after", "subsequently", "earlier", "later", or an explicit
  statement that a sequence carries no causal or inferential force. Read such a
  pair as "Precedence" when the source comes first in time and "Succession" when
```

```
  it comes second. Do not silently upgrade one to "Cause-Effect" or "Evidence"
  because the two events sit next to each other in an argument.
- Relate a pair only ONCE unless the response genuinely draws two different
  connections between them. Emitting many relations does not make the recovery
  more complete; a pair the response merely mentions in the same paragraph is not
  related.
- "coupling" is not an independent choice: it must be the coupling that the
  "sense" you chose compiles to. Cause-Effect, Effect-Cause, Evidence,
  Condition and Instantiation are "entailment"; Restatement is "equivalence";
  Contrast and Concession are "contradiction"; Alternative is "exclusive";
  Disjunction is "co_necessity"; Precedence and Succession are "none".
- Pairs the response leaves unrelated must be omitted entirely.

The function returns a JSON list. Note that your response will be passed to the
python interpreter, SO NO OTHER WORDS!

recover_relations(response="[RESPONSE_PLACEHOLDER]", atoms=[ATOMS_PLACEHOLDER])
```

A.7 V2 — PROSE AUDIT

Input: the response. **Output:** three 1–5 scores and three span lists, every span quoted verbatim. **Decision:** majority per facet. 2,455 characters. Placeholder: `[RESPONSE_PLACEHOLDER]`. The list of spans that must *not* be flagged as leakage is the narrow definition described in §6.2.

```
Let's define a function named check_response(response: str).
The function audits a generated response and returns a JSON object:
  {"fluency": 1-5, "formality": 1-5, "organization": 1-5,
   "leakage": [list of quoted spans], "hedging": [list of quoted spans],
   "artifacts": [list of quoted spans]}

- fluency, formality, organization: 5 = indistinguishable from expert human
  prose; 1 = badly degraded.
- leakage: any span revealing the response was generated from an annotation
  plan. Read this NARROWLY: naming a relation type or sense; naming the plan,
  the strength bands, or the validity labels; referring to "claims", "atoms" or
  "statements" as the objects the text was assembled from; or any label, tag,
  numeric index, or bracketed marker. Do NOT flag a connective or coordination
  that merely realizes a relation - "either X or Y", "at least one of X or Y",
  "at least one of these was true", "perhaps both", "one or both",
  "the two cannot both be true, and one of them must hold",
  "one of these accounts must be wrong", "although X, Y",
  "before", "after", "subsequently", "earlier than", "which postdates",
  "as indicated by", "as shown by", "which indicates", or an
  explicit statement that an ordering carries no causal or inferential force. A
  writer is REQUIRED to use those, and a coupling over two claims' joint truth
  cannot be expressed without one. Every span you report must be QUOTED VERBATIM
  from the response; if you cannot copy it out of the text, it is not there and
  must not be reported. Do NOT
  flag ordinary scholarly vocabulary used about the subject matter, such as
  "claim", "outcome", "reading", "account" or "finding".
- hedging: a span using one of these EXACT words: "assume", "might",
  "possibly", "allegedly", "supposedly", "reportedly", "it is claimed". The list
  is CLOSED - quote the word you found, and flag nothing that does not contain
  one of them. Uncertainty a writer is REQUIRED to express is not hedging: "at
  least one of X or Y, perhaps both" and "either X or Y" state the logical
  strength of a relation, not doubt about the content, so do NOT flag them.
- artifacts: enumerated list structure, headings, repeated template phrasing, or
  a sentence that reads as a stitched-together claim rather than prose.

Return the JSON object only. Note that your response will be passed to the
python interpreter, SO NO OTHER WORDS!

check_response(response="[RESPONSE_PLACEHOLDER]")
```

A.8 V3 — CLAIM COVERAGE

Input: the response and its claims. **Output:** a status and quoted span per claim. **Decision:** majority per claim over the gating statuses, with altered also gating on a planted-false claim. 1,403 characters. Placeholders: `[RESPONSE_PLACEHOLDER]`, `[ATOMS_PLACEHOLDER]`.

```
Let's define a function named check_atom_coverage(response: str,
atoms: list[str]).
For each atom, the function decides whether the response asserts it, returning a
JSON list of objects:
  {"index": i,
```

```
   "status": one of ["asserted", "altered", "missing", "merged"],
   "span": the quoted span of the response that asserts it, or null,
   "position": the 1-based index of that span among the response's sentences}

- "asserted": the response asserts the atom's full content, in any wording.
- "altered": the response asserts something related but changes the atom's
  content (a different number, date, entity, or polarity).
- "missing": the response does not assert the atom.
- "merged": the atom's content has been absorbed into another atom's, so it is no
  longer separately assertable - you could not negate this atom alone and leave
  the other standing. Reserve this for genuine loss of content.
  Two atoms joined by a connective are NOT merged: "either X or Y", "at least one
  of X or Y", "although X, Y", "X, whereas Y", and "X; that is, Y" each assert
  BOTH atoms, because each remains independently true or false. Mark both
  "asserted", with the span covering the part that asserts each one.

Return the JSON list only. Note that your response will be passed to the python
interpreter, SO NO OTHER WORDS!

check_atom_coverage(response="[RESPONSE_PLACEHOLDER]", atoms=[ATOMS_PLACEHOLDER])
```

A.9 Retry feedback

When an attempt fails, the prompt is re-issued with one of two notes appended, the failure reason substituted for `{reason}`. The *format* note follows a parse failure; the *content* note follows a gate failure and tells the model its format was correct. Neither note is a separate prompt: each is appended to the stage's own prompt, and a test asserts that neither contains text that could route the retry to another prompt's parser.

Format note (parse failure).

```
NOTE: your previous attempt could not be read by the automated reader.
Reason: {reason}
Emit only the output described above, in exactly the required format.
```

Content note (gate failure).

```
NOTE: your previous attempt was well-formed, but its content did not meet a quality
↪  requirement.
Problem: {reason}
Keep the same output format and revise the content so this is fixed.
```

B Relation-mining and baseline prompts, verbatim

These prompts are not part of the benchmark: they are the prompts of the measures we evaluate on it in §10, reproduced so that those results can be replicated. Placeholders here use `{{name}}` (mining) or `{name}` (judges) and are filled per call.

B.1 Mining prompt A — sense and coupling

Issued once per candidate claim pair, with the full response as context, so that the model asserts only relations the response draws. The bracketed `[coupling=...]` tag is read from the output, and its token probabilities give the type confidence. The released mined arms use this variant with the full sense menu.

```
Instructions:
You are given the full RESPONSE a model produced, and two atomic claims, A and B, both taken
↪  FROM THAT RESPONSE, in their order of appearance (A comes before B). Your task is to
↪  decide the discourse/logical relation FROM A TO B, following the steps below.

IMPORTANT -- ground your decision in the response. Assert a coupling ONLY if the response
↪  ITSELF draws that connection between A and B (as written, or as a clear step in the
↪  author's argument/narrative). Do NOT assert a relation that is merely plausible in
↪  general but that the response does not actually make. If A and B both appear in the
↪  response yet the response draws no logical or discourse dependence between them, the
↪  answer is None.

```

```
1. Reason step by step, referring to the response: does the response present A as causing, enabling, providing evidence for, restating, elaborating, temporally preceding, contrasting with, or contradicting B? Is B a claim the response later withdraws or that a holding resolves? Consider the direction (A to B). If the response links A and B only indirectly through other claims, or not at all, that is None.

2. Name the DISCOURSE SENSE, one of: Cause-Effect, Effect-Cause, Evidence, Condition, Restatement, Instantiation, Contrast, Concession, Alternative, Disjunction, Precedence, Succession, None.
   - Cause-Effect: A causes/leads to B. Effect-Cause: A is the effect, B its cause.
   - Evidence: A provides evidence for B. Condition: A is a condition for B.
   - Restatement: A and B assert the same thing. Instantiation: A is a general claim, B a specific instance (or vice versa).
   - Contrast: A and B are in opposition but NOT exhaustive (they need not cover all possibilities; both could conceivably be false).
   - Concession: A and B are in tension but the text concedes/resolves it ("although A, still B", or a holding settles it).
   - Alternative: A and B are EXHAUSTIVE competing options -- EXACTLY ONE holds (they are mutually exclusive AND together cover the possibilities: not both, and not neither). E.g. "no one was harmed" vs "three people died"; "the cause was pilot error" vs "the cause was a metallurgical defect".
   - Disjunction: AT LEAST ONE of A and B holds (they may both hold, but the response rules out neither being true) -- e.g. two supporting findings at least one of which must be present.
   - Precedence/Succession: A and B are ordered in time with no truth dependence.
   - None: the response draws no logical or discourse dependence between A and B.

3. Map the sense to a COUPLING, one of: entailment, contradiction, equivalence, exclusive, co_necessity, none.
   - Cause-Effect, Effect-Cause, Evidence, Condition, Instantiation -> entailment
   - Restatement -> equivalence
   - Contrast, Concession -> contradiction
   - Alternative -> exclusive       (exactly one of A, B is true)
   - Disjunction -> co_necessity     (at least one of A, B is true)
   - Precedence, Succession, None -> none
   Prefer "exclusive" over "contradiction" when the two claims are not just incompatible but EXHAUSTIVE (one of them must be true); prefer "contradiction" when they merely cannot both hold but could both be false.

4. Give your final answer as two bracketed tags on ONE line, sense first:
[sense=Cause-Effect] [coupling=entailment]
A JSON object {"sense":"Cause-Effect","coupling":"entailment"} is also acceptable.

Use the following examples to better understand your task.

Example 1 (the response makes the causal link):
RESPONSE: The company launched a flawed product last quarter. Reviewers panned it, returns spiked, and the company's stock price fell 15 percent over the same period.
A: The company launched a flawed product last quarter.
B: The company's stock price fell 15 percent last quarter.
1. Reasoning: the response presents the flawed launch as the head of a chain (panning, returns) that ends in the stock decline, so the response itself draws a causal link from A to B.
2. Discourse sense: Cause-Effect.
3. Coupling: A causing B is a positive inferential link, i.e. entailment.
4. Final answer:
[sense=Cause-Effect] [coupling=entailment]

Example 2 (both claims present, but the response draws NO connection -> None):
RESPONSE: The quarterly report was published in April. Separately, the annual audit was scheduled for December. The two processes are run by different teams and were not related this year.
A: The quarterly report was published in April.
B: The annual audit was scheduled for December.
1. Reasoning: both claims appear in the response, and one might imagine a reporting-to-audit link in general, but this response explicitly treats them as separate and unrelated. The response draws no dependence from A to B.
2. Discourse sense: None.
3. Coupling: no dependence the response asserts, i.e. none.
4. Final answer:
[sense=None] [coupling=none]

Example 3 (the response states an EXHAUSTIVE alternative -> exclusive):
RESPONSE: The official statement said no one was harmed in the incident. However, the coroner's report confirmed that three people died in the incident.
A: No one was harmed in the incident.
B: Three people died in the incident.
1. Reasoning: the response sets A and B against each other ("However, ...") and they cannot both be true; but they also cannot both be false -- either people were harmed or they were not -- so exactly one holds. This is exhaustive, not a mere contrast. No holding resolves it.
2. Discourse sense: Alternative.
```

```
3. Coupling: exactly one of A, B is true, i.e. exclusive.
4. Final answer:
[sense=Alternative] [coupling=exclusive]

Example 4 (at least one must hold -> co_necessity):
RESPONSE: The defect was caught in review: at least one of the two independent checks -- the vibration analysis or the metallurgical assay -- flagged it.
A: The vibration analysis flagged the defect.
B: The metallurgical assay flagged the defect.
1. Reasoning: the response asserts the defect WAS caught by at least one check, so A and B cannot both be false; but both could hold (both checks may have flagged it). This is a disjunction, not an exclusion.
2. Discourse sense: Disjunction.
3. Coupling: at least one of A, B is true, i.e. co_necessity.
4. Final answer:
[sense=Disjunction] [coupling=co_necessity]

Your task:
RESPONSE: {{response}}
A: {{atom_a}}
B: {{atom_b}}
```

B.2 MINING PROMPT B — CONDITIONAL STRENGTH

Issued for each pair that prompt A coupled. The strength is the renormalized probability of the answer token, $P(\text{Yes})/(P(\text{Yes}) + P(\text{No}))$.

```
Instructions:
You are given the full RESPONSE, two claims A and B drawn from it, and a {{coupling}} relation that holds from A to B. Assuming A is TRUE, judge -- IN THE CONTEXT OF THIS RESPONSE -- whether the relation's implication about B is credible: at least plausible / more likely than not, NOT whether it is certain.

- entailment or equivalence: given A and how the response uses it, is B at least plausibly TRUE (more likely than not)?
- contradiction or exclusive: given A and how the response uses it, is B at least plausibly FALSE (more likely than not)? (For "exclusive", A and B are exhaustive alternatives, so A being true makes B false.)
- co_necessity: A and B are a pair of which at least one holds. Given the response rules out "neither", is it at least plausible that B holds when A does NOT?

Answer with a SINGLE WORD, the very first word of your reply: Yes or No.
- Answer "Yes" if the implication is credible/plausible (even if not certain).
- Answer "No" only if the implication is implausible or the response does not actually support it.
Do not output anything before the word Yes or No.

Example 1 (response supports a near-certain entailment):
RESPONSE: The new alloy is chemically identical to the certified reference alloy, so it meets the certified reference specification.
A: The new alloy is chemically identical to the certified reference alloy.
B: The new alloy meets the certified reference specification.
coupling: entailment
Answer: Yes

Example 2 (weak but plausible, and the response draws the link -- still Yes):
RESPONSE: The company launched a flawed product last quarter, and its stock price fell 15 percent over the same period.
A: The company launched a flawed product last quarter.
B: The company's stock price fell 15 percent last quarter.
coupling: entailment
Answer: Yes

Example 3 (clear contradiction stated by the response):
RESPONSE: The official statement said no one was harmed, but three people died in the incident.
A: No one was harmed in the incident.
B: Three people died in the incident.
coupling: contradiction
Answer: Yes

Your task:
RESPONSE: {{response}}
A: {{atom_a}}
B: {{atom_b}}
coupling: {{coupling}}
Answer:
```

B.3 G-EVAL COHERENCE JUDGE

The SummEval coherence rubric as used by G-Eval (Liu et al., 2023; Fabbri et al., 2021), adapted from summaries to responses, rated 1–5 and normalized.

```
You will be given one response written for a query.

Your task is to rate the response on one metric.

Evaluation Criteria:

Coherence (1-5) - the collective quality of all sentences. We align this dimension
with the DUC quality question of structure and coherence, whereby the response
should be well-structured and well-organized. The response should not just be a
heap of related information, but should build from sentence to sentence to a
coherent body of information about a topic.

Evaluation Steps:

1. Read the response carefully and identify the main topic and key points.
2. Check whether the response presents them in a clear and logical order, and
   whether later statements are consistent with earlier ones.
3. Assign a score for coherence on a scale of 1 to 5, where 1 is the lowest and 5
   is the highest.

Response:

{response}

Reason briefly, then give your rating as a single digit in square brackets, for
example [3].
```

B.4 DIRECT COHERENCE JUDGE

The no-rubric control.

```
Read the following response and rate its logical coherence -- how well it hangs
together as an argument, and whether it contradicts itself -- on a scale of 1 to 5,
where 1 is the lowest and 5 is the highest.

Response:

{response}

Give your rating as a single digit in square brackets, for example [3].
```

C THE PROMPT SPECIFICATION AND SUBSTITUTION CONTRACT

Each pipeline prompt's numbered-instruction count and required placeholder set are declared in a specification table, and a test asserts the declaration against the prompt strings.

Prompt	Characters	Instructions	Required placeholders
P1	1,902	9	`TOPIC`
P2	3,752	7	`QUESTION`
P3	7,778	10	`CLAIMS`, `QUESTION`
P4	5,571	8	`PLAN`, `QUESTION`
P5	2,592	5	`OPERATOR`, `PLAN`, `RESPONSE`
V1	4,086	0	`ATOMS`, `RESPONSE`
V2	2,455	0	`RESPONSE`
V3	1,403	0	`ATOMS`, `RESPONSE`

Instructions are counted by matching lines that begin at column one with a digit, a period and a space, so indented sub-items do not inflate the count. Placeholders are extracted by matching `[A-Z_]+_PLACEHOLDER` in brackets. Substitution enforces a *bidirectional* exact match between declared placeholders and supplied values: a missing name and an extra name are both errors, since a typo in either direction would otherwise ship a prompt containing a literal `[FOO_PLACEHOLDER]`, which a model will treat as content.

A validator that was specified and removed. The original design included a further validator to adjudicate exhaustiveness. It is deliberately absent: exhaustiveness is derived from the sense by the compile map, so the validator would have re-derived from prose what the scheme fixes by construction, and could have contradicted it. Exhaustiveness is instead checked by V1's sense-recovery rate, and by the validation study's both-false and both-true questions (§11.2).

D HUMAN-STUDY MATERIALS

D.1 PILOT ANNOTATOR BRIEF

The pilot's full written brief, given to every annotator before the first screen.

> **The task.** Each screen shows two responses, A and B, answering the same question. You decide which one is *more logically coherent*, or whether they are about equal.
>
> **What "logically coherent" means here.** A response is logically coherent when the things it says can all be true together, and the connections it draws between them hold up. Ask yourself: if I believed every claim in this response, would I be believing something consistent? Look for contradictions; connections that do not follow; tensions the text raises and then leaves hanging — versus tensions it raises and then settles, which is fine; and broken chains of reasoning.
>
> **What is not the target.** Do not reward length; fluency, style or polish; vocabulary or sophistication; or topic knowledge — a response can be perfectly coherent and factually wrong. Sentence order in itself does not make a response less coherent, unless the new order actually breaks a connection. That last point is a real question we are trying to answer, so please judge it honestly.
>
> **The three answers.** A; B; or about equal — you cannot tell them apart on coherence, or any difference is about style, length or wording. "About equal" is a real answer, not a cop-out; but if you see a difference in consistency, however small, pick a side and mark low confidence.
>
> **Confidence and reason.** Rate confidence low, medium or high, and give one sentence saying what drove the choice, ideally naming the specific claims or connection.

D.2 VALIDATION-STUDY BRIEFS

The two briefs of the validation study (§11.2), condensed. The full text, the Label Studio configurations and every task file are in the release's `docs/human_study_release`.

Core session (annotators C1–C7, about 2 hours 45 minutes).

> **The task.** Each task is one response, with its claims numbered underneath. Read it once, carefully; every question is about it, and needs no outside knowledge.
>
> **Part 1: claim pairs.** For each pair of numbered claims: does the response draw a logical link between them? A link means the text connects them — one causes, supports, explains, restates, contrasts with or rules out the other; mentioning the same topic is not a link. Then, *whatever the answer*: could both claims be true at the same time, and could both be false at the same time? Answer these two from the claims' meaning alone. A worked table covers the four cases: two different dates for one event (not both true; both could be false); a claim and its negation (neither); two possible causes (both); "at least one inverter failed" and "not every inverter failed" (both could be true; not both false).
>
> **Part 2: claims as stated.** For each numbered claim: does the response assert it as written, change or contradict it, or not state it? Judge only what the text says. Some claims are false in the real world on purpose; a false claim the response states is "asserted as written".
>
> **Part 3: edited versions.** Each version is the response with a change: removed text struck through in red, added text marked in blue. Is the version more coherent, about as coherent, or less coherent than the response you read? A response is coherent when what it says can all be true together and the connections it draws hold up. Did the change alter anything besides a link between claims — a fact, a number, what a claim means, or another connection? If so, choose minor or material and quote it. Some versions may change nothing.
>
> **Not the target.** Length, fluency, style, vocabulary, and real-world truth.

Instruction test (annotators I1–I3, about 2 hours 50 minutes).

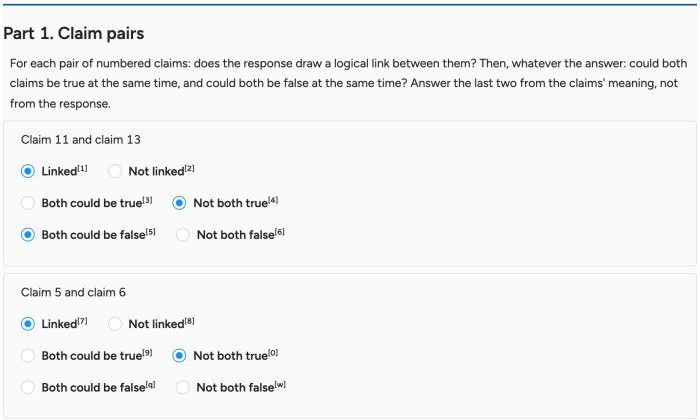
> **The task.** Each screen shows two responses with much the same content, and a question above them, highlighted. *The question is not the same on every screen*: read it each time and answer exactly what it asks. One asks which response would leave you believing something consistent, if you believed everything it says — whether the claims fit together, not how the response reads. The other asks which hangs together better as an argument — the response as a piece of writing. The two can have different answers for the same pair; that is what is being measured, so do not give the same answer to both on principle.
>
> **Your answer.** Response A, Response B, or about equal, which is a real answer; then confidence (low, medium or high) and one sentence saying why.

D.3 PILOT SCREEN SELECTION

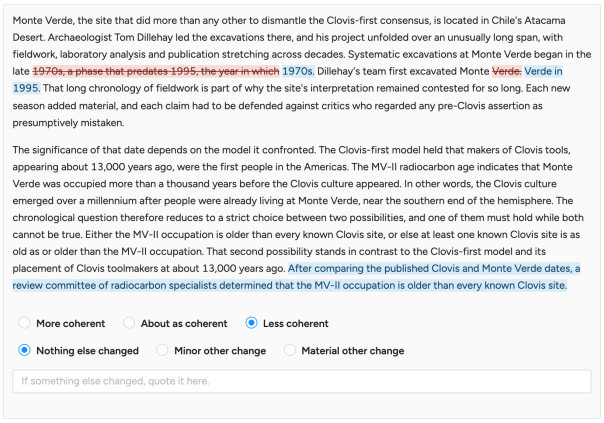
Family identifiers refer to the preliminary 14-family corpus of the pilot (§11.1), not to the release.

Family	Ladder	Topic	Pairs
f014, f013, f012	CONFLICT	Education, Economics, Cybersecurity	(ρ_0, ρ_1), (ρ_0, ρ_4)
f009, f005	CHAIN	Civil Engineering, Astronomy	(ρ_0, ρ_1), (ρ_0, ρ_4)
f007	ORDER	Botany	(ρ_0, ρ_1), (ρ_0, ρ_4)
f003	CONTROL	Art	(ρ_0, ρ_2), (ρ_1, ρ_2)

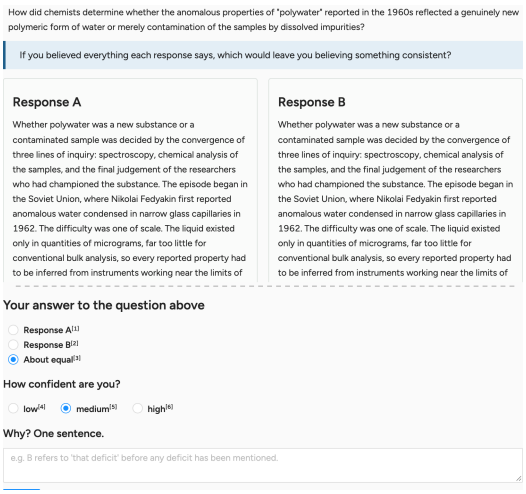
Families were chosen where the gold-graph measure and the strongest evidence-free baseline disagree most, including one (f012) on which the baseline wins. The CONTROL family's endpoint pair is byte-identical by design, so its two screens compare each meaning-preserving edit with the untouched base.



(a)



(b)



(c)

Figure 2: The validation study's interfaces in Label Studio, on the practice families, which are outside the study sample. (a) Two claim pairs from a core task, each with its link question and the two truth questions that recover its coupling. (b) An edited version, shown as the base response with removed text struck through and added text highlighted; only the paragraphs carrying a change are shown here. The highlight marks every textual change, including rewording that alters no relation — here the date clause is restructured as well as a resolving sentence added — which is why each version also asks whether anything besides a link between claims changed. (c) An instruction-test screen, with the screen's instruction in the band above the two responses; the responses are cut at the dashed line.